

From Causal Production to Structural Constraint:

Explanatory Roles, Emergence, and Ontological Stability



Lumen Feng

(冯鲁门)

2026-05

Contents

Copyright Page	3
Abstract	5
Chapter 1	7
<i>Introduction: Why Explanation Should Not Be Reduced to Causal Production</i>	
Chapter 2	13
<i>Explanatory Monism in Contemporary Debates: Literature Background and Problem Diagnosis</i>	
Chapter 3	47
<i>The Differentiation of Explanatory Roles: Causal Production, Structural Constraint, and Permissibility</i>	
Chapter 4	81
<i>The Reinterpretation of Emergence: Structural Novelty Rather Than New Causal Power</i>	
Chapter 5	100
<i>The Legitimacy of Higher-Level Explanation: Beyond Exclusion and Mechanistic Monopoly</i>	
Chapter 6	123
<i>The Scientific Standing of Non-Predictive Frameworks</i>	
Chapter 7	148
<i>Quantum Entanglement and the Misplaced Demand for Nonlocal Causation</i>	
Chapter 8	178
<i>Unknown Phenomena, Ontological Robustness, and Explanatory Discipline</i>	
Chapter 9	206
<i>Restrained Layered Structural Realism: From Explanatory Roles to Ontological Levels (with Discussion of “Radiance-Will Ontology”)</i>	
Chapter 10	225
<i>Conclusion: From Causal Centrism to the Multi-Role View of Explanation</i>	
Note to the Reader:	241
References	243
Appendix	256
<i>Appendix A: Criteria Table for the Three Explanatory Roles</i>	
Glossary	265
Postscript	270
About the Author	274

Copyright Page

Title:

From Causal Production to Structural Constraint: Explanatory Roles, Emergence, and Ontological Stability

Author:

Lumen Feng (冯鲁门)

ORCID:

<https://orcid.org/0009-0003-3716-6557>

GitHub:

<https://github.com/lumen-feng>

Edition:

Revised Edition 2.1

Publication Format:

Self-published open-access electronic monograph

Completion Date:

May 5, 2026

Copyright Notice:

This book is released under the Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License (CC BY-NC-ND 4.0). Reproduction and sharing are permitted on the premise of attribution, non-commercial, and no adaptation. Any commercial use, adaptation, translation, compilation, or republication requires the author's prior written permission. The cover image was generated and produced by the author and is distributed with this book under the same license.

Artificial Intelligence Assistance Notes:

This book was completed under the direction of a human author, with artificial intelligence tools used as assistance in the processes of writing, revision, structural organization, language refinement, logical checking, and format editing. Artificial intelligence tools do not constitute the author of this book and bear no intellectual, academic, legal, or ethical responsibility for it. All views, conceptual decisions, argumentative arrangements, textual review, and any possible errors in this book are the responsibility of the author alone.

Disclaimer:

This book is a theoretical work in philosophy and philosophy of science, intended to discuss such topics as theories of explanation, the problem of emergence, higher-level explanation, structural constraint, permissibility, and ontological stability. Its contents do not constitute professional advice in the natural sciences, law, medicine, investment, or any other field of practice. Readers should exercise independent judgment in light of their own disciplinary background.

Suggested Citation Format:

Lumen Feng (冯鲁门). From Causal Production to Structural Constraint: Explanatory Roles, Emergence, and Ontological Stability. Revised Edition 2.1, 2026.

Contact:

lumen_feng@pm.me

Abstract

This book examines a recurring pattern of explanatory contraction in contemporary philosophy of science: in debates over emergence, higher-level explanation, mechanistic analysis, quantum entanglement, and anomalous phenomena, explanation is often tacitly understood as the disclosure of a process of causal production. The present work does not deny the central importance of causal explanation. What it rejects is the elevation of causal explanation into the only model of explanation.

It begins by diagnosing this tendency as a form of explanatory monism, and distinguishes within it three levels: evaluative monism, role monism, and competitive monism. It then proposes a multi-role view of explanation, according to which explanation comprises three irreducible roles: causal production, structural constraint, and permissibility. Causal production answers how a result comes about; structural constraint answers which paths are limited or excluded; permissibility answers which patterns can stably exist under given constraints.

On this basis, the book reexamines emergence, higher-level explanation, non-predictive frameworks, quantum entanglement, and unknown phenomena, arguing that many traditional difficulties do not in fact require the postulation of new causal powers or ontological inflation, but instead arise from explanatory miswriting between distinct explanatory roles. Finally, it further argues that when structural constraint and permissibility display stability, indispensability, counterfactual constraint, and realizational support across multiple theoretical contexts, they should not be reduced to mere pragmatic convenience, but should instead support a restrained layered structural realism.

This book adopts “restrained layered structural realism” as the principal name of that position. “Radiance-Will Ontology” is retained only as a rhetorical secondary designation, meant to characterize the stable manifestation and organizational directionality of reality across different scales of organization.

The aim of this book is to show that theoretical maturity does not come from forcing all forms of explanation into a single model of production, but from rigorously differentiating among explanatory roles, methodological standards, and ontological commitments.

Chapter 1

Introduction: Why Explanation Should Not Be Reduced to Causal Production

In contemporary philosophy of science and foundational theory, many of the most persistent and recurrent disputes appear, on the surface, to belong to different domains. In fact, however, they share a deeper presupposition. Whether in debates over emergence and physicalism, in controversies surrounding higher-level explanation and causal exclusion, in defenses of mechanistic explanation as an explanatory paradigm, or in the ontological anxieties prompted by quantum entanglement, anomalous phenomena, and even unknown experiential objects, there is a widespread tendency to assume the following: explanation counts as explanation, ultimately, because it reveals a chain of causal production underlying some event. It is as if an account deserves the name of explanation only when it tells us how a certain result was produced; and any account that does not directly fit this model is treated either as a description of background conditions, as a merely provisional cognitive placeholder, or as something that must eventually be reduced to a more basic causal narrative.

This book begins from a suspicion toward precisely this default assumption. What is being challenged here is not the legitimacy of causal explanation itself, still less the central place of mechanisms, dynamical processes, and analyses of event-generation in scientific practice. What truly calls for scrutiny is a stronger claim: that causal production constitutes the only standard form of explanation, and that all legitimate explanations must therefore compete for standing within one and the same productive framework. This tendency may be characterized as a kind of causal centrism, or, more generally, as explanatory monism. It is not always advanced in the form of an explicit theoretical thesis, yet it plays a decisive role in many disputes. It determines which questions count as genuine questions, which

answers count as genuine explanations, and what kinds of theoretical work are regarded as methodologically admissible.

Once this assumption is accepted, a series of familiar difficulties follows. Emergence seems forced into a choice between two equally unsatisfying paths: if higher-level phenomena are genuinely new, then they must possess new causal powers; if they do not introduce new causal powers, then their alleged novelty is nothing more than an epistemic illusion. Likewise, in Kim-style causal exclusion arguments (Kim, 1998, 2005), once the causal sufficiency of the physical level is granted, higher-level properties seem to be left with no explanatory space at all; otherwise, one must appeal either to systematic overdetermination or to a weakening of physical closure. Much the same is true of many defenses of mechanistic explanation: mechanisms acquire an almost paradigmatic status precisely because explanation is tacitly presupposed to be the activity of revealing productive chains. In the case of quantum entanglement, the reason long-distance correlations continue to be treated as evidence for some kind of nonlocal influence is that correlation is automatically assumed to require explanation in terms of a productive connection propagating between relata. Even when confronted with unknown phenomena, anomalous experiential objects, or empirical novelty, people often slide quickly toward ontological inflation, as if the moment our existing causal picture cannot immediately absorb such objects, we must introduce new entities, new powers, or new ontological levels.

This book argues that these difficulties do not arise merely from a lack of empirical knowledge, nor are they simply the result of particular theoretical details. They recur because the roles of explanation have been contracted in advance. Many issues that ought to be understood in terms of structure, constraint, stability, permissibility, and modes of organization are instead miswritten as questions about how events are produced. As a result, accounts that do not belong to the same level are forced into the same competitive space; conceptual tension is

mistaken for ontological crisis, and explanatory mismatch is mistaken for a contradiction in the structure of the world itself.

To avoid this mismatch, the book proposes a multi-role view of explanation. Its basic idea is this: causal production is indeed an indispensable explanatory role, but it does not exhaust explanation as such. In addition to answering the question of how a given event is produced, many explanations in science and philosophy answer a different kind of question, namely, which configurations, trajectories, forms of organization, or outcomes are permissible for a system. Such accounts do not perform their explanatory work primarily by tracing productive chains. Rather, they explain by delimiting spaces of possibility, excluding unstable configurations, revealing relations of structural dependence, or showing why certain forms of organization possess stability. If such accounts are simply treated as incomplete causal analyses, then not only is their function misunderstood, but their independent standing in scientific understanding is obscured.

By structural constraint and permissibility, I do not mean to invoke some mysterious “non-causal force,” nor do I mean to deny physical closure, mechanistic analysis, or event-level causal inquiry. On the contrary, what this book seeks to defend is a stricter division of explanatory labor: different kinds of questions require different explanatory resources, and the legitimacy of an explanation should not be judged solely by whether it participates in one and the same productive competition. Causal explanation concerns how events occur; structural constraint concerns which paths, configurations, and states are limited or excluded; permissibility concerns which patterns can stably exist under given constraints. These three are not substitutes for one another, nor are they simply related as higher and lower levels. They cooperate across different explanatory roles. Only on this basis can we preserve the centrality of causal analysis while avoiding the compression of every problem into the causal paradigm, thereby generating many unnecessary theoretical dilemmas.

From within this general orientation, the book pursues several interconnected aims. First, it diagnoses how explanatory monism operates in a number of contemporary disputes and explains why it constitutes a problem deeper than any particular disagreement between substantive positions. Second, it introduces and clarifies a set of core distinctions, above all the difference between causal production, structural constraint, and permissibility, and argues that these are not rival forms of explanation but explanatory resources belonging to different roles. Third, it shows the explanatory yield of this framework through a series of representative problems, including emergence, the dispute over higher-level explanation and exclusion, the methodological legitimacy of non-predictive frameworks, the demand for nonlocal causation in quantum entanglement, and the impulse toward ontological inflation in response to unknown phenomena. Finally, it seeks to elevate the preceding analysis into a more systematic layered ontological expression, in order to clarify how the differentiation of explanatory roles requires a clearer ontological articulation of the relations among levels, structure, manifestation, and stability.

The central claims of this book may be stated on three levels. First, at the level of theory of explanation, causal production does not exhaust explanatory activity; structural constraint and permissibility also possess independent explanatory roles. Second, at the methodological level, the generation of novel predictions cannot serve as the sole criterion for evaluating every theoretical framework; some frameworks contribute instead by reorganizing problems, clarifying roles, and limiting improper inferences. Third, at the ontological level, if certain structures, boundaries, and stable patterns repeatedly display indispensability in explanation, then they should not simply be demoted to pragmatic conveniences, but should instead receive restrained ontological recognition.

The book does not, however, claim three things. First, it does not deny the central importance of causal explanation, mechanistic explanation, and

interventionist explanation. Second, it does not claim that structural constraint is an additional causal power. Third, it does not attempt to propose a new empirical theory to compete with existing physics, but instead seeks to clarify the relation between forms of explanation and ontological commitment.

This book does not purport to offer an empirical theory in competition with contemporary physics, nor does it seek to replace scientific inquiry with metaphysical construction. Its aim is more limited, yet also more fundamental: to clarify several conceptual dislocations that continue to trouble contemporary theoretical discussion. Its central judgment is that many so-called deep problems do not arise because the world itself forces us to choose between additional causal powers and explanatory nihilism, but because we assume too quickly that only causal production can count as genuine explanation. Once this assumption is relaxed, a more layered picture of explanation becomes possible, and many ontological anxieties can be correspondingly repositioned.

The argument of this book proceeds in four steps. First, it diagnoses how explanatory monism operates in contemporary disputes. Second, it introduces the threefold explanatory roles of causal production, structural constraint, and permissibility. Third, it applies this framework to emergence, higher-level explanation, non-predictive frameworks, quantum entanglement, and unknown phenomena. Fourth, it argues that these explanatory roles support a restrained layered structural realism rather than ontological inflation.

The structure of the book is as follows. Chapter 2 surveys the relevant contemporary literature and shows how explanatory monism is embedded in exclusion arguments, the prioritization of mechanistic explanation, and certain methodological positions, while also explaining why current research on non-causal explanation, constraint, and structural forms of account still awaits further integration. Chapter 3 presents the conceptual core of the book, offering a systematic distinction among causal production, structural constraint, and

permissibility, together with the corresponding criteria and boundary clarifications. Chapters 4 and 5 address, respectively, emergence and higher-level explanation, showing why the reality and explanatory standing of higher-level phenomena do not depend on the introduction of new causal powers. Chapter 6 defends this framework at the methodological level by arguing that non-predictive frameworks can still play a legitimate role in scientific understanding. Chapters 7 and 8 apply the framework to quantum entanglement and unknown phenomena, respectively, showing how it can ease both the demand for nonlocal causation and the tendency toward ontological inflation. Chapter 9, on that basis, attempts to articulate a fuller stratified ontological formulation in order to integrate the relations among structure, level, stability, and manifestation. Finally, Chapter 10 reviews the book's principal arguments, assesses both their theoretical gains and their limitations, and points toward possible directions for further research.

If the central thesis of this book had to be compressed into a single sentence, it would be this: several characteristic deadlocks in contemporary philosophy of science arise not merely from a lack of empirical data or from theoretical conflict, but from reducing explanation to causal production; once we acknowledge that structural constraint and permissibility possess independent explanatory roles, problems such as emergence, higher-level explanation, quantum entanglement, and ontological stability can be redescribed and handled more clearly without appealing to additional causal powers or ontological inflation. The task of the chapters that follow is to argue for this claim step by step.

Chapter 2

Explanatory Monism in Contemporary Debates: Literature Background and Problem Diagnosis

If the core claim of this book is that, in a number of key debates, once the standing of an explanation comes under pressure, a productive standard of adjudication tends to regain priority, such that an account is easily treated as incomplete, subordinate, or still awaiting reduction if it cannot show how a result is produced, then the first thing to clarify is that this judgment is not a simple rebuttal of one particular school, nor a denial of the overall value of the causal-explanatory tradition. The problem runs deeper, and is more concealed. It is not that philosophers or scientists explicitly maintain that only causal explanation counts as explanation. Rather, the problem is that the dominant disputes across many different fields, however different their overt aims may be, share at decisive moments a more basic explanatory orientation: any explanation with genuine weight must ultimately tell us how a given result is produced; and if an account cannot perform this task, then its status can be, at most, auxiliary, provisional, or still awaiting reduction.

This orientation is especially worth diagnosing precisely because it usually does not appear as a central thesis defended in the open, but operates instead as a background condition of dispute. It determines which questions will be raised, which answers will be judged insufficient, and why many accounts that do not belong to the same level are nevertheless forced into the same competitive space. For that very reason, the difficulties that recur across different topics, though different in form, often share the same structure: once some phenomenon is acknowledged as real, stable, and explanatorily significant, the next question is immediately compressed into whether it possesses an independent causal power;

and if the answer is no, then its explanatory status, ontological status, and even theoretical status are correspondingly weakened. This is true of higher-level properties, of emergence, of structural regularities, of quantum correlations, and likewise of ontological reactions to empirical novelty.

To understand how this problem arises, we must first return to several representative traditions in contemporary philosophy of science. The first unavoidable point of reference is the discussion of higher-level causation and exclusion centered on Jægwon Kim (Kim, 1998, 2005). Kim's work has had enduring influence not merely because it posed a sharp challenge to nonreductive physicalism, but because it systematized a particular explanatory intuition. On this line of thought, if a physical state is already causally sufficient for a result, then any corresponding higher-level property, if it is still said to explain that result, comes under pressure: is it doing redundant work, or is it claiming some unacceptable additional causal efficacy? What gives the exclusion argument its real force here is not only the principle of causal closure, but the fact that it converts causal sufficiency into explanatory sufficiency with almost no argument. Once the productive chain at the physical level is treated as complete, higher-level accounts seem to lose any remaining space.

It is precisely this feature that allows Kim's challenge to extend far beyond philosophy of mind or disputes over higher-level causation as such. In effect, it provides a paradigm for a broader explanatory monism: once one level has already delivered a causally sufficient account, accounts at other levels are left to choose among redundancy, competition, or the status of weak heuristic language. Many debates over higher-level autonomy, higher-level patterns, and even the stratified organization of science find it difficult to escape the dilemma of either additional causal power or merely superficial language for exactly this reason: they have inherited this explanatory schema.

A second major line of literature that must be brought into view is the tradition of mechanistic explanation, especially the developments associated with Craver and those influenced by him (Craver, 2007; Machamer, Darden, and Craver, 2000). The rise of mechanistic explanation greatly heightened philosophers' sensitivity to organization, components, activities, and hierarchical composition in scientific practice, and it successfully advanced a renewed understanding of what it means for explanation to stay close to actual scientific work. In many contexts, mechanistic explanation is attractive because it seems to avoid empty law-like accounts and instead demands that explanation show how a concrete structure produces a phenomenon through the organized activity of parts and operations. In this respect, the mechanistic tradition has genuine value in correcting overly abstract conceptions of explanation.

Yet the tradition of mechanistic explanation also introduces another risk. Mechanists need not explicitly claim that all explanation should be reduced to mechanistic account, but in actual discussion mechanistic explanation is often treated as a default paradigm, as if any account that cannot take the form of displaying a production mechanism were thereby explanatorily weaker. In that way, accounts that might instead be understood in terms of constraint, boundary conditions, structures of stability, or organizational permissibility are easily treated as incomplete mechanistic analyses or as merely peripheral supplements to real explanation. The problem is not that mechanistic explanation is unimportant. The problem is that it is tacitly treated, in many discussions, as the standard form of explanation. As mechanistic analysis becomes more successful, the independent standing of other kinds of explanation becomes harder to see.

A third indispensable line of literature is the causal-explanatory tradition represented by Woodward (Woodward, 2003). Compared with the first two, Woodward's work offers more refined and more powerful resources, because it does not rest content with understanding explanation as a simple causal narrative

in temporal sequence, but instead seeks to clarify explanatory structure through intervention, counterfactual dependence, and manipulability. Precisely through these contributions, causal explanation in contemporary philosophy of science has taken on a much more sophisticated form than the older covering-law model. Many scientific accounts do indeed gain explanatory force by showing how intervention on certain variables would alter others.

But it is precisely because the causal-explanatory tradition has developed in a more mature way at this stage that its boundary problem becomes all the more important to examine carefully. Even if one grants that counterfactual dependence and intervention analysis are vital for a great many scientific explanations, one must still ask: must every legitimate explanation ultimately be rewritten as some picture of interventionally tractable dependence? When scientists appeal to symmetry, conservation, topological limitation, phase-space structure, boundary conditions, or global forms of organization, are they always engaged in causal work that remains incomplete? Or do such accounts themselves perform distinct explanatory roles? If the latter is true, then the issue is no longer whether causal explanation is correct, but whether it has been generalized too far.

It is at this point that another group of literary resources becomes important, namely, work on non-causal explanation, constraint explanation, symmetry explanation, and structural forms of account. Represented by Batterman, Lange, and related discussions (Batterman, 2002, 2021; Lange, 2017), this line of research reminds us that scientific understanding does not always depend on tracing the generative chain of particular events. Sometimes an explanation succeeds because it shows that certain results are governed by higher-level limiting conditions, or because it demonstrates that a space of possibilities has itself already been structurally filtered. Conservation laws do not “push” events into occurring, yet they can explain why certain evolutions cannot occur; symmetry principles do not intervene as local causes, yet they can explain why certain relations must hold;

certain topological properties do not function as mechanistic components, yet they can explain why a system exhibits a particular stable form.

This body of work is crucial, because it shows that suspicion toward causal centrism is not without precedent. Existing literature has repeatedly touched upon the fact that explanatory practice in science is far richer than any single productive model. Yet the judgment of this book is that these important works, while opening conceptual space, have not fully accomplished a further task: integrating these scattered insights into a cross-topic, cross-level explanatory framework. They often concentrate on a particular type of account, such as symmetry explanation, mathematical explanation, or non-causal explanation in specific physical cases, but they less often show systematically why debates over emergence, higher-level explanation, quantum entanglement, ontological inflation, and methodological legitimacy may all be re-understood under the single issue of explanatory roles.

That is precisely where this book seeks to intervene. Rather than inventing an entirely new type of explanation, it is better understood as reorganizing a number of resources that have already been partially recognized but not yet systematically connected. More specifically, the book argues that contemporary debates repeatedly fall into deadlock not only because the relevant concepts lack sufficiently clear definition, but because the boundaries among different explanatory tasks have not been adequately marked. On the one hand, people acknowledge the importance of structure, constraint, level, and stability for scientific understanding; on the other hand, at the decisive moment they still habitually ask: through what causal mechanism do these things operate? It is precisely when this question is tacitly granted the highest adjudicative authority that explanatory monism returns to center stage.

Accordingly, the task of this chapter is not simply to divide the existing literature into a “mistaken camp of causal explanation” and a “correct camp of non-causal explanation.” Such a division would be neither accurate nor helpful. A more

precise approach is to note that contemporary literature has already supplied several powerful but disconnected resources. Kim reveals how the tension between higher-level explanation and causal closure can be sharpened; the tradition of mechanistic explanation displays an important form of productive account in scientific practice; Woodward gives causal explanation a highly refined counterfactual analysis; and Batterman, Lange, and related work show that constraint, symmetry, and structural conditions can themselves perform indispensable explanatory functions. What this book does is to show, without denying the contribution of any of these traditions, that they remain insufficient to dissolve several central disputes because they lack a more general framework of explanatory roles, one capable of showing how different kinds of account can be legitimate, coexist, and yet need not compete within a single causal space.

The analysis that follows proceeds from this diagnosis. First, it is necessary to explain more precisely what explanatory monism consists in and whether it is merely a habitual preference or a deeper structure with genuine theoretical consequences. Second, it must be explained why existing discussions of non-causal explanation, while highly valuable, still do not suffice to address the systematic confusions surrounding higher-level explanation, emergence, and ontological stability. Finally, it must be shown why a book that seeks to connect these issues cannot rest content with securing legitimacy for structural accounts in a few isolated cases, but must instead offer a more stable set of conceptual criteria for distinguishing between cases in which we are genuinely dealing with problems of causal production and cases in which what really requires explanation is permissibility, form of organization, and structural constraint.

2.1 The Three Levels of Explanatory Monism

If many contemporary disputes do indeed share an excessive contraction of explanation, the next question is where exactly this contraction occurs. Simply

saying that explanation has been understood in an overly causalized way remains too vague. In many cases, the relevant theorists may not explicitly deny the importance of structure, constraint, stability, or forms of organization; nor may they directly declare that no mode of account exists other than causal production. The real problem is that, once a dispute enters its decisive stage, one type of question is repeatedly granted final adjudicative authority over all others: can this account show how something is produced? If not, then however useful, insightful, or even indispensable it may otherwise be, its standing tends to be lowered.

To bring this structure into clearer focus, this book distinguishes three levels at which explanatory monism operates: evaluative monism, role monism, and competitive monism. These three levels are related, but not identical. They differ in force, in function, and in the way they shape philosophical judgment. Distinguishing them is important because many theoretical confusions persist precisely by moving back and forth among these levels without making the transitions explicit.

The first level is evaluative monism. This is the tendency to assess all explanations by a single standard of explanatory excellence, namely, the extent to which they reveal a process of causal production. On this view, even if one allows that different kinds of account may be useful in practice, the real explanatory ideal remains singular: the best explanation is the one that most fully shows how a result is generated. Other forms of account may be tolerated, but only as weaker substitutes, heuristic devices, or temporary stand-ins pending a more genuinely productive explanation. Evaluative monism does not yet say that all explanations are literally the same kind of thing. What it does say is that they are all to be ranked by reference to one privileged model.

The second level is role monism. Here the contraction goes further. It is no longer merely that all explanations are judged by one evaluative standard; rather, explanation itself is treated as if it performed only one basic task. That task is to

answer how some event, state, or outcome is brought into reality. Once explanation is understood in this way, accounts concerned with constraint, boundary, organization, stability, or permissibility are no longer seen as performing distinct explanatory functions. They are reinterpreted instead as incomplete attempts to answer the same production-oriented question. Role monism is therefore more than a preference for causal explanation. It is a prior narrowing of what explanation is taken to be.

The third level is competitive monism. This is the view that all legitimate explanations must compete for the same explanatory position. If causal production is treated as the sole genuine explanatory role, then any higher-level, structural, or permissibility-based account can survive only by claiming to do the same work as a lower-level causal account. The result is a zero-sum space of competition: either a higher-level account introduces something causally additional, or it is explanatorily redundant. Either a structural account discloses a real productive factor, or it is merely decorative. Competitive monism is what transforms differences in explanatory role into apparent conflicts over who has the right to explain.

These three levels often reinforce one another. Evaluative monism privileges causal-production accounts as the highest standard; role monism recasts every explanatory task in production-oriented terms; competitive monism then forces all candidate explanations into a single arena of rivalry. Yet analytically they should be kept apart. A theorist may, for example, admire causal explanation as the highest ideal without explicitly denying that other explanations serve different roles. Another may verbally acknowledge multiple forms of account, yet still treat them as explanatorily inferior because they fail to satisfy a single production-based measure. Still another may accept that different explanatory tasks exist in principle, while continuing in practice to stage all disputes as if explanations had to compete for a single slot. Only by separating these levels can we see exactly where the contraction occurs and how it exerts pressure.

This tripartite distinction also helps explain why explanatory monism can persist even in theoretically sophisticated environments. A philosopher need not openly proclaim that “only causal explanation is real explanation” in order to reproduce the structure of explanatory monism. It is enough that, when a dispute becomes pressing, causal production reappears as the ultimate court of appeal. Once that happens, structural accounts are demoted to background description, permissibility-based accounts are treated as merely modal glosses, and higher-level explanations are pressed to justify themselves either by claiming independent causal power or by conceding inferiority. The monism may therefore remain implicit while still exercising real theoretical force.

It is important to see that the problem here is not the prominence of causal explanation as such. Causal explanation is indispensable, and in many domains it rightly occupies a central place. The problem is the unmarked slide from indispensability to exclusivity, from centrality to monopoly. Once that slide occurs, distinctions among explanatory tasks are flattened, and disagreements that should concern explanatory division of labor are recast as struggles over explanatory legitimacy itself.

For this reason, the present book does not seek merely to add one more kind of explanation to an already crowded field, as if the point were simply to enlarge an inventory. Its more basic aim is to resist the contraction by which different explanatory roles are collapsed into one, and by which distinct questions are compelled to answer to a single productive tribunal. The purpose of distinguishing evaluative monism, role monism, and competitive monism is thus diagnostic as much as theoretical: it reveals that many of our familiar deadlocks are not generated only by substantive disagreement about the world, but by a prior and often unnoticed narrowing of what explanation is allowed to be.

The sections that follow will build on this diagnosis. Once explanatory monism is disaggregated into these three levels, it becomes easier to understand why a

number of recurring philosophical pressures arise with such regularity, and why the solution cannot consist merely in defending this or that favored case of non-causal explanation. What is required instead is a more systematic rearticulation of explanatory roles themselves.

2.2 The Advances and Limits of Non-Causal Explanation, Constraint, and Structural Accounts

If the analysis of the previous section has already shown that explanatory monism is not an accidental prejudice of some isolated theory, but a deep structure that repeatedly operates across multiple disputes, then the next question is whether contemporary literature already contains resources sufficient to break that structure. In one sense, the answer should be yes. Over the past several decades, research on non-causal explanation, constraint explanation, symmetry explanation, mathematical explanation, and structural accounts has significantly broadened the philosophy of science's understanding of explanation. At the very least, these studies jointly suggest one important point: scientific explanation does not always do its work by tracing the generative chain of events. In many contexts, an explanation succeeds because it reveals why certain results are necessary, why they are stable, or why they are confined within a particular space of possibilities, rather than because it shows how some local influence is transmitted step by step to the outcome.

The importance of this line of work should not be underestimated. It means that suspicion toward causal centrism is not an external challenge to scientific practice, but something that grows from within scientific practice itself. Scientists have long made extensive use of forms of account that cannot be adequately described simply as "the search for generative mechanisms." Conservation laws explain why certain processes cannot occur; symmetry principles explain why certain relations remain invariant across a wide range of conditions; topological

features explain why certain systems exhibit stable behavior that does not correspond one-to-one with particular microscopic paths; boundary conditions and global forms of organization explain why certain local dynamics can unfold only within limited structural ranges. All of these show that scientific understanding comes not only from knowing what drives what, but also from knowing what is permitted, what is excluded, and what can stably exist.

For that reason, the major contribution of Batterman, Lange, and related discussions lies not merely in providing alternative descriptions for a number of special cases. More fundamentally, they unsettle a deeper assumption: not every explanatory problem must be translated into a problem of event-production. Some explanations are powerful precisely because they bypass local productive chains and appeal directly to higher-level limiting relations, global structures, or conditions of invariance. In this respect, these studies have already succeeded in showing that the actual terrain of scientific explanation is broader than any single causal or mechanistic picture.

Yet the judgment of this book is that, although these important works open conceptual space, they have not fully completed another, more systematic task. They typically concentrate on certain specific classes of cases, such as mathematical explanation, symmetry explanation, universality explanation, or non-mechanistic accounts in particular areas of physics, and use these cases to rebut the overly strong thesis that all explanations must be causal explanations. This is of course a crucial advance. But it still leaves two interconnected problems behind.

The first problem is that these studies usually enter the discussion in the form of exceptional cases. They successfully show that, in some settings, non-productive accounts also possess explanatory force; but they often do not go on to ask whether this means that explanatory activity itself contains several fundamentally distinct roles. In other words, they are usually arguing for the legitimacy of certain special forms of account rather than systematically reconstructing the task-structure

of explanation. As a result, these achievements may loosen evaluative monism without fully overturning role monism. A reader may readily accept that some explanations are not causal, while still regarding them as marginal or exceptional, without thereby altering their overall understanding of what explanation itself is.

The second problem is that these studies generally do not respond in any unified way to the disputes that most readily generate ontological pressure and pressure about levels. Even if one grants that symmetry, conservation laws, or mathematical structure can explain certain phenomena, one may still continue to think that, in the case of emergence, the real issue is whether higher-level phenomena possess independent causal power; that, in the case of exclusion, the real issue is whether higher-level properties compete with lower-level physical causes; that, in the case of quantum entanglement, the real issue is whether long-distance correlations correspond to some action at a distance; and that, in the face of unknown phenomena, the real issue is whether our existing ontological inventory is no longer sufficient. In other words, research on non-causal explanation has shown the possibility of alternative forms of account, but it has not automatically dissolved the recurring “causal adjudicative authority” that reappears across these cross-domain disputes.

This is exactly where the present book seeks to advance the discussion. It is not satisfied merely with pointing out that explanations in some special cases are non-causal. It goes further and argues that a number of long-standing deadlocks in the philosophy of science persist precisely because they are all placed within an insufficiently examined structure of explanatory monism. Put differently, the issue is not simply whether non-causal explanations exist, but rather this: when we discuss emergence, higher-level explanation, structural stability, quantum correlation, or unknown phenomena, what kind of question are we actually asking? If these questions are not primarily asking from the outset how some result is produced,

then continuing to use a productive standard as the final tribunal will systematically generate distortion.

From this perspective, the limits of existing research do not show that it is misguided. They show, rather, that it has not yet been organized into a more general theory of explanatory roles. It has already supplied important resources for this book: first, it shows that explanatory practice does indeed extend beyond causal production; second, it shows that structure, limitation, invariance, and forms of organization can have genuine explanatory force; third, it shows that our understanding of what explanation is must itself be expanded. But if these resources are to intervene effectively in disputes over emergence, higher-level autonomy, exclusion, entanglement, and ontological inflation, two further tasks remain to be completed. The first is to provide clearer conceptual criteria for distinguishing what productive accounts and constraint-based accounts are respectively answering. The second is to dissolve the presupposed competitive relation between explanations at different levels.

It is precisely because these two tasks have not yet been fully carried out that the existing literature, although it has made “explanation is not merely causal production” into a proposition that can be taken seriously, has not yet turned it into a theoretical framework capable of reorganizing multiple domains of dispute. As a result, explanatory monism continues to survive in a more concealed form in many crucial debates. One may acknowledge that structure matters, that constraint matters, that boundary conditions matter, and yet still ask, at the decisive moment: through what mechanism do these factors operate? Once that question is tacitly treated as the most fundamental one, structural accounts are once again pulled back into the productive framework.

Accordingly, the relation between this book and existing work on non-causal explanation should not be understood as one of simple opposition. The book does not seek to reject the achievements of that work. On the contrary, it moves forward

within the space those achievements have already opened. What it tries to do is to reorganize insights that were previously scattered across different cases and discussions into a more stable framework of explanatory roles, and then to use that framework to address systematically several of the most characteristic disputes in contemporary philosophy of science. Only when that systematic work has been completed will “structural constraint” and “permissibility” cease to be merely linguistic choices in a handful of exceptional cases and instead become explanatory roles that stand alongside causal production without competing with it.

This also shows why the next chapter must turn to conceptual reconstruction. Without clearer criteria, it is difficult to say in a genuinely rigorous way when an account is primarily tracing event-generation, when it is primarily characterizing a space of possibilities, stable structures, or organizational limits, when the two kinds of account can work together, and when we have merely miswritten one kind of problem as another. In exactly this sense, Chapter 3 is not an optional supplement to the preceding review of literature, but a necessary turning point in the argument as a whole. The task of the first two sections has been to show that the problem is real, and that existing literature, while providing crucial resources, has not yet fully resolved it. The task of the next chapter is to propose a framework of distinctions clear enough to organize those resources into a continuous theoretical path.

2.3 Kim and the Problem of Higher-Level Causal Exclusion

In contemporary discussions of higher-level properties, nonreductive physicalism, and explanatory stratification, Jaegwon Kim’s exclusion argument undoubtedly functions as a central point of reference. Its importance lies not only in the fact that it poses a technical difficulty for higher-level causation, but also in the fact that it displays, in an especially clear form, the tension among causal closure, reductionist pressure, and explanatory space. More importantly, Kim’s argument has remained

persuasive over time not simply because it relies on certain metaphysical premises, but because it organizes a deeper explanatory intuition into a systematic argument. In this sense, the exclusion argument is not merely one position within the debate over higher-level causation; it is also a paradigmatic expression of explanatory monism within questions of levels.

Kim's basic line of thought is well known. Suppose that a higher-level property *M* is realized by some physical property *P*, and that *M* is taken to have causal efficacy with respect to some result *E*. If the physical domain is causally closed, then *P* is already sufficient to bring about *E*. At that point, the familiar difficulty emerges: if *M* is distinct from *P*, and *M* is also said to be a cause of *E*, then *E* appears to be systematically overdetermined. In order to avoid that result, one must either deny the independent causal efficacy of *M*, weaken physical closure, or in some way reduce *M* to *P*. The sharpness of Kim's challenge lies precisely in the fact that it forces nonreductive physicalism to choose among several options, none of them especially attractive.

The force of this argument is usually taken to derive from the principle of physical closure itself. But if that were all there was to it, the influence of the exclusion argument would not in fact be so wide-ranging. For causal closure, taken by itself, shows first of all only this: along the productive chain of events, the physical level has already supplied sufficient conditions. Yet the point at which the exclusion argument generates much broader philosophical pressure lies in the fact that it tacitly completes a further step, namely, the slide from causal sufficiency to explanatory sufficiency. Once that slide is accepted, the standing of higher-level accounts changes immediately. The question is no longer merely whether higher-level properties can still exist as additional causes; rather, all higher-level accounts begin to appear suspect: if the lower-level productive chain is already complete, what explanatory work is left for higher-level explanation to do?

It is precisely here that Kim's argument reveals its deep connection with explanatory monism. If explanation is understood as a single type of activity, and if the core task of that activity is to reveal how some result is produced, then the conclusion of the exclusion argument follows almost naturally. Once the productive task has already been completed at the lower level, it becomes difficult not to regard accounts at other levels as redundant, competitive, or merely pragmatically convenient forms of expression. Put differently, the ease with which the exclusion argument expands into a general suspicion toward the standing of higher-level explanation does not arise only because it is concerned with causation. It arises because it presupposes a strong isomorphism between causal production and explanatory task.

This can be seen in many subsequent discussions. Higher-level properties are often required to show why they are not redundant; and the most direct way of showing that they are not redundant is then understood as proving that they make some independent causal contribution. As a result, the whole structure of the debate is compressed into a narrow set of options: either higher-level properties have causal powers of their own, or they are merely convenient labels for lower-level physical states; either one accepts some strong form of higher-level causation, or one admits that higher-level language lacks genuine explanatory content. The problem is that this compression does not follow directly from the question of whether higher-level phenomena are real. It is driven instead by a more prior explanatory setup, one that assumes in advance that if some level is to have theoretical standing, it must occupy a place within productive explanation.

That requirement, however, is itself open to question. The reason higher-level accounts continue to play an irreplaceable role in scientific and philosophical practice is not obviously that they are always adding some further productive chain. Many higher-level accounts matter because they characterize organizational structure, pattern stability, systemic constraint, functional dependence, or

conditions of realization, and thereby answer another class of question: why a certain pattern can stably exist over a given physical realization, why certain micro-level differences do not destroy a higher-level property, and why certain higher-level behaviors retain systemic consistency despite variation in details of implementation. If these questions are not themselves questions about how a result is additionally produced, then requiring higher-level accounts to secure legitimacy by winning independent causal power is a misdirected demand.

From this perspective, Kim's exclusion argument certainly captures a real problem, but at the same time it magnifies a deeper confusion. The real problem is this: given physical closure, how are we to understand the standing of higher-level accounts without turning them into containers for mysterious causal powers? But the confusion it magnifies is the suggestion that, unless higher-level accounts take on an additional productive task, they no longer possess genuine explanatory standing. The position of this book is that the former is a problem worth taking very seriously, whereas the latter is an additional pressure imposed on that problem by explanatory monism. In other words, the exclusion argument forces us to clarify higher-level explanation, but it does not show that higher-level explanation can gain legitimacy only in the form of an "additional cause."

This also explains why many responses to Kim, though each illuminating in its own way, have often found it difficult fully to dissolve the pressure of the exclusion argument. Some responses try to redefine realization in order to weaken the contrast between M and P; some try to articulate a noncompetitive understanding of higher-level causation; others shift emphasis toward the indispensability of multilayered description in explanatory practice. Yet so long as these responses do not explicitly challenge the slide from causal sufficiency to explanatory sufficiency, they are often pulled back into the original framework. Higher-level accounts may be acknowledged to have value in heuristic work, description, or model

construction, but their “genuine” explanatory standing remains suspended, because the standard of adjudication still lies in the productive picture.

It follows that Kim occupies an important place in this book not because the book must refute each of his technical premises one by one, but because his argument reveals a more general structure: when explanation is presupposed to be a single productive task, stratified accounts are almost inevitably driven into competition. Once lower-level explanation is taken to be complete, higher-level accounts can do little but oscillate between redundancy and usurpation. The exclusion argument is therefore not an accidental theoretical conclusion, but a natural product of explanatory monism within questions of levels.

This does not mean that Kim’s work is simply mistaken. On the contrary, its philosophical value lies precisely in the fact that it forces us to see something important: unless explanatory roles are redefined, higher-level explanation will indeed continue to face recurrent crises of legitimacy. For that very reason, the present book’s response to exclusion will not proceed along the simple line that “higher-level explanation is useful in practice, therefore exclusion must be wrong.” Practical usefulness by itself cannot automatically settle the issue of theoretical standing. What must really be shown is that higher-level explanation is legitimate not because it competes with lower-level physical explanation for the same productive position, but because it answers a different class of questions concerning organization, stability, realizational constraint, and the persistence of patterns. Only when that is made clear does the deadlock generated by the exclusion argument begin genuinely to loosen.

Kim’s place in this book is therefore double. On the one hand, he represents the strongest form of pressure faced by higher-level explanation. On the other hand, he helps us identify where that pressure actually comes from. Without the sharpening effect of the exclusion argument, the role played by explanatory monism in debates over levels might have remained harder to discern. The next

section, which turns to the tradition of mechanistic explanation, will make this even clearer. Unlike the exclusion argument, mechanistic explanation does not always appear by excluding higher-level accounts outright, but it can reinforce the same tendency along a different path: the tendency to think that genuine explanation simply is the display of a productive process. It is precisely between Kim and mechanistic explanation that what this book calls explanatory monism reveals itself as capable of taking both adversarial forms and apparently milder forms that nonetheless have the same contracting effect.

2.4 The Tradition of Mechanistic Explanation and the Expansion of Productive Accounts

If Kim's exclusion argument pushes higher-level explanation into the dilemma of either having causal power or lacking standing, then the tradition of mechanistic explanation exerts an equally far-reaching, though more moderate, influence on explanatory monism along a different path. Unlike the adversarial structure of the exclusion argument, mechanistic explanation does not seek to eliminate higher-level accounts. More often, it presents itself as offering a more faithful characterization of the nature of explanation precisely because it remains closer to actual scientific practice. For that very reason, its shaping effect on our picture of explanation is often harder to notice as a form of restriction.

Associated above all with Craver and later developments, mechanistic explanation emphasizes that the core of scientific explanation lies in showing how a system composed of entities, activities, and their organization produces a phenomenon. Within this framework, a satisfactory explanation is expected to display how the parts of a system work together through specific organizational relations in order to generate the behavior to be explained. Compared with the traditional covering-law model, this orientation is plainly closer to real scientific

work, and it better captures why certain explanations possess the advantage of rich detail and explicit structure.

Yet it is precisely within this success, this success in being “close to scientific practice,” that mechanistic explanation gradually comes to occupy the position of a default paradigm. Mechanists do not usually declare outright that all explanations must be mechanistic explanations, but in many concrete discussions mechanistic structure is treated as the standard form of explanation: an account is judged incomplete because it has not yet displayed the relevant mechanism, and once a mechanism has been identified, the explanation appears to have reached a kind of endpoint. This tendency does not arise from an explicit theoretical manifesto. It is reinforced instead through patterns of evaluative practice.

Within this setting, explanatory monism no longer needs to sustain itself by denying the existence of other forms of account. It can operate, rather, through absorption. Structural constraint, boundary conditions, stability analysis, and even certain abstract patterns are reinterpreted as peripheral components of mechanistic explanation. A given constraint matters because it limits the paths a mechanism may take; a given stable structure matters because it reflects the steady-state outcome of a mechanism under a particular organization. In this way, material that might otherwise constitute an independent explanatory role is recoded as supplementary to mechanistic analysis.

The problem is not that this integration is always mistaken. In many cases, bringing constraints and structures into a mechanistic framework genuinely helps us understand how a system operates. The problem begins when this integration is generalized. At that point, an important difference starts to disappear: not every account of constraint and structure is serving the more detailed characterization of a productive process. Some such accounts function primarily not to tell us how a system generates a result, but to tell us why only certain generative paths are possible, or why certain outcomes remain stable across a wide range of conditions.

If this difference is neglected, then the boundaries between explanatory roles are once again flattened.

The pressure mechanistic explanation generates here differs in form from Kim's exclusion argument, yet it points toward a similar result. The exclusion argument uses a competitive structure to force higher-level accounts either to become additional causes or to lose standing; mechanistic explanation, by contrast, makes it difficult for other kinds of account to appear as independent explanations at all, since they are absorbed into the components of mechanistic analysis. In the former case, higher-level explanation is compressed into the role of a causal competitor. In the latter, it is transformed into just another layer of mechanistic description. What the two cases share is the maintenance, to varying degrees, of the assumption that genuine explanation must ultimately revolve around productive structure.

This is especially visible in discussions of scientific stratification. Mechanistic explanation often acknowledges the existence of multilevel structure and emphasizes compositional relations between levels. Yet such acknowledgement does not automatically imply that different levels possess different explanatory roles. On the contrary, the levels are often understood as descriptions of one and the same mechanism at different degrees of granularity. As a result, the standing of a higher-level account still depends on the extent to which it can be embedded within a lower-level mechanism. If some higher-level description cannot be aligned with a clear mechanistic structure, it is easily treated as not fully "genuinely explanatory." Within this framework, differences of level are translated into differences of descriptive resolution rather than differences in the kinds of question being asked.

It is precisely here that the connection between mechanistic explanation and explanatory monism becomes clear. The issue is not whether mechanistic explanation is correct, but whether it has been overgeneralized into the only model of explanation. Once that overgeneralization occurs, the space of explanatory roles

is contracted once again: all accounts are understood as participating, to varying degrees, in a single task, namely, the disclosure of some generative structure. As a consequence, accounts primarily concerned with constraint, permissibility, and stability can no longer easily be recognized as performing another task. They are instead interpreted as incomplete mechanistic analyses.

This diagnosis also helps explain why, despite its great methodological success, mechanistic explanation has not resolved the theoretical difficulties faced by higher-level explanation and structural accounts. In some cases, it even intensifies them. Once mechanism becomes the default template of explanation, higher-level accounts find it harder to escape the demand that they either display their corresponding mechanistic basis or concede that they are merely descriptive or heuristic. This demand differs in form from the pressure generated by the exclusion argument, but in effect the two pressures reinforce one another.

Accordingly, the present book's attitude toward the tradition of mechanistic explanation is not one of simple rejection or replacement. Mechanistic analysis unquestionably occupies a central position in scientific explanation, and this book does not deny its importance. What truly needs to be questioned is the tendency to elevate mechanistic explanation into a universal paradigm. If explanatory activity is genuinely multi-role, then mechanistic explanation should be understood as one crucial role among others, namely, a way of answering productive questions, rather than as the sole model of explanation. Once this is acknowledged, we can reexamine those forms of account that have been placed "outside mechanism" and ask whether they are in fact answering another class of equally basic questions.

It is against this background that the next section turns to the causal-explanatory tradition associated with Woodward. Unlike mechanistic explanation, Woodward does not emphasize the composition of concrete structures, but rather counterfactual dependence and relations of intervention. The question to be

examined next is whether this refinement genuinely breaks with explanatory monism, or instead extends it in a more abstract form.

2.5 The Refinement of Causal Explanation and Its Boundaries: Starting from Woodward

If the tradition of mechanistic explanation is persuasive because it stays closely aligned with scientific research into generative processes, then the greater philosophical weight of the causal theory of explanation associated with Woodward lies in the fact that it offers a far more refined analysis than the traditional production model of what counts as causal relevance and what makes an explanation effective. Rather than understanding causal explanation simply as a temporal sequence in which one event produces another, Woodward seeks to clarify explanation through intervention, counterfactual dependence, and manipulability. In doing so, he gives causal explanation a more disciplined and conceptually articulate form. This is why his work has been taken, in contemporary philosophy of science, as one of the most powerful resources for thinking about explanation after the decline of the covering-law model.

The attraction of this approach is easy to see. It does not require that explanation always consist in narrating an underlying mechanism in concrete detail. Instead, it asks whether changing one variable under appropriate interventions would systematically change another. An explanation succeeds, on this view, when it identifies patterns of dependence that support relevant counterfactuals. This allows the theory to accommodate a wide range of scientific practice, including cases in which researchers cannot yet fully specify the microstructure of the process involved. It also captures an important intuition: many explanations are indeed informative because they tell us what would happen if things were different.

For these reasons, Woodward's framework is in many respects more flexible and more philosophically sophisticated than both older production-based pictures and narrower mechanistic models. It recognizes that explanation is not merely about listing prior events, but about identifying difference-making relations that support understanding, intervention, and inquiry. Many scientific explanations genuinely owe their force to this kind of structure. A causal theory of explanation in this form is therefore not something the present book seeks to dismiss or replace.

But precisely because this tradition has become more refined, its limits become more important to examine. The question is no longer whether causal explanation matters. It plainly does. The question is whether every legitimate explanation must ultimately be rewritten in terms of interventionally tractable dependence. Once that stronger claim is made, the framework begins to risk functioning as another route back to explanatory monism. For then explanation, however sophisticated its formulation, is still treated as if its core task were to reveal a pattern of causal dependence whose content can be cashed out through intervention and counterfactual variation.

This becomes especially important when we turn to the kinds of cases that concern this book. When scientists appeal to symmetry, conservation, topological restriction, phase-space structure, boundary conditions, or large-scale organizational form, are they always describing causal dependence in a more abstract idiom? Or are they sometimes answering a different sort of question altogether? If an account explains why certain outcomes are impossible, why only some trajectories are available, or why certain stable patterns persist across many realizations, then what is doing the explanatory work may not be interventionally manipulable dependence in Woodward's sense. It may instead be the articulation of structural limitation, organization of possibility, or the selective openness of a system under given constraints.

The importance of this point lies in the fact that Woodward's framework, precisely because it is more abstract and more accommodating, can absorb other forms of account even more effectively than mechanism can. Mechanistic explanation at least announces itself as a model focused on the organized production of phenomena. Interventionist causal explanation, by contrast, can appear broad enough to subsume almost any genuine explanatory relation, so long as one can redescribe that relation in terms of difference-making under some hypothetical intervention. The danger, then, is not crudeness but overextension. A framework of great explanatory power can begin to function as though it had no real boundary.

Once this happens, the same pressure reappears in a more sophisticated form. A structural account may be granted importance, but only insofar as it can ultimately be reconstrued as indicating some pattern of difference-making dependence. A symmetry principle may be treated as explanatory only because it tracks what would vary under certain changes. A constraint may be admitted only insofar as it can be redescribed as a higher-order determinant of interventional space. None of these moves is always illegitimate. But if they become mandatory, then the independent standing of structural constraint and permissibility again becomes hard to see. They are not denied, but they are translated back into the causal idiom before they are allowed full explanatory status.

The issue, then, is not whether Woodward's theory is too broad to be useful. On the contrary, its breadth is one reason for its usefulness. The issue is whether this breadth conceals a boundary problem. Not every explanatory question is a question about what would change under intervention. Some are questions about why the space of possible change is already organized in a certain way. Some are questions about why certain forms are ruled out before any local process even begins. Some are questions about why a pattern remains stable across many causal implementations. These are not obviously reducible to the same explanatory role

as causal dependence, even when causal dependence remains relevant within the broader picture.

This is precisely where the distinction proposed in this book becomes necessary. Causal production, structural constraint, and permissibility are not three rival accounts of the same explanatory job. They are three different explanatory roles. Woodward's framework gives us powerful resources for understanding one of those roles, and in some contexts it may partially illuminate the others. But it does not follow that the role of structural constraint or the role of permissibility can simply be collapsed into interventionist dependence. If we force that collapse, then the same contraction that appeared in older production models and in mechanistic expansion reappears in a more refined register.

Seen from this angle, the contribution of Woodward's tradition is double. Positively, it shows that causal explanation need not be understood naively, and that explanation can be analyzed through counterfactual structure rather than merely through temporal sequence. Negatively, it also reveals how explanatory monism can survive even after explanation has become more conceptually sophisticated. Monism need not insist that explanation is a crude chain of production. It can persist by allowing only one ultimate format of explanatory legitimacy, however abstractly stated.

For that reason, the present book does not oppose Woodward by denying the reality or force of interventionist explanation. It instead asks a boundary question: where does this kind of explanation end, and what lies outside it without thereby falling outside genuine explanation? Once that boundary question is raised, it becomes possible to see why structural and permissibility-based accounts should not be treated merely as derivative causal reconstructions. They may instead be answering different questions, with different criteria of success, within a genuinely differentiated explanatory landscape.

This, in turn, helps prepare the transition to the next part of the chapter. Once we have seen how Kim, mechanistic explanation, and Woodward each sustain, in different ways, the pull toward a single productive tribunal, we are in a better position to clarify why existing advances in non-causal explanation remain insufficient unless they are embedded in a fuller framework of explanatory roles. What remains to be shown is not simply that non-causal explanation exists, but why its existence points toward a reorganization of explanation itself.

2.6 Why the Existing Literature Still Lacks a Unified Framework of Explanatory Roles

After the preceding sections, it should now be clear that contemporary philosophy of science does not lack important resources relevant to the concerns of this book. The exclusion tradition forces us to confront the pressure faced by higher-level accounts under the background of causal closure; the tradition of mechanistic explanation reminds us that uncovering generative pathways is indeed a highly central scientific practice; Woodward's theory of causal explanation elevates that practice into a more refined counterfactual and interventionist analysis; and meanwhile, work on non-causal explanation, symmetry, constraint, and structural accounts has already succeeded in undermining the overly strong claim that all explanations must be locally productive accounts. Taken individually, these contributions are all of considerable importance.

And yet the judgment of this book is that, even so, the existing literature still has not provided a unified framework of explanatory roles capable of spanning these domains of problem. By this "lack," I do not mean that the existing work has achieved little. I mean, rather, that it has for the most part developed in scattered, local, or dispute-specific ways, without yet yielding an overall picture capable of answering, in a stable manner, questions such as the following: what distinct tasks different kinds of explanation are performing; how these tasks are to be

distinguished; why they may all be legitimate; and why they need not be adjudicated within one and the same competitive structure. In other words, the literature has already shown us that the space of explanation is broader than a single causal model, but it has not yet sufficiently clarified how that broader space itself ought to be divided.

The first reason for this situation is that many studies, while successfully demonstrating the explanatory force of some non-productive form of account, still proceed primarily by way of case-by-case defense. Some works argue that symmetry principles can explain certain phenomena; others show that mathematical structure or limiting behavior plays an irreducible explanatory role in certain contexts; still others demonstrate that boundary conditions or topological properties are not mere background descriptions, but key explanatory elements. All of this is important. But the typical form of advance remains the same: to show that some specific kind of account ought not to be excluded from explanation. That can effectively loosen explanatory monism, but it does not by itself force the more general question of whether explanatory activity, from the outset, contains several different roles. As a result, such studies often succeed in expanding the list of “admissible explanatory types” without fully rewriting the more basic question of what explanation is.

The second reason is that many existing disputes, though they do touch on level, structure, and constraint, still preserve an implicit order of adjudication. That is to say, they may acknowledge that certain non-productive accounts are useful in inquiry, and even admit that they are irreplaceable, while still tending, when it comes to deciding whether an explanation is “ultimately successful,” to retain causal production as the higher-order standard. The result is that structural accounts often occupy an unstable position: they may be described as illuminating, organizing, or theoretically valuable, while still being tacitly denied an equal final standing with causal explanation. So long as that order of adjudication remains in

place, evaluative monism has not truly been abandoned, and explanatory diversity is only tolerated in a limited way.

The third reason is the lack of sufficiently strong lateral connections among the different traditions of literature. Discussions of higher-level explanation often remain confined to the contexts of philosophy of mind, realization, and nonreductive physicalism. Discussions of mechanistic explanation are more deeply rooted in the practices of the life sciences and neuroscience. Theories of counterfactuals and intervention are largely framed around causal modeling and explanatory relevance. And work on non-causal explanation, symmetry, and constraint is often concentrated in cases drawn from physics or mathematical explanation. There is nothing wrong with such division of labor in itself. But one consequence is that the same deep structure may recur across multiple fields without being recognized as the same problem. Emergence, higher-level autonomy, entanglement, methodological legitimacy, and ontological inflation therefore continue to be treated separately, rather than being brought back under a single framework of explanatory roles.

The fourth reason, and the most crucial one, is that the existing literature generally lacks sufficiently explicit criteria for distinguishing what productive accounts and constraint-based accounts are respectively answering. Many disputes remain deadlocked not merely because participants disagree, but because they are not always answering the same kind of question. One person may be asking how a result occurs; another may really be asking why only this class of result is possible, or why some pattern remains stable across differences of implementation. If these two kinds of inquiry are not clearly distinguished, then any structural account can easily be retranslated within debate into an incomplete productive story, and any productive account can equally be misrequired to bear the burden of explaining overall organization and the boundaries of permissibility.

The result of lacking criteria is not merely theoretical vagueness. It is that the very types of problem keep covering over one another.

From the standpoint of this book, this gap matters especially because it directly affects whether several later core issues can be reorganized at all. Without a more stable framework of explanatory roles, the discussion of emergence will repeatedly return to whether there are new causal powers; the discussion of higher-level explanation will again and again be compressed into whether it competes with the physical level; the discussion of entanglement will repeatedly be redirected toward whether there is action at a distance; and the discussion of unknown phenomena will continue to slide toward whether new ontology must be introduced. These questions continue to be asked in these ways not because no alternative is available, but because the existing literature has not yet sufficiently supplied a broadly usable alternative way of asking them.

For this reason, the “unified framework” proposed in this book does not mean reducing all explanation to some single structure. Quite the contrary: it means offering a unified principle of division capable of accommodating difference. The unity in question is not homogeneity of content, but clarity of method. This book does not aim to invent some new master-model that replaces all existing theories of explanation. Rather, it aims to provide a more general way of distinguishing when we are facing a problem of causal production, when we are facing a problem of structural constraint, when an account is primarily characterizing a boundary of permissibility, and when accounts at different levels are working together rather than competing with one another. Only when clarity is achieved at this level can the various resources surveyed in the preceding sections be genuinely organized, rather than merely placed side by side.

This also shows why Chapter 3 plays a pivotal role in the book as a whole. The task of Chapter 2 is not to provide the solution directly, but to establish two points. First, explanatory monism is indeed deeply embedded in contemporary disputes

in multiple forms. Second, although the existing literature has already supplied a large number of important resources, it has not yet integrated them into a theory of explanatory roles capable of spanning multiple issues. The conceptual reconstruction that follows is therefore not an arbitrary addition of new terminology, but a direct response to a gap that the preceding literature has jointly exposed.

2.7 A Clarification of the Target of the Book's Critique

What this chapter calls “explanatory monism” does not mean that contemporary philosophy of science contains a unified school openly proclaiming that only causal explanation is explanation. On the contrary, this book fully acknowledges that, over recent decades, discussions of non-causal explanation, mathematical explanation, symmetry explanation, topological explanation, and constraint explanation have significantly broadened the scope of explanatory theory. The target of critique here is more specific: when certain disputes enter the stage of normative adjudication, productive explanation is often tacitly treated as the highest standard. Structural accounts may thus be acknowledged as useful, yet they are still required to show how they ultimately serve, supplement, or reduce to accounts of causal production. It is this adjudicative structure, rather than any single school, that constitutes the core of what this book calls explanatory monism.

2.8 Chapter Summary: From a Map of the Literature to Conceptual Reconstruction

The point of departure of this chapter has been a diagnosis of a phenomenon that appears dispersed but in fact runs through multiple topics: in contemporary philosophy of science and related theoretical disputes, explanation is often presupposed to be a single kind of activity, whose core task is understood as revealing how some result is produced. It is this presupposition that forces

accounts belonging to different levels into one and the same competitive space, and that causes persistent theoretical tensions to be mistaken for conflicts in the structure of the world itself rather than mismatches among explanatory tasks.

In order to grasp this more precisely, the chapter first distinguished three interconnected levels of explanatory monism: evaluative monism, role monism, and competitive monism. The first determines what counts as a genuinely sufficient explanation; the second determines what kind of task explanation is taken to perform; and competitive monism pulls accounts from different levels into one and the same exclusionary structure of adjudication. Through this distinction, the chapter has tried to show that the problem is not merely that some particular theory happens to neglect structure, constraint, or level, but that explanatory space itself has been contracted at a deeper level. This is also why problems such as emergence, higher-level explanation, quantum entanglement, non-predictive frameworks, and ontological stability repeatedly reproduce similar deadlocks in apparently different domains.

On that basis, the chapter then examined several key lines of literature. First, Kim's exclusion argument displays in its sharpest form the pressure borne by higher-level accounts under the framework of causal closure. Its importance lies not only in raising a causal problem, but also in systematizing the slide from causal sufficiency to explanatory sufficiency. Second, the tradition of mechanistic explanation reinforces the same contracting effect along a different path. Precisely because it closely tracks a great deal of scientific practice and effectively demonstrates the force of generative accounts, mechanistic explanation is easily taken as the standard form of explanation, so that other forms of account, even when not openly excluded, can often achieve standing only in subordinate or transitional form. Third, the causal theory of explanation associated with Woodward carries this productive orientation into a more refined level of counterfactual and interventionist analysis. For that very reason, its boundary problem becomes

especially important: even within so sophisticated a framework, there remain problems whose core is not how variables alter outcomes, but how a system's space of possibilities is organized, constrained, and stabilized.

Set against these causally centered traditions, the chapter also reviewed work on non-causal explanation, constraint, symmetry, and structural accounts. The relevant literature has already shown clearly that scientific understanding does not always depend on tracing local productive chains; in many settings, explanation succeeds because it reveals why certain possibilities are excluded, why certain relations remain invariant, and why certain patterns possess stability. At the same time, however, the chapter argued that these studies, though extremely important, do not yet automatically amount to a unified framework capable of spanning multiple domains of dispute. They have mostly succeeded in securing legitimacy for particular forms of account, but they have not yet sufficiently answered what different kinds of explanation are respectively doing, how those tasks are to be distinguished, and why they need not fall into one competitive structure.

The conclusion reached here is therefore not merely that non-causal explanation exists. It is the stronger claim that the existing literature has already exposed the problem of explanatory monism and has supplied important resources for overcoming it, yet still lacks a more stable theory of explanatory roles capable of showing what kinds of problem correspond respectively to causal production, structural constraint, and permissibility, and why all three may constitute legitimate explanation. Up to this point, the work of the chapter has remained primarily diagnostic. It shows that many tensions in contemporary debates are not only the product of substantive disagreement, but also of a persistent slippage caused by unclear boundaries among explanatory tasks.

That also means that the next step cannot remain at the level of literary comparison or reconciliation among positions. If this book is to respond seriously to the foregoing problems, it must move into more rigorous conceptual

reconstruction. It is not enough to say that there is more than one kind of explanation. We must also say what the different explanatory roles are, how they are distinguished, what questions they answer, and under what conditions accounts at different levels are working together rather than excluding one another. Only when clarity is achieved at this conceptual level can the later discussions of emergence, higher-level explanation, quantum entanglement, and ontological stability escape their repeated orbit around the single axis of whether there are additional causal powers.

Accordingly, Chapter 3 turns to the theoretical core of the book and introduces three key distinctions: causal production, structural constraint, and permissibility. Chapter 2 has shown that without such distinctions, however rich the existing literature may be, it remains insufficient to prevent different kinds of problem from being repeatedly miswritten as one and the same productive task. The goal of Chapter 3 is to supply a clearer conceptual framework exactly at this point of lack. Only once that step has been completed can the book's preceding diagnosis of the literature genuinely be transformed into a theoretical proposal with systematic explanatory power.

Chapter 3

The Differentiation of Explanatory Roles: Causal Production, Structural Constraint, and Permissibility

If the analysis of the previous chapter is correct, then the reason so many contemporary theoretical disputes repeatedly fall into deadlock is not simply that our grasp of certain phenomena remains incomplete, nor simply that different positions render different judgments on the same facts. The deeper problem is that different kinds of explanatory task have long been compressed into a single problem-framework. Any account of theoretical importance is required to answer one and the same question: how was some result produced? For that very reason, analyses that are in fact closer to the explanation of stability, limiting conditions, forms of organization, and the boundaries of a space of possibility are continually retranslated into narratives of event-generation; and once they fail to complete that translation, their explanatory standing begins to look suspect.

If this diagnosis is sound, then the response cannot consist merely in securing exceptional status for a few special forms of account within the existing linguistic framework, nor can it rest content with repeating the general claim that explanation is not limited to causal explanation. What the book must now undertake is stricter conceptual work: to show how different explanatory tasks are to be distinguished, what questions they respectively answer, and why they may all count as legitimate explanation. What we need, in other words, is not a loose acknowledgment of explanatory diversity, but a framework capable of stably distinguishing different explanatory roles. Only if clarity is achieved at this point can the problems treated in the earlier chapters, such as emergence, higher-level explanation, quantum entanglement, and ontological stability, be reorganized rather than continuing to circle around the single axis of whether there are additional causal powers.

The central claim of this chapter is that explanatory activity contains at least three elements that are related but must not be conflated: causal production, structural constraint, and permissibility. Causal production concerns how an event, state, or result is brought about in a concrete process; structural constraint concerns which organizational relations, boundary conditions, or global structures in a system limit the range of changes that can occur; and permissibility concerns which configurations, trajectories, patterns, or results can stably appear under given constraints, and which are excluded as impossible, unstable, or inconsistent. The three are connected, but they should not be compressed into one and the same explanatory activity. Most importantly, structural constraint and permissibility are neither failed substitutes for causal production nor forms of account that appear only temporarily when causal knowledge is insufficient. They answer another class of equally basic questions, and their explanatory legitimacy should therefore not depend on whether they can be completely rewritten as productive narratives.

The task of this chapter is to develop that claim into a more operational conceptual framework. To do so, it will answer four questions in sequence. First, why must explanatory roles be differentiated, rather than being understood as mere differences of granularity within a single activity? Second, how exactly should causal production be defined, and what task does it perform in explanation? Third, what is the role of structural constraint, and how does it differ from mere background conditions or peripheral description? Fourth, in what sense does permissibility constitute an explanatory target rather than simply receiving explanation passively as a result? In the end, the chapter will propose a number of criteria of distinction to help determine when we are dealing with a productive problem and when what truly requires explanation is organization, limitation, and the boundaries of a space of possibility.

3.1 Why Explanatory Roles Need to Be Differentiated

The claim that explanatory roles need to be differentiated may initially provoke an intuitive doubt. Why not say that all explanation is, at bottom, just one and the same activity unfolding at different levels? Some explanations may be more detailed, others more general; some may lie closer to local mechanisms, others closer to the overall picture; but are they not all ultimately trying to explain why the world is as it is? If so, then emphasizing the differentiation of roles might seem to make a simple matter unnecessarily complicated. Rather than introducing new distinctions, why not place all accounts on a single spectrum, where better explanations are simply those that are more complete, more precise, and closer to the real productive structure?

This view is attractive because it captures something true: many scientific explanations are indeed related to one another, and they often support one another across different levels of abstraction. A higher-level account may guide us toward the search for a lower-level mechanism, and an account of a lower-level mechanism may in turn help us understand why a higher-level phenomenon appears. In a great deal of research practice, explanations are not isolated modules, but nested, mutually supportive, and mutually illuminating networks of work. For that reason, when this book argues that explanatory roles need to be differentiated, it is not trying to split explanation into several disconnected kinds, nor to deny continuity and cooperation among different forms of account.

The problem is that continuity is not identity, and cooperation does not mean that all explanations answer the same question. Precisely because different accounts can work together, we need all the more to understand what task each of them performs. Without such a distinction, the cooperative relation among explanations is easily miswritten as a hierarchical relation, and we slide toward a picture in which every higher-level, more global, or more structural account is merely a rough preview of the real explanation, while the real explanation always

lies at the level best able to show how a result was produced. In such a picture, explanation is indeed unified, but at the cost of losing our ability to judge why different forms of account are legitimate.

Explanatory roles need to be differentiated, first, because the questions we actually ask in theory and scientific practice are not of only one kind. Sometimes we ask: how did some result occur? Which events, conditions, or activities together brought it into reality? This is plainly a productive kind of question. It asks about generative pathways, difference-making factors, causal chains, or mechanistic organization. In such cases, explanatory success generally depends on whether the explanation reveals how the result was brought about, and how it would vary if certain conditions were different.

At other times, however, what we are asking is not primarily how some single result was produced, but why certain patterns can stably exist, why certain configurations are systematically excluded, why some organizational form can be maintained despite changes in details of implementation, or why a system allows one class of trajectories rather than another. In such questions, the focus is no longer simply where an event comes from, but within what boundaries the system itself operates. What matters here is often not a particular generative path, but the structure of conditions that make such paths possible, sustainable, or limited in the first place. If we treat such questions simply as “not-yet-completed causal questions,” we miss their true explanatory object.

Explanatory roles need to be differentiated, second, because without such differentiation the miswriting of one problem as another will continue to generate theoretical deadlock. The earlier chapters have returned to this point repeatedly: when questions about emergence are rewritten as questions about whether there are new causal powers, structural novelty is forced into a choice between mysterious causation and explanatory nihilism; when questions about higher-level explanation are rewritten as questions about whether they compete with physical

causes, accounts at different levels are compressed into the same exclusionary structure; when questions about entanglement are rewritten as questions about what influence propagates long-distance correlations, the structural problem of the permissibility of joint outcomes disappears; when questions about unknown phenomena are rewritten as questions about whether new entities must be introduced, reflection on explanatory discipline and ontological robustness is skipped over. These difficulties appear across different domains, but they all arise from the same mistake: fixing the problem too early as a productive task.

Explanatory roles need to be differentiated, third, because the standards of legitimacy are not entirely the same. For a productive explanation, we typically ask: has it identified the relevant difference-making factors? Has it revealed the generative pathway of the result? Has it shown how the result would change under intervention? But for a constraint-based or permissibility-based explanation, the focus of evaluation is often different. What matters more is whether it accurately reveals system boundaries, organizational structure, conditions of stability, or invariant relations; whether it explains which possibilities are excluded and why; and whether it helps us understand why certain patterns continue to persist across multiple realizational conditions. If we insist on evaluating all explanations by one single standard, the result is usually not unity but systematic misjudgment. Accounts not meant to trace productive chains will be judged incomplete because they do not provide such chains, while problems that ought to concern structural boundaries will be distorted because they are forced to tell stories of event-generation.

Finally, explanatory roles need to be differentiated not in order to exaggerate difference, but in order to understand cooperation more accurately. Causal production, structural constraint, and permissibility are not isolated from one another. A mechanism is often able to produce some result only within certain structural boundaries; a pattern often counts as permissible only because it can be

realized through concrete processes. The issue is not whether the three are related, but whether the relation among them should be understood as one of ultimate reduction or, rather, as one of answering different kinds of question. Only if the roles are distinguished in this sense can we understand why a complete theoretical picture often requires multiple kinds of account to coexist without being treated as though they were competing for the same explanatory position.

Accordingly, the distinctions introduced in this chapter should not be misunderstood as an artificial cutting-up of explanatory space. They should be understood as conceptual clarification of real differences in explanatory practice. If the work of the first two chapters was to show why explanatory monism generates systematic mismatch, then from this point onward the task of the book is to offer an alternative framework in which that mismatch is no longer invisible. The next section therefore begins by defining the meaning and scope of causal production itself. This is not because the book seeks to weaken causal explanation, but precisely because only by getting clear about causal production first can we secure for structural constraint and permissibility a position that is both clear and noncompetitive, without diminishing the importance of causal explanation.

3.2 Causal Production: Generative Pathways and Difference-Making

If this book is to argue that explanatory activity has more than one role, it must first avoid a common misunderstanding: namely, the thought that making space for structural constraint and permissibility must somehow weaken the standing of causal explanation. The truth is just the opposite. Precisely because causal production has irreplaceable importance in scientific explanation, we need first to state accurately what it is, in what kinds of problem it operates, and when its explanatory task has been completed. Only if clarity is achieved at this point will the later discussion of structural constraint and permissibility not appear to be an

attempt to escape rigorous explanatory demands by means of vague language, but rather a legitimate distinction among different explanatory tasks.

Within the framework of this book, causal production refers first of all to the explanatory task of showing how some event, state, or result is brought into reality through a concrete process. It is concerned with generative pathways rather than mere static correlation. It asks which factors, conditions, activities, or mechanisms work together so that a given result occurs. It asks, "How did this result come about?" rather than simply, "What structures make this class of result possible?" In the broadest sense, causal production deals with questions of occurrence in reality, not questions about the boundaries of a space of possibility.

It is not difficult to see why this role is so central in science. Many explanations are satisfying precisely because they show that a result does not arise out of nowhere, but is generated along some path that can be traced. Whether in describing how forces act on the state of a system in a physical process, explaining in biology how some function is jointly realized by a particular structure and its activities, characterizing in neuroscience how a cognitive performance depends on a specific neural mechanism, or explaining in everyday life why some phenomenon occurred under a sequence of conditions, what people seek is some understanding of a pathway of occurrence. It is this path-like understanding that makes explanation more than simply subsuming a result under a category; it reveals why the result appears in this concrete world.

Yet if causal production is to be defined more rigorously, it is not enough merely to say that it concerns "pathways of occurrence." Not every description of a pathway constitutes a genuine causal explanation. A casually arranged sequence of events does not automatically become explanatory, and a condition that merely accompanies a result does not thereby become a relevant cause. To count as a productive account, an explanation must satisfy at least two interconnected requirements.

First, it must point to some generative connection. That is, the factors included in the explanation are not merely juxtaposed with the result by accident, but play the role of making the result occur. This role may be characterized in terms of mechanism, process, organized activity, counterfactual dependence, or difference-making relation; but whatever the vocabulary, the core thought is the same: without the appropriate intervention of these factors, the result would not occur in the same way. Causal production does not simply gather every correlated factor into the explanation; it attempts to identify those conditions that are genuinely generatively significant for the appearance of the result.

Second, it typically also involves some structure of difference-making. A good causal explanation does not merely state how the result was brought about; it also shows how the result would differ if the relevant factors were different. It is precisely here that modern theories of causal explanation, especially counterfactual and interventionist approaches, have offered important clarification. Causal production is not merely a narrative of an accomplished fact; it also reveals a structure of dependence: what changes what, which factors are the key difference-makers for why the result is as it is, and which are merely accompanying conditions or background circumstances. For this reason, causal production not only explains results that have already occurred, but also provides a basis for control, prediction, prevention, and intervention.

A particularly important point here is that generative pathways and difference-making are not two separate definitions, but two complementary aspects of the same explanatory role. The former emphasizes how a result is realized in the world; the latter emphasizes why that realization has explanatory relevance. If there is only a pathway without difference-making, the explanation risks collapsing into a long narrative. If there is only difference-making without any generative understanding, the explanation may become so abstract that it no longer shows through what sort

of process the result actually comes to appear. A mature causal-productive explanation typically achieves a balance between these two aspects.

From this, we can see the boundary of causal production more clearly. It is best suited to answer questions centered on event-generation, state-transition, the occurrence of outcomes, and variable-change. Why did some reaction occur? Why did some system enter a given state? Why did some behavior appear under these conditions? Why did some result vary with a given factor? What all such questions have in common is that they require us to explain how a real result was brought about, and to identify the generative and difference-making structure relevant to it. For problems of this kind, causal production is not merely one optional explanatory style among many; it is often the central explanatory task itself.

But precisely because causal production has such a clear and powerful domain of task, we must not let it expand without limit. It is excellent at answering how a result occurs, but it is not automatically equally well suited to answering which results are themselves possible, why certain configurations are systematically excluded, or why some pattern can remain stable across differences of realization. Such questions may indeed be related to causal processes, but they are not always asking first about a generative chain. If we rewrite all of them as "Which factor caused this result?" we risk wrongly compressing questions of boundary, organization, and stability into the productive model. For that reason, marking out the scope of causal production is not a way of weakening it, but of preventing its success from being mistaken for the whole of explanation.

One further clarification is also necessary: causal production is not equivalent to a reductive narrative at the most micro-level. So long as an explanation takes up the task of generative pathway and difference-making, it may be causal at different levels. In other words, this book does not reserve "causal production" for the lowest level of physics, nor does it hold that only the finest-grained descriptions qualify as causal explanation. Mechanistic explanation can be causally productive, and some

higher-level process models can also be causally productive, so long as they genuinely explain how a result is brought about and which factors constitute the relevant difference-makers. This point matters because the book's later argument is not that higher-level explanation as such should be moved outside the causal domain, but rather that we must distinguish when a higher-level account is doing productive work and when it is doing structural work. Without first separating causal production from "micro-reduction," the later distinction would be misunderstood from the outset.

On this account, causal production has at least the following marks. It answers questions of occurrence rather than questions of boundary. It traces generative pathways rather than merely describing stable relations. It identifies relevant difference-making factors rather than treating every accompanying condition as a cause. And its success typically lies in revealing why the result appeared under these conditions and how it would vary if those conditions changed. Any account whose explanatory work is performed mainly in this way belongs within the scope of causal production.

Once this has been clarified, we can see more accurately that the book's later argument is not meant to deny the explanatory standing of causal production, but to block a stronger slide: the slide from "causal production is an important explanatory role" to "only causal production is genuine explanation." The former is entirely correct; the latter wrongly demotes many accounts that do not take generative pathways as their primary target. It is on this basis that the next section can turn to structural constraint and show why some explanations are not primarily answering how a result is brought about, but rather another equally basic question: what kinds of change, configuration, and pattern are themselves allowed, stable, or sustainable within a system.

3.3 Structural Constraint: Boundaries, Organization, and Conditions of Exclusion

If the work of the previous section was to define the kinds of question answered by causal production, then the task of this section is to show that not all legitimate explanations have the disclosure of generative pathways as their core. In many theoretical and scientific practices, what we seek is not how some result is brought about step by step, but why a system allows only certain changes, certain configurations, certain forms of organization, or certain outcomes, while excluding others. What stands at the center of explanation here is not the productive “process of occurrence,” but the limiting “structure of boundaries.” The book groups such accounts under the heading of structural constraint.

Structural constraint, first of all, does not mean just any vague “background condition” or “environmental factor.” If the concept is to have theoretical force, it must be stricter than “all conditions that influence a result.” What structural constraint names is an explanatory function of the following sort: it works not primarily by showing how some result is generated, but by revealing which relations, boundaries, organizational principles, or global conditions in a system limit the range of changes that can occur, thereby explaining why some states, trajectories, or patterns are feasible while others are impossible, unstable, or inconsistent. Structural constraint, in other words, does not directly bring a result into reality; it specifies the boundary conditions under which that result can appear.

This can be grasped at the most basic level. Causal production deals with the question of how something occurs; structural constraint deals more nearly with the questions under what limitations it can occur and what cannot occur at all. The two are of course related, because no real generative process can exist independently of some boundary conditions and organizational structure. But relation does not mean that they perform the same explanatory task. An account may be fundamentally important to the appearance of a result without thereby being the

producer of that result. Structural constraint is easily misunderstood precisely because it matters explanatorily without mattering by way of “pushing” or “making” a result.

To define structural constraint more clearly, at least three features must be emphasized.

First, structural constraint is limiting rather than directly productive. Its role is not to push some result forward, but to delineate which paths of change, forms of organization, or combinations of states are allowed within a system, and which are excluded. A boundary condition may not “cause” a system to enter some particular state, yet it may render a great many other states impossible from the outset. A certain global organizational relation may not appear as the cause of any single event, yet it may determine that the system can operate only within a limited set of patterns. A symmetry condition or topological relation may not intervene as a local driving force, yet it may stably restrict the structural forms available to the system. The explanatory force of structural constraint lies in this function of exclusion and delimitation.

Second, structural constraint is organizational, not merely externally restrictive. By “structure” here I do not mean a set of accidental peripheral conditions sitting outside the process. I mean the relational arrangement internal to the system or operative at the level of the system itself. It may take the form of the organization among components, the pattern of dependence between whole and parts, or boundary conditions, symmetries, invariances, topological relations, or conditions of realization. The crucial point is that such factors have explanatory force not because they happen to accompany some result, but because they define the way the system can operate. A genuine account of structural constraint does not add one more background factor outside the result; it explains how the very space of possibilities to which the result belongs is constructed.

Third, structural constraint has a dimension of stability. It is usually not a chance limitation on some isolated event, but an enduring delimitation on a class of possible changes. For that reason, structural constraint is often tied to questions of pattern maintenance, realizational stability, and consistency across variation. We invoke structural constraint not because we want to explain why one particular event happened not to occur, but because we want to explain why, across a wide range of conditions, one class of states is persistently excluded or one class of patterns is persistently maintained. Without this dimension of stability, what we call "constraint" easily degenerates into an arbitrary list of background conditions and loses its independent explanatory significance.

These features also help distinguish structural constraint from ordinary background conditions. Not every background condition is a structural constraint. A factor may be a contingent prerequisite for some result without thereby organizing the system's space of possibility. An environmental variable may be necessary in one context without thereby constituting a structurally bounded condition with independent explanatory standing. Structural constraint is different because it is not merely there; it continuously and organizationally limits what can happen in a system, what cannot happen, and what can be stably sustained. In other words, it is not decorative material added at the edge of an explanation, but part of the internal structure that defines the system's operability.

In this sense, the relation between structural constraint and causal production can be stated provisionally as follows: causal production explains how a result is brought about along some path, whereas structural constraint explains which paths, forms of organization, and combinations of states are open, closed, stable, or excluded for the system. The former focuses on the process of realization; the latter focuses on the boundaries of realization. The former answers what makes the result occur; the latter answers what determines that only this class of result can occur, or that certain results cannot occur at all. Without this distinction, structural

constraint will continually be misread as an incomplete causal narrative, and its real explanatory function will remain obscured.

The importance of this distinction can be tied directly to several of the book's later themes. In the problem of emergence, many key accounts are not mainly telling a story about how higher-level patterns are additionally "produced," but about what kind of organizational relation allows a pattern to remain stable across different micro-level realizations. What is being explained there is a constraining structure over a space of realization, not merely a generative chain for a single event. The same holds for higher-level explanation: the importance of some higher-level accounts does not lie in adding new lower-level causes, but in articulating stable relations at a level of organization, thereby helping us understand why certain changes matter at that level while others are systematically ignored. In the case of quantum entanglement, the point becomes sharper still: if entanglement primarily reflects the inseparability of joint configurations, then the explanatory focus lies not on what sort of superluminal influence propagated the outcomes, but on what structure allows some combinations of joint outcome from the start while excluding others.

Of course, acknowledging the independence of structural constraint does not mean that it stands wholly apart from causal process. Structural constraint is not a set of abstract rules floating outside real processes. Any real system-constraint must operate in some realizational structure and ultimately connect with concrete processes. Nor does this book claim that structural constraint and causal production constitute two sealed worlds. The question has never been whether they are related, but whether they should be treated as the same explanatory role. Structural constraint may shape the boundaries within which causal processes unfold, and causal processes may realize concrete outcomes within given constraints; but it does not follow from this that structural constraint is itself a kind

of causal production, any more than setting the stage and performing an action on the stage are the same act simply because they are interdependent.

From this perspective, the explanatory standing of structural constraint can be defended neither by mystifying it nor by preserving it only through complete reduction to causal production. What this book proposes is a more restrained and more rigorous understanding: structural constraint matters not because it introduces a new hidden force, but because it answers another class of questions about boundaries, organization, and limiting conditions. It does not replace causal production; it joins with it to form a more complete explanatory picture. Without causal production, we cannot understand how outcomes are realized in the world; without structural constraint, we equally cannot understand why only certain realizational paths, patterns, and stable forms are open to the system.

At this point the next question arises naturally. Even if we grant that structural constraint characterizes boundaries, organization, and limiting conditions, we still need to ask what exactly such constraint explains. Does it merely provide background for later causal analysis, or does it itself point toward an independent explanatory target? That question leads directly to the core concept of the next section: permissibility. For only once we see that the question of which configurations, trajectories, and patterns are allowed can itself be an explanatory object does the role of structural constraint fully come into view.

3.4 Permissibility: Stable Patterns and Realizable Space

If causal production explains how a result is brought about, and structural constraint explains which boundaries, organizational relations, and limiting conditions define the range within which a system can operate, then the next question is this: toward what explanatory target do such constraints point? Without a further answer, structural constraint may still be misunderstood as a merely ancillary kind of background analysis, as though its true significance lay only in

preparing the way for productive explanation. This book argues that, to avoid that misunderstanding, permissibility must be explicitly introduced as an independent explanatory object. In many contexts, what requires explanation first is not how some result occurs, but why certain results, configurations, trajectories, or patterns are feasible, stable, and consistent for a system, while others are excluded.

Permissibility, as the term is used in this book, does not refer to subjective acceptability, nor does it mean merely that something “can happen” in a loose sense. More strictly, it refers to a structural fact of the following kind: under given boundary conditions, organizational relations, and system-constraints, which combinations of states, pathways of change, or forms of pattern can be stably borne by the system, and which cannot. Permissibility is therefore not concerned with whether some result actually occurred, but with whether that result belongs within the space of possibility opened by the system. It concerns the organization of that space, not the advancement of a sequence of actual events.

This concept is necessary because without it the explanatory function of structural constraint still cannot be fully stated. If we say only that structural constraint “limits certain changes” without further specifying what such limitation is explaining, then constraint still looks like a peripheral condition, as though it merely tells us which causes are more likely to operate and which paths are more likely to succeed. Once permissibility is introduced, however, the center of gravity shifts. Structural constraint no longer functions merely as scenery for causal process; it directly explains why certain configurations fall outside the system’s realizable range from the outset, why certain patterns can persist despite local differences of implementation, and why certain imaginable changes are in fact not open to that system.

From the standpoint of explanatory task, permissibility involves at least three dimensions.

First, it involves exclusion. A system has a structure of permissibility, first of all, because it is not open to every logically conceivable state. Some states fail to appear not merely because no cause has yet triggered them, but because they are incompatible with the system's mode of organization, its boundary conditions, or its global relations. Their nonappearance is therefore not an accidental absence, but a structural exclusion. To explain permissibility is often to explain why certain outcomes are impossible in principle, why certain configurations cannot be stably maintained, or why certain trajectories do not belong to the kinds of change the system can undergo.

Second, it involves stability. Permissibility is not a casual listing of momentary possibilities, but an account of which states or patterns can persist in a system. If a configuration only flashes into existence briefly, cannot be maintained, and cannot preserve coherence through variation, then even if it appears once, it may still fail to count as a permissible state of the system. Permissibility therefore concerns not simply whether something can happen once, but whether it can be sustained as a form the system allows. This is also why permissibility is naturally tied to emergence, organization, and pattern-retention: in many cases, what really needs explanation is not how a state happened to appear, but why a certain class of pattern can become a stably existing form in that system.

Third, it involves consistency. To say that a state or pattern is permissible means not only that it is not causally blocked, but that it is compatible with the rest of the system's organization. A local change may be impermissible not because there is no sufficient cause for it, but because if it were realized, it would disrupt the system's overall consistency. In this sense, permissibility is not merely a matter of whether there are other causes stopping something; it is a matter of whether the form belongs to what the system can coherently bear. The focus here is no longer local pushing, but systemic coordination.

With these three dimensions in view, we can understand more clearly why permissibility is not a mere byproduct of causal production, but an independent explanatory target. Causal production asks how, among the open paths, the actual result is brought about along one particular path. Permissibility asks why the open paths are these rather than those, why some patterns are open to the system while others are excluded from the outset. The former concerns the process of realization; the latter concerns the space of realization itself. Without this distinction, we are easily led to treat every case of nonoccurrence as merely a case where the right cause has not yet been found, and every case of stable existence as merely something that has happened many times, thereby missing the explanation of the system's structure itself.

This point is crucial for a number of the issues treated later in the book. In the case of emergence, what we often need to explain is not which extra cause produces a higher-level pattern, but why a certain organizational form can be stably borne across many different micro-level realizations. That is fundamentally a question of permissibility: which micro-level differences still allow the higher-level pattern to remain, and which differences would destroy it. The same holds for quantum entanglement. If entanglement reflects a kind of structural inseparability within a joint system, then the key explanatory target is not simply how an outcome is "transmitted," but which combinations of joint outcomes are open to that system and which are structurally excluded. In the case of higher-level explanation as well, some higher-level descriptions have independent value not because they invent new causal chains, but because they articulate which changes are meaningful at that level, which patterns are sustainable, and which forms of organization are open to the system.

One possible misunderstanding should also be cleared away here: emphasizing permissibility does not mean retreating from the real world into a purely abstract modal space. Permissibility is not empty modal talk detached from

reality. On the contrary, it is always relative to a concrete system: given this system's organizational relations, boundary conditions, and structural constraint, what is open and what is not. For that reason, permissibility is not opposed to realization. It is one of the conditions that make realization intelligible. To ask what is permissible for a system is not to stop caring about what actually occurs, but to ask what sort of reality the system is structurally able to sustain.

3.5 The Cooperative Relation Among the Three Explanatory Roles

Up to this point, the chapter has defined separately the problems addressed by causal production, structural constraint, and permissibility. Causal production explains how a result is brought into reality; structural constraint explains the boundaries, organizational relations, and limiting conditions under which a system operates; and permissibility explains which states, trajectories, and patterns are open, stable, and consistent for that system under those constraints. If we stopped there, the differences among the three would already be clear enough. But if the book is genuinely to propose a multi-role view of explanation, distinction by itself is not enough. We must also explain how these three explanatory roles work together within one and the same theoretical picture, and why they are neither to be conflated nor to be treated as mutually exclusive.

The first point that must be emphasized is that these three roles are not three independent explanatory systems, each managing its own separate territory. They do not usually deal with three wholly unrelated objects. More often, when confronted with one and the same system, one and the same phenomenon, or one and the same theoretical problem, they answer different questions from different directions and at different levels. That is to say, their relation is not one in which one explanation is completed and another is then optionally added afterward. It is closer to this: if we are to achieve a full understanding of certain complex

phenomena, we often need to know at once how the result occurs, within what structural boundaries that occurrence is possible, and why these rather than other patterns are open to the system. For that reason, what this book calls the differentiation of roles is not the disassembly of explanation, but the rendering visible of its cooperative structure.

In terms of logical order, the relation among the three may be expressed provisionally as follows: structural constraint determines the operative boundaries of a system; permissibility characterizes which states, configurations, and patterns are open within those boundaries; and causal production explains how some of those open states or patterns are realized in concrete circumstances. More tersely put, structural constraint shapes a space of possibility, permissibility describes the allowed forms within that space, and causal production actualizes some of those forms. This formulation is of course schematic, but it captures a crucial point: causal production does not unfold in a vacuum. It always operates within some structurally organized space of possibility. Likewise, structural constraint and permissibility are not abstract pictures detached from real process. They have explanatory significance precisely because they define why real processes unfold as they do, and why they cannot unfold otherwise.

This relation shows, first, that although causal production occupies a central place in the explanation of events, it does not by itself always constitute a complete explanation. A result's generative pathway may be described in great detail, yet if we do not know why that pathway is open to the system, why other pathways are excluded, or why that result can stably exist within the relevant organizational structure, the explanation may still remain incomplete. Causal production tells us how something happens, but it does not always by itself tell us why only this way of happening is possible here. It is precisely at this point that structural constraint and permissibility supplement tasks not directly borne by causal production.

Conversely, structural constraint and permissibility also cannot, without causal production, yield a sufficient account of the real world. If we know only which patterns a system permits, which configurations it excludes, and which stable forms it sustains, yet know nothing about how a concrete result is realized in a particular case, then what we have is still only a framework of boundaries and possibilities, not a full understanding of the actual process of occurrence. Especially in questions involving concrete events, mechanistic operation, or state-transition, permissibility by itself does not bring a result into reality, nor does structural constraint automatically substitute for a generative pathway. They explain what is open to the system and what is excluded, not which one of the open outcomes is in fact realized this time. Without causal production, the multi-role view of explanation would lose its contact with real processes of occurrence and degenerate into purely structural description.

It follows that the relation among these three roles is neither reductive, nor simply parallel, nor competitive. They are not three steps in a ladder, as though one could always absorb the other two in principle. Nor are they three unrelated language-games arbitrarily describing the same object. A more accurate understanding is that they perform complementary rather than exclusionary tasks within explanation. Structural constraint answers questions of boundary. Permissibility answers questions of openness and stability. Causal production answers questions of realization. Only by keeping all three classes of question on the table can we avoid a long-standing misjudgment, namely, that all theoretical significance must be compressed onto the single line of how a real result is produced.

This cooperative relation also helps clarify a common confusion: are structural constraint and permissibility merely abstract generalizations of causal explanation? That is, can they ultimately be reduced to "more complex" or "more global" causal descriptions? The answer given in this book is no. The reason is not that the three

have no overlap, but that the types of problem they address are themselves different. Even if we can find the concrete physical process through which some structural constraint is realized, it still does not follow that the problem of constraint has been replaced by the problem of production. For the question how this constraint is realized is not the same as the question what that constraint opens and excludes for the system. Likewise, even if some permissible pattern must ultimately be realized through concrete causal processes, it still does not follow that permissibility is nothing more than a causal chain not yet unfolded. The question how this pattern is actualized is still different from the question why it belongs to the system's open space in the first place.

This will be especially important in the book's later core cases. In the case of emergence, if we retain only causal production, we are easily led to think of higher-level patterns either as possessing additional causal powers or as mere byproducts of lower-level events. But if structural constraint and permissibility are brought in as well, we can say more clearly that the reality of higher-level patterns lies not in what additional thing they push, but in the fact that they correspond to some stable and open form of organization, while concrete micro-level processes actualize those forms. The same is true of higher-level explanation: the legitimacy of a higher-level account does not always lie in supplementing some lower-level causal chain not yet revealed, but may lie instead in its articulation of structural boundaries and permissible variation at that level, thereby telling us what kinds of difference matter there. In the case of quantum entanglement, similarly, structural constraint and permissibility help explain why the space of joint outcomes is organized in an inseparable way, while the problem of causal production moves to a different place and no longer bears the whole explanatory burden.

A further point must be added: the cooperation of the three explanatory roles does not require that they be equally prominent in every case. Some problems are primarily productive, with structural constraint and permissibility supplying only

necessary but nonprominent background. Other problems are primarily questions of constraint and permissibility, where causal production remains indispensable but is not the explanatory focus. In other words, the multi-role view does not require that every explanation contain three equal parts. It requires, rather, that we be able to identify which role is dominant in a given problem, which is supportive, and which, though present, is not the center of explanation. Without that ability, we will still drag all accounts back under one and the same standard of evaluation, and the point of the multi-role framework will be lost.

A provisional conclusion can therefore be drawn. The reason this book distinguishes causal production, structural constraint, and permissibility is not to divide explanation into three competing territories, but to recover a division of labor that is already there in explanatory practice. Structural constraint shapes the boundaries of a system, permissibility articulates its open space, and causal production realizes concrete results within that space. What they jointly constitute is not an unnecessary and cumbersome tripartite scheme, but an explanatory picture capable of preserving at once event-generation, organizational limitation, and the stability of patterns. For that reason, the next task is no longer simply to add more concepts, but to provide a more operational set of criteria by which we can judge, in concrete analysis, when we are dealing primarily with a productive problem and when what truly requires explanation is structural boundary and open space.

3.6 Criteria of Distinction and Steps of Analysis

At this stage, the chapter has introduced and distinguished three explanatory roles: causal production, structural constraint, and permissibility. They can be separated conceptually. But if this framework is to possess genuine analytical force, definition alone is not enough. The next question is therefore this: in concrete theoretical disputes, how do we determine what kind of problem we are mainly

facing? How do we avoid miswriting questions about boundary, organization, and stability as questions of event-generation, or conversely, prematurely elevating a result that ought to be explained through a generative pathway into an abstract structural issue? Without relatively stable criteria of distinction, the multi-role view of explanation will easily degenerate into a retrospective labeling game rather than becoming a truly operational framework of analysis.

The basic position of this book is that the key to distinguishing explanatory roles lies not in what words a given account uses, but in what question it is primarily answering. That is to say, we cannot judge what role an explanation belongs to simply by asking whether it contains terms such as cause, structure, constraint, or possibility, since the same term may perform very different functions in different contexts. What matters is where the explanatory center of gravity lies: is the account explaining how a result is brought about, or what kinds of result are open to the system? Is it tracing a generative pathway, or characterizing boundary conditions? Is it identifying difference-makers, or showing which differences are irrelevant at the relevant level? In exactly this sense, the distinction among explanatory roles is, first of all, a distinction among types of question.

On that basis, a preliminary set of criteria may be proposed.

The first is the criterion of generativity. If the core task of an account is to identify the activities, processes, mechanisms, or difference-making factors through which a result comes into being, and by means of those factors to explain how the result is brought into reality in the concrete case, then the account belongs primarily to causal production. The emphasis here falls on brought into reality: explanatory success depends on whether the account reveals how the result occurred, and how it would differ if the relevant conditions differed. If an account loses its main explanatory content once the generative pathway is removed, then it is a productive account.

The second is the criterion of exclusion. If the emphasis of an account lies not on telling how a result was produced, but on explaining why a large range of other configurations, trajectories, or states is simply not open to the system, why certain forms are impossible, unsustainable, or inconsistent, then the account belongs primarily to structural constraint. What is doing the real work here is not which factor pushed the result, but what boundary systematically excludes certain results. If the primary force of an account lies in exclusion rather than production, then it should not be misread as an unfinished causal narrative.

The third is the criterion of openness. If an account mainly answers the question what kinds of state, pattern, or form of change are allowed for this system, especially when it explains why some class of pattern can remain stable across local variation, then its explanatory target is permissibility. The concern here is not that some particular result happened to occur, but why a given form belongs to the system's open space in the first place. If the center of an account lies in defining open space rather than tracing a single realization, then it should not simply be assimilated to causal production.

The fourth is the criterion of level-relevance. Many confusions arise not from the concepts themselves, but from a failure to distinguish what counts as a relevant difference at different levels. In productive problems, we usually care about which differences would change whether the result occurs. In permissibility-problems, by contrast, we are more often concerned with which differences would not disrupt whether some pattern or structure still holds. The former requires us to identify difference-makers; the latter to identify a range of tolerance for difference. If the focus of an account lies in characterizing under what variations a pattern still holds, then it is closer to structural constraint and permissibility than to causal production alone.

The fifth is the criterion of loss upon failure. Sometimes the clearest way to judge what an account is doing is to ask: if this account failed, what understanding

would we lose? If what we would lose is our grasp of how the result occurs, then the account is mainly productive. If what we would lose is our grasp of why only these rather than those forms are open to the system, then it is mainly a constraint-based or permissibility-based account. This criterion is useful because it forces us to judge the type of explanation from its function rather than from the surface vocabulary it happens to employ.

On the basis of these criteria, the book recommends a more explicit sequence of analysis. When confronted with a concrete theoretical problem, one may proceed by asking the following questions in order.

First: What is the explicit form of the question being asked?

Is it framed as "How does some result occur?" or as "Why is some class of result, pattern, or organizational form possible, stable, or sustainable?" A great many theoretical misjudgments arise already at this stage, because the problem has often been rewritten before analysis even begins.

Second: What is the explanatory center of the account at issue?

Is it mainly identifying a generative pathway, or is it mainly defining boundaries, limitations, and open space? If most of the force of the account comes from excluding impossible states rather than tracing a process of realization, then it more likely belongs to structural constraint.

Third: What type of differences are relevant here?

Are they differences that would change whether the result occurs, or differences that would not destroy whether a pattern still holds? The former usually point toward causal production; the latter often point toward structural constraint and permissibility.

Fourth: Would rewriting the account as a productive explanation distort the original problem?

If, after such rewriting, we lose our grip on stability, organization, boundary conditions, or open space, then the original problem is not primarily a productive

one. This step is especially useful for identifying the miswritings generated by explanatory monism.

Fifth: Which of the three roles is dominant in the present case?

Not every case distributes the three roles equally. In some, causal production is central and structural constraint merely supportive; in others, the reverse is true. The point of analysis is not to force every problem into three equal pieces, but to identify dominant and supportive roles.

If we look back at the disputes discussed in the earlier chapters through this analytical sequence, the necessity of the multi-role framework becomes more visible. If the problem of emergence is framed only as “What extra cause produces the higher-level pattern?”, then the original problem of stable pattern and space of realization has already been rewritten. If the exclusion problem looks only at whether lower-level causes are already sufficient, then explanatory roles at different levels are compressed into one competitive position. If the problem of entanglement asks only what propagates the correlation, then the structural inseparability of the space of joint outcomes is obscured. These disputes are persistent not only because the answers are difficult to find, but because the types of question at stake are often never properly distinguished.

Of course, these criteria are not a mechanical algorithm. They cannot guarantee that every boundary case is resolved at once, nor can they automatically classify complex theoretical problems. Their function is more like that of a diagnostic toolkit, aimed at reducing two of the most common errors: first, compressing all explanatory tasks prematurely into productive tasks; second, hastily elevating any account not easily reducible to a productive pathway into an independent ontological level. The multi-role view is meant to avoid exactly these two opposite yet equally crude reactions.

What is proposed in this section, then, is not a rigid partition of explanation, but a more restrained analytic discipline. It asks us, before entering a concrete

dispute, first to distinguish what question we are actually asking; before evaluating an account, first to see what task it chiefly performs; and before claiming that two accounts conflict, first to determine whether they are in fact competing for the same explanatory position. Only in that way will the distinction among causal production, structural constraint, and permissibility cease to remain at the level of conceptual description and become an analytical tool usable in the chapters that follow.

3.7 Boundary Cases and Common Confusions

The preceding sections have introduced and distinguished the three explanatory roles of causal production, structural constraint, and permissibility, and have also provided the corresponding criteria and steps of analysis. In actual theoretical and scientific contexts, however, explanations do not always appear in “pure form.” Many important accounts simultaneously contain productive, constraining, and permissibility-based components, and may even perform different roles at different levels of analysis. If the multi-role framework is to avoid being misunderstood as a rigid classificatory system, it must confront these boundary cases directly.

The aim of this section is precisely to explain how, when an account exhibits more than one feature at once, we should identify its dominant explanatory function; and under what conditions explanatory roles are wrongly conflated in ways that generate theoretical misjudgment.

3.7.1 Hybrid Explanations: The Normal Form of Multi-Role Coexistence

The first point that must be made clear is that mixture is not the exception, but the normal condition of explanatory practice. In many mature accounts, causal production, structural constraint, and permissibility do not appear separately. A mechanistic explanation typically not only exhibits a generative pathway, but also

relies implicitly on organizational relations, boundary conditions, and stratified arrangement. An account centered on stable pattern may have permissibility as its core, yet still require concrete processes to bear the task of realization. The real issue therefore is not whether an explanation is “pure,” but which role constitutes its dominant function when multiple roles are present. The point of the multi-role framework is not to hack complex explanations into pieces, but to help us identify their functional center of gravity.

3.7.2 First Confusion: Miswriting Structural Constraint as Hidden Causation

One of the most common errors is to translate structural constraint back into some as-yet-undiscovered causal mechanism. Boundary conditions are then treated as “causes not yet found,” structural limitations as “hidden driving forces,” and impossibility as if it were merely “the absence of some triggering condition.” The consequence of this rewriting is that the original question, why some situations are not open to the system at all, is dragged back into the question what exactly caused them to happen. The fundamental error here is to mistake exclusion for nonrealization. But what structural constraint often really expresses is not that some result happened not to occur, but that it cannot occur in that system in the first place. Once this distinction is flattened, explanation is systematically dragged back into the old framework of causal production.

3.7.3 Second Confusion: Reducing Permissibility to Frequency or Statistical Outcome

Another common mistake is to understand permissibility as whatever occurs often, displays statistical stability, or has a high probability of appearing. This may seem natural, but in fact it rewrites structural openness as empirical frequency. Permissibility is concerned not with how many times a pattern has already

appeared, but with whether it belongs to the system's open space in the first place. A form may occur frequently and yet be unstable, and therefore fail to count as permissible in the structural sense; conversely, it may occur infrequently and yet still fall within the structural range allowed by the system. If permissibility is reduced directly to statistical outcome, its significance as an independent explanatory target disappears.

3.7.4 Third Confusion: Expanding Causal Production into the Sole Final End of Explanation

The third confusion is also the deepest one. Even while admitting that structure and permissibility matter, one still insists that genuine explanation must ultimately return to causal production. This view often appears in milder form, for example, in the claim that structural explanation is merely higher-level causal explanation, that constraints are abstract summaries of causal relations, or that every mature account must finally explain how it is realized. The problem is not that such claims are wholly false, but that they mistake a relation of realization for a relation of explanatory belonging. A structure may indeed need to be realized through causal process, but it does not follow that its explanatory function is identical with that process. It is precisely this slide that allows explanatory monism to persist in a more concealed form even after it has outwardly accepted diversity.

3.7.5 Boundary Case: The Same Account Performing Different Roles at Different Levels

A more complex and more common situation is that one and the same account performs different roles at different levels of analysis. A given mechanism may, at the local level, function as causal production; at the global level, it may function as structural constraint; and at a still higher level, it may help define which patterns a system can stably sustain. If these levels are not distinguished, two misjudgments

are easy to make: either one treats such multilayered articulation as redundant repetition of the same explanation, or one mistakes it for conflict among multiple explanations. In reality, the situation is better understood as the same system presenting itself in multiple ways across different dimensions of questioning. The advantage of the multi-role framework lies precisely in the fact that it permits such stratified difference, rather than demanding that all accounts be flattened into one and the same function.

3.7.6 A Methodological Warning: Avoiding Two Symmetrical Errors

On the basis of the foregoing analysis, two extremes that must be simultaneously avoided come into sharper view. The first is causal reductionism, which compresses all explanation into generative pathways and demotes structure and permissibility to subordinate language. The second is structural inflationism, which treats structural constraint as an independent force and hastily reifies permissibility into a new ontological level. The reason this book insists on the multi-role view of explanation is precisely to avoid these two symmetrical errors: neither dissolving structural explanation into causation, nor mystifying it into a new entity, but understanding both as different forms of account corresponding to different kinds of problem.

3.7.7 Summary: From Classification to Analytic Discipline

What these boundary cases show is that the point of the multi-role framework does not lie in producing a static table of classifications, but in establishing an analytic discipline. It requires that we not presuppose that all problems are productive problems, that we not evaluate all accounts by one single standard, and that we not forcibly drag different roles into the same competitive structure. More importantly, it continually compels us to return to the most basic question: what

exactly am I explaining right now? Only if that question is seriously preserved will the differentiation of explanatory roles cease to remain at the level of conceptual maneuver and begin genuinely to alter the way theoretical judgment is made.

3.8 Chapter Summary: From Explanatory Roles to the Reconstruction of Problems

The central task of this chapter has been to reconstruct explanatory activity itself. Unlike the diagnostic analyses of the first two chapters, this chapter has no longer remained at the level of identifying the problem of explanatory monism. It has instead attempted to propose a more stable framework of distinctions in order to show what kinds of question different kinds of explanation are answering, and why they can coexist within the same theoretical picture without excluding one another.

On the basis of the foregoing analysis, the central conclusions of the chapter may be summarized in three points.

First, explanation is not a single task, but contains at least three basic roles that must not be conflated: causal production, structural constraint, and permissibility. Causal production explains how a result is brought into reality; structural constraint explains under what boundaries and organizational relations a system operates; and permissibility explains which states, configurations, and patterns are open to the system under those constraints. The three correspond respectively to the dimensions of generation, limitation, and openness, and together they constitute the basic structure of explanatory space. Any theory that attempts to compress these three into one single activity will inevitably produce systematic mismatch in certain kinds of problem.

Second, the relation among these three roles is not one of reduction or competition, but one of cooperation. Structural constraint shapes the boundaries of a space of possibility, permissibility articulates stability and exclusion within that space, and causal production actualizes concrete results within it. A complete

explanatory picture often requires the joint operation of all three roles: one must know not only how something happens, but also why only these ways of happening are open to the system, and why certain patterns remain stable across variation. In precisely this sense, the multi-role view of explanation is not a weakening of causal explanation, but a redefinition of its place.

Third, the distinctions proposed in this chapter are not merely conceptual divisions, but a form of analytic discipline. Through the criterion of generativity, the criterion of exclusion, the criterion of openness, and the corresponding steps of analysis, we become better able to identify, in any concrete problem, which explanatory role is dominant and which are supportive; which accounts are answering questions of occurrence and which are articulating boundaries and spaces of possibility; and under what conditions different accounts are not genuinely in conflict, but merely pulled mistakenly into the same evaluative framework. It is in this respect above all that the chapter provides an operational methodological basis for what follows.

If these conclusions are compressed into a more general judgment, then one may say this:

Many contemporary theoretical deadlocks arise not because the structure of the world itself contains irreconcilable conflicts, but because different kinds of problem are repeatedly miswritten as one and the same explanatory task.

Once the differentiation of explanatory roles is acknowledged, those deadlocks can be redescribed. The issue is no longer whether some level possesses additional causal power, nor whether different accounts are competing for the same explanatory position, but whether, in a given case, we are asking about a generative pathway, a structural boundary, or the open form of a space of possibility. It is precisely in this redescription that many disputes that previously seemed entrenched begin to show the possibility of reorganization.

This also provides a direct path for the chapters that follow. The next task is no longer to continue expanding the conceptual framework, but to apply it to concrete problems in order to test its explanatory force. Chapter 4 turns to the problem of emergence, for it is there that the consequences of explanatory monism are displayed most concentratedly: higher-level patterns are either required to possess independent causal power or dismissed as lacking genuine standing. Yet from the standpoint of differentiated explanatory roles, the structure of the problem changes fundamentally. Emergence no longer appears first of all as a problem of “new causes,” but more nearly as a problem of stable patterns and permissible structure; and the cooperative relation among causal production, structural constraint, and permissibility will be exhibited there in its most direct form.

The conclusion of this chapter is therefore not merely a conceptual distinction, but a reconstruction of problems. What it provides for the subsequent analysis is not only a new vocabulary, but a new way of asking questions. In this sense, Chapter 3 constitutes the turning point of the book as a whole: from here onward, explanation is no longer understood as a single productive task, but is reopened as a multidimensional structure, and all the concrete topics that follow will be repositioned within that structure.

Chapter 4

The Reinterpretation of Emergence: Structural Novelty Rather Than New Causal Power

The problem of emergence retains its special place in contemporary philosophy of science not merely because it concerns complex systems, layered structures, or higher-level patterns, but because it repeatedly touches a deeper difficulty: how can we acknowledge the reality and explanatory importance of higher-level phenomena without giving up physical continuity and causal closure? Whenever we confront phase transitions, self-organization, biological organization, cognitive function, or social structure, we encounter the same tension again and again. On the one hand, such phenomena clearly depend on lower-level realization. On the other hand, they do not seem to be mere lists of lower-level processes, but exhibit new forms of organization, new stable relations, and new theoretical significance. The question is thus forced upon us: if this “newness” is real, in what does it consist?

The standard answer tends to push this question quickly into causal language. If higher-level phenomena are to count as genuinely new, then they seem to have to possess some additional causal efficacy; if they do not introduce new causal power, then they are easily treated as descriptive convenience, computational compression, or merely surface phenomena from the standpoint of an observer. The central claim of this chapter is that this very move is open to question. Emergence has long appeared mysterious not because the world forces us to choose between “additional causal power” and “explanatory nihilism,” but because novelty has from the outset been miswritten as causal novelty. Once that presupposition is relaxed, we can turn to a more robust understanding: higher-level phenomena are new not because they add new movers to the world, but because they form new structural facts, new stable boundaries, and new

permissible forms of organization. Emergence, in other words, is first of all not a problem about new causes, but a problem about structural novelty.

This chapter develops that line of thought. First, it explains why emergence continues to exert theoretical pressure rather than being a pseudo-problem generated merely by terminological preference. Second, it surveys the standard dilemma between strong emergence and weak emergence, and identifies the shared presupposition underlying both. Third, it introduces structural novelty as an alternative concept and, with the aid of the framework of causal production, structural constraint, and permissibility developed in the previous chapter, shows how higher-level phenomena can both depend upon lower-level realization and yet not be reduced to mere descriptive convenience. The final aim of the chapter is neither to deny the reality of the higher-level nor to appeal to additional causal powers, but to restate what higher-level novelty ought to mean.

4.1 Why the Problem of Emergence Continues to Exert Theoretical Pressure

The reason emergence keeps returning to the center of philosophy of science is not that the word itself is vague, but that it points to a class of real and stubborn phenomena. Across many domains, we find that simply grasping the local details of lower-level processes is not sufficient automatically to yield a full understanding of the behavior of the system as a whole. Certain higher-level patterns are not only identifiable, but play an irreplaceable role in explanation, classification, prediction, and intervention. When we speak of “emergence,” we are often not indulging a taste for exaggerated complexity. We are responding to the fact that some phenomena do display the following feature: they depend on the lower level, and yet cannot simply be understood as a term-by-term sum of lower-level processes.

This pressure arises first from the coexistence of two demands. We are unwilling to deny the reality of higher-level phenomena, because scientific inquiry

repeatedly shows that macroscopic patterns, organizational structures, and systemic behavior are not dispensable linguistic surfaces. At the same time, we are also unwilling lightly to give up scientific continuity and the principle of physical closure, or to mystify higher-level phenomena into additional existents independent of lower-level realization. The core tension of the problem of emergence thus becomes visible: how to preserve dependence while also preserving novelty, how to insist on a basis of realization without demoting higher-level patterns to mere observational convenience.

From the standpoint of scientific practice, this tension is not abstract. A phase transition is not a simple sum of the positions of microscopic particles; biological organization is not an arbitrary label attached to chains of chemical reaction; cognitive function is not merely an intuitive nickname for neural activity. In such cases, what we identify are higher-level forms of behavior, stable patterns, and organizational relations. They matter not merely because they are “more convenient to describe,” but because if we retreat entirely to lower-level detail, we actually lose our grip on certain crucial structures. It is in exactly this sense that the problem of emergence keeps returning: it forces us to acknowledge that some theoretically indispensable objects are grasped at the level of higher-order structure and pattern.

But once this is acknowledged, a new question follows. If such higher-level patterns do indeed possess independent explanatory standing, what exactly does that independence amount to? Must it depend on some additional causal power? Or can it be grounded in something else, such as new organizational boundaries, new stable relations, and new forms of system-level relevance? So long as this question remains unclear, emergence will continue to occupy an unstable position: it refuses to be absorbed into a simple reductive picture, yet has not yet found a positive formulation that avoids mystification.

The reason emergence continues to exert theoretical pressure, then, is not that the world has suddenly produced a class of incomprehensible phenomena, but that it forces us to answer a more basic question anew: what does “newness” in reality actually mean? If that question continues to be presupposed as “whether new causes have been added,” emergence will inevitably remain troublesome. If the question can be reformulated, the subsequent theoretical terrain will change with it.

4.2 The Standard Dilemma Between Strong and Weak Emergence

The standard literature usually organizes the problem of emergence around two main positions. The first is strong emergence. This view attempts to preserve the reality and irreducibility of higher-level phenomena, and so tends to hold that certain higher-level properties possess not only descriptive distinctness but also some irreducible ontological or causal standing. On this way of thinking, if higher-level phenomena are genuinely new, then they cannot be mere coarse-grained summaries of lower-level realization; they must in some way exceed the full resources of lower-level account.

The attraction of strong emergence lies in the fact that it refuses lightly to dissolve the higher level. It captures an important intuition: if higher-level patterns repeatedly become theoretically central in scientific practice, then it seems too hasty to demote them to mere cognitive shortcuts. Yet the problems of strong emergence are equally evident. It often requires appeal to additional causal powers, independent laws, or some stronger form of higher-level ontological standing, and those commitments sit uneasily both with physical closure and with the demand for empirical support. Strong emergence succeeds in preserving novelty, but often at the cost of too much.

The second position is weak emergence. It aims to preserve realization-dependence, physical closure, and scientific continuity, and therefore rejects additional causal power. Weak emergence typically holds that higher-level phenomena seem novel mainly because systems are complex, because computation is difficult, because levels of description differ, or because macroscopic behavior cannot be directly read off from a limited cognitive standpoint. Higher-level novelty is thereby retained, but mainly in an epistemic or methodological sense rather than as a stronger form of structural reality.

The advantage of weak emergence is that it avoids the risk of mystification faced by strong emergence, while preserving the general framework of naturalism and physicalism. But its difficulty is that it too easily weakens the reality of higher-level phenomena into “we cannot calculate it yet” or “it is more convenient to speak at this level.” Once that happens, the independent explanatory standing of higher-level patterns becomes fragile: they begin to appear as projections of lower-level complexity rather than as organizational facts with theoretical weight of their own.

Thus the standard dilemma takes shape. Strong emergence preserves reality, but at too high a cost; weak emergence preserves continuity, but at the risk of sacrificing reality. More importantly, this dilemma is not merely an empirical comparison between two options. It reflects a deeper shared presupposition: both sides assume that if some phenomenon is genuinely new, then in some sense it must be causally new. Strong emergence accepts this head-on, while weak emergence, in rejecting it, is forced to attenuate novelty. The next section argues that this common presupposition itself is open to question. For once novelty is no longer presupposed to be causal novelty, the force of the standard dilemma is substantially weakened, and structural novelty can emerge as a genuine third path.

4.3 The Miswriting of Novelty: From a Structural Problem to a Causal Problem

The standard opposition between strong and weak emergence appears, on the surface, to offer two different answers to one and the same question. The former emphasizes the irreducible reality of higher-level phenomena; the latter emphasizes their dependence on lower-level realization and the continuity of theory. Yet the persistence of this opposition does not arise only from the fact that each side captures part of what is right. It persists because both share a deeper presupposition that is rarely made explicit: if higher-level phenomena are genuinely new, then that newness must in some sense be causal newness. Strong emergence accepts this premise and draws the stronger conclusion from it; weak emergence, unwilling to accept the conclusion, preserves continuity by weakening the status of novelty. The result is a familiar dilemma whose real force lies not merely in the alternatives it offers, but in the common frame within which both alternatives are posed.

From the standpoint of this book, it is precisely this framing that needs to be challenged. Why should novelty be understood from the outset as a matter of causal addition? Why assume that, if something is theoretically irreducible, it must therefore add a new producer, a new power, or a new push within the causal order? This assumption may feel natural only because explanatory monism has already taught us to regard whatever is genuinely important in theory as something that must ultimately occupy a place in productive explanation. Once that background is in place, newness seems to have only two options: either it enters the world by adding causal efficacy, or it shrinks back into epistemic appearance.

But this is exactly where the problem is miswritten. Many higher-level phenomena appear novel not because they introduce a new causal ingredient into the world, but because they instantiate new forms of organization, new relations of

stability, new structural boundaries, and new permissible patterns. Their novelty lies not in pushing the world differently, but in structuring it differently. What changes at the higher level is not necessarily the inventory of causes, but the pattern of organization under which processes become intelligible, stable, and theoretically significant. If this is right, then the real question is not whether higher-level phenomena add new causal powers, but what kind of structural fact they embody.

This point can be stated more sharply. To ask whether emergence involves “new causal power” is already to translate the problem into the idiom of causal production. Yet many of the cases that motivate talk of emergence are not first of all cases in which we discover a mysterious new producer. They are cases in which we encounter a level of organization that cannot be adequately captured by term-by-term listing of lower-level details alone. A phase transition is not puzzling because a new force suddenly appears in the system; it is puzzling because a new stable regime of behavior emerges. Biological organization is not important because it introduces a new occult driver into chemistry; it is important because it establishes forms of regulation, integration, and persistence that are not exhausted by a catalogue of local reactions. The same is true of cognition, social organization, and many other higher-level phenomena. The novelty at stake is organizational and structural before it is anything else.

Once this is seen, the standard dilemma begins to loosen. Strong emergence looks compelling only so long as one thinks that reality requires causal addition. Weak emergence looks deflationary only so long as one thinks that, absent causal addition, what remains can be only epistemic convenience. But if novelty can be structural rather than causal, then these are no longer the only options. One can say that higher-level patterns are real and theoretically indispensable, not because they intervene as extra causes, but because they mark stable organizational forms, delimit new structural boundaries, and define new spaces of permissibility. Such

patterns may be fully dependent on lower-level realization and yet still count as genuinely new in the only sense that matters here: they disclose structural features that are not captured by merely restating lower-level events one by one.

This is why the book proposes to rewrite emergence in terms of structural novelty. Structural novelty does not mean that a new kind of substance has appeared, nor that some higher level floats free of its realizational basis. It means that, under certain conditions, a system comes to exhibit a form of organization, stability, or pattern-governed openness that is not adequately understood merely by tracing local production alone. The point is not to oppose structure to process, but to refuse the reduction of all novelty to process alone. Causal production remains relevant, because any higher-level pattern must still be realized through concrete processes. But what is novel may lie in the structural configuration those processes sustain, not in the introduction of a further producer over and above them.

To put the matter in the terms developed in the previous chapter: what has been miswritten as a question of causal production is often more properly a question involving structural constraint and permissibility. We ask whether some higher-level pattern is “real,” and then immediately demand to know what extra cause it contributes. But perhaps the better question is this: what new constraints, stable forms, and permissible patterns characterize the system at that level? If those are the explanatorily salient features, then causal novelty is not the right measure of reality. The relevant novelty lies instead in the emergence of a new structured space within which processes are organized and outcomes are stabilized.

This also explains why emergence so often seems both undeniable and elusive. It is undeniable because higher-level patterns really do matter: they pick out forms of organization, stability, and relevance that cannot simply be ignored. It is elusive because we keep asking the wrong question about them. If we ask what

new causal power they add, we either inflate ontology or deflate the higher level into convenience. If instead we ask what new structural facts they disclose, then the phenomenon becomes much less mysterious. We no longer need to choose between metaphysical excess and explanatory emptiness. We can say that higher-level phenomena are new because they instantiate new stable organizations of the space of possibility.

The rest of this chapter develops this proposal in more concrete form. The next step is to say more precisely what structural novelty amounts to, and how it allows us to preserve both dependence and explanatory significance without appealing to additional causal powers.

4.4 Structural Novelty as an Alternative Concept

If the novelty at issue in emergence has been miswritten as causal novelty, then the central task is to formulate more precisely what an alternative notion of novelty would look like. This book proposes structural novelty for that purpose. By structural novelty I mean not the appearance of a new force, a new producer, or a new ontological ingredient added to the world, but the appearance of a new form of organization, a new regime of stability, or a new configuration of constraints and permissibilities that cannot be adequately captured by a mere itemized restatement of lower-level goings-on.

The point of this concept is not to deny dependence on the lower level. On the contrary, structural novelty presupposes such dependence. A structurally novel phenomenon does not float free of its realization base; it is realized through it. But realization does not exhaust significance. A higher-level pattern may depend entirely on lower-level processes and yet still be novel because what matters explanatorily is not simply that those processes occur, but that they sustain a new stable organization, a new boundary-structure, or a new open pattern-form. In such a case, the novelty lies in the fact that the system has become intelligible under a

new structural description that is not replaceable by a simple enumeration of local events.

This is precisely where the distinction among causal production, structural constraint, and permissibility becomes useful. Structural novelty does not primarily mean that a new result is produced by a new cause. It means that a system now exhibits a new configuration of structural constraint and permissibility. Certain pathways are now stabilized, certain configurations excluded, certain patterns sustained, and certain forms of variation tolerated in ways that mark a distinct level of organization. The novelty is therefore not first a novelty of causal source, but a novelty of structured space.

This can be put in another way. A lower-level description may tell us what processes occur, and a causal account may explain how a given outcome is realized. But neither of these by itself always tells us why the system now supports a stable pattern with its own organizational integrity, why certain variations no longer matter at the higher level, or why a new class of behavior becomes possible only under a certain form of global organization. Structural novelty names this additional fact: the fact that the system has acquired a new organization of relevance, stability, and openness.

For that reason, structural novelty avoids the two main difficulties of the standard options. It avoids the inflationary burden of strong emergence because it does not require the addition of new causal powers or independent substances. And it avoids the deflationary weakness of many weak-emergence views because it does not reduce higher-level reality to mere epistemic difficulty or descriptive convenience. A pattern can be structurally novel while remaining fully realized by lower-level processes. Its novelty lies in the emergence of a new stable organizational fact.

This also helps explain why higher-level patterns can be both dependent and irreducible in the sense that matters. They are dependent because they require

lower-level realization. They are irreducible not because nothing lower-level is involved, but because the explanatory role they play cannot be discharged simply by listing lower-level events one by one. What must be captured is the stable organizational structure sustained across those events. In that respect, irreducibility here is not causal independence, but structural indispensability.

The concept of structural novelty also clarifies why emergence is so often tied to robustness across variation. A pattern is often treated as emergent not because it appears once as an isolated anomaly, but because it persists across changes in micro-level detail, so long as certain structural conditions are preserved. That is exactly the kind of case in which permissibility and structural constraint become central. What needs explanation is not merely how one token outcome occurred, but why a type of pattern remains open and stable across a range of realizations. Structural novelty therefore belongs less to the logic of extra causation than to the logic of stabilized organization.

None of this means that causal production becomes irrelevant. Any emergent pattern still has to be realized somehow, and that requires concrete processes. But those processes are not the whole story. They explain how the pattern is instantiated in a given case; they do not by themselves explain why that pattern counts as a stable higher-level form, why it tolerates certain lower-level differences, or why it occupies a distinctive place in the system's organization. Structural novelty is the name for that further explanatory fact.

In this sense, the reinterpretation of emergence proposed here is neither eliminativist nor mysterious. It does not eliminate the higher level, because it grants real explanatory standing to higher-level structure. And it does not mystify the higher level, because it does not turn that standing into an extra causal power. It instead relocates novelty from the domain of additional production to the domain of new structural organization.

The next step is to show more concretely how this account lets higher-level phenomena be both dependent on lower-level realization and yet more than mere descriptive convenience. That requires spelling out how structural novelty supports explanatory legitimacy without collapsing back into causal competition.

4.5 Dependence on Lower-Level Realization Without Reduction to Descriptive Convenience

Once emergence is reinterpreted in terms of structural novelty, the next task is to show how higher-level phenomena can remain fully dependent on lower-level realization without being reduced to mere descriptive convenience. This is the point at which many accounts falter. If one stresses dependence too strongly, the higher level begins to look dispensable; if one stresses novelty too strongly, the higher level begins to look ontologically inflated. What is needed, therefore, is a way of understanding dependence that does not automatically collapse into elimination.

The first point is that realization-dependence is not the same thing as explanatory redundancy. To say that a higher-level pattern is realized by lower-level processes is to say that it does not exist independently of those processes. But it does not follow from this that every explanatory question about the higher-level pattern has already been answered once the lower-level processes are listed. Realization tells us what bears the pattern; it does not by itself tell us what makes that pattern explanatorily salient, what structural boundaries stabilize it, or what space of permissible variation it occupies. Dependence, in other words, is a relation of ontological bearing, not automatically a relation of explanatory exhaustion.

This distinction matters because many reductive arguments move too quickly from the claim that the lower level is sufficient for realization to the claim that the lower level is sufficient for explanation. But sufficiency for realization does not by

itself yield sufficiency for every explanatory task. A system may be fully realized by local processes while still exhibiting a stable higher-level organization that must be characterized at its own level if we are to understand why some differences matter and others do not, why certain forms persist, and why certain patterns are open across multiple realizations. The higher level is not explanatory surplus simply because it is realized below.

A second point is that descriptive convenience is too weak a category for many of the cases at issue. A merely convenient description is one we may adopt because it is easier to handle, shorter to state, or better suited to our practical limitations, while the underlying structure of the world itself remains fully captured elsewhere. But many higher-level patterns are not like that. They are not merely shorthand for an already understood lower-level story. They identify stable structures, recurrent organizations, and ranges of relevant variation without which the lower-level story itself does not yet tell us what matters. In such cases, the higher-level description is not merely useful because of our limitations; it is useful because it tracks a genuine structural feature of the system.

This is where structural constraint and permissibility do crucial work. What makes a higher-level pattern more than convenience is not that it introduces a second causal chain over and above the lower-level one, but that it marks a stable regime of organization. It shows which lower-level differences are absorbed without destroying the pattern, which are boundary-crossing and therefore disruptive, and what class of forms remains open to the system under a given organization. A higher-level phenomenon can thus be fully dependent on lower-level realization while still carrying independent explanatory weight, because what it contributes is not extra production but structured relevance.

This also clarifies why higher-level explanation often survives variation in micro-level detail. If higher-level patterns were merely coarse abbreviations for one exact lower-level sequence, then any change in micro-level realization would threaten

their identity and standing. But in many scientific domains, higher-level patterns remain explanatorily stable across many such changes. That stability is not magic. It reflects the fact that the explanatory target is a structurally organized pattern whose permissibility extends across a range of realizations. The higher level therefore has standing not because it floats above realization, but because it is borne across multiple realizations in a way that itself requires explanation.

At this point, the contrast with both strong and weak emergence becomes sharper. Strong emergence takes reality to require additional causal contribution; weak emergence often treats higher-level novelty as a consequence of epistemic limitation. The present account rejects both moves. A higher-level pattern can be real because it corresponds to a stable structural fact, and novel because that fact marks a new organization of constraints and permissibilities, while still being fully dependent on lower-level realization. What matters is not whether the higher level causes over and above the lower level, but whether it identifies a structurally indispensable way the system is organized.

This is also why the language of “mere appearance” is too quick in such cases. A phenomenon is not demoted to appearance simply because it depends on something more basic. Many real features of the world are dependent in that sense. The crucial question is whether the feature marks a stable and explanatorily indispensable structure, or whether it is merely an arbitrary way of grouping lower-level details. Where structural novelty, stable constraint, and permissibility are present, the case is no longer one of arbitrary grouping. It is a case of a higher-level form that the system genuinely sustains.

What this gives us is a more disciplined account of higher-level reality. The higher level is not a second world, not an extra ontological inventory, and not a ghostly layer of causes. But neither is it just a convenient fiction. It is a real level of structural organization, borne by lower-level processes and yet not explanatorily exhausted by them. In that sense, emergence can be understood as the

appearance of a new stable form within the realized organization of the world, not as the irruption of a new causal force into it.

The next step is to draw out what this means for the legitimacy of higher-level explanation itself. If emergence is structurally novel in this sense, then higher-level explanation need not defend itself by claiming independent causal power. It can instead claim legitimacy by identifying stable organizational facts, constraint-structures, and permissible patterns that matter explanatorily in their own right.

4.6 The Explanatory Standing of Higher-Level Phenomena

If higher-level phenomena can be dependent on lower-level realization without being reduced to mere descriptive convenience, then the next question is what exactly grounds their explanatory standing. On the view defended in this chapter, the answer is not that they contribute an additional causal force, but that they identify a stable form of organization that is explanatorily indispensable for understanding the system. Their legitimacy lies in the kind of question they answer.

A higher-level explanation is often legitimate because it tells us what organizational pattern is in play, what structural boundaries define the relevant regime, what forms of variation the system tolerates, and what patterns remain open under those conditions. These are not trivial supplements to a more fundamental lower-level story. In many cases, they are the very features that make the system intelligible as the kind of system it is. Without them, the lower-level story may still describe processes, but it does not yet tell us which processual differences matter, which may be ignored, and why a stable pattern persists as a unit of scientific understanding.

This means that the explanatory standing of the higher level is not parasitic in the pejorative sense. It does not merely borrow all of its legitimacy from a lower-level productive account while adding nothing of its own. Rather, it occupies a different role within explanation. The lower level may tell us how a result is realized;

the higher level may tell us why the result belongs to a stable pattern, why the pattern can survive local variation, and why certain forms of organization are relevant for explanation in the first place. These are not rival answers to one and the same question. They are answers to different questions about one and the same system.

Seen in this way, the higher level gains explanatory standing through structural indispensability. If removing the higher-level description would leave us unable to characterize the stability, organization, and permissible variation of the phenomenon, then that description is not merely convenient. It is explanatorily doing work that is not already performed by a simple recounting of lower-level processes. What matters is not that the higher level adds another causal push, but that it marks another explanatory dimension of the same reality.

This also explains why higher-level explanations can be both selective and objective. They are selective because they do not track every lower-level detail; they abstract away from many micro-differences. But they are objective because this abstraction is not arbitrary. It is guided by the actual structural organization of the system, by the fact that certain differences are irrelevant to the persistence of the pattern while others are not. The higher level is therefore not just a perspective we happen to adopt. It is a level at which the system itself sustains a stable explanatory pattern.

At this point, the pressure that originally drove emergence toward the language of new causal power becomes easier to diagnose. We reach for that language when we assume that only causal production can confer explanatory dignity. If that assumption is relaxed, a different possibility becomes visible. Higher-level phenomena do not need to enter the world as new producers in order to count as real or explanatorily serious. They need only correspond to stable structural facts that cannot be omitted without loss of understanding. Their

explanatory standing is thus secured not by causal addition, but by organizational necessity.

This does not mean that every higher-level description automatically has such standing. The framework defended here is not a blanket license for all macro-level talk. Some higher-level descriptions may indeed be loose heuristics, pragmatic summaries, or temporary classificatory devices. The point is rather that we now have a clearer standard for telling the difference. Where a higher-level account identifies stable structural novelty, organizes a space of permissible variation, and captures constraints that matter across realizations, it has genuine explanatory standing. Where it does not, it may remain merely useful language. The distinction must therefore be earned, not assumed.

This makes the emergence debate far less mysterious than it often appears. We do not need to choose between saying that higher-level patterns are only convenient appearances and saying that they are new causal agents. We can say instead that they are real insofar as they mark stable forms of organization that matter for explanation, and that their novelty is structural rather than causal. That is enough to secure both their reality and their theoretical importance without forcing us into ontological inflation.

In this sense, the reinterpretation of emergence offered here does more than solve a local dispute. It shows how a multi-role view of explanation can reorganize a whole field of argument. Once explanatory roles are differentiated, many of the pressures that once seemed metaphysically deep turn out to have been generated by a misdescribed explanatory demand. Emergence is no longer the name of an impossible choice between extra causation and deflation. It becomes the name of a problem about how new stable structures arise within a realized world.

The transition to the next chapter is therefore straightforward. If higher-level phenomena can possess explanatory standing without claiming independent causal power, then the broader question immediately follows: how should higher-

level explanation itself be understood? Chapter 5 takes up that question directly, asking how higher-level explanation can be legitimate once we move beyond exclusion arguments and the monopoly of mechanistic explanation.

4.7 Chapter Summary: Reinterpreting Emergence Through Structural Novelty

The central task of this chapter has been to reinterpret emergence. Against the standard tendency to understand emergence through the language of additional causal power, the chapter has argued that what is genuinely at stake is better captured as structural novelty. The problem of emergence is not first of all whether a higher level introduces a new producer into the world, but whether a system comes to exhibit a new stable form of organization, new structural boundaries, and new patterns of permissibility that cannot be adequately understood through a mere itemized recitation of lower-level processes.

Three main conclusions follow.

First, the standard dilemma between strong and weak emergence rests on a shared presupposition: both sides assume that genuine novelty must in some sense be causal novelty. Strong emergence preserves novelty by positing stronger ontological or causal commitments; weak emergence preserves continuity by attenuating novelty into epistemic or methodological terms. The difficulty with this opposition is not only that each side pays too high a price, but that both begin from the same miswritten question.

Second, once novelty is no longer presupposed to be causal novelty, a third path becomes available. A higher-level phenomenon may be genuinely new because it marks a new structural fact: a new stable organization, a new regime of constraint, a new space of permissible pattern, or a new way in which relevant and irrelevant variation are sorted. In that sense, novelty can be real without requiring additional causal powers. Structural novelty preserves dependence on lower-level

realization while resisting the reduction of higher-level reality to mere descriptive convenience.

Third, the explanatory standing of higher-level phenomena can be grounded in structural indispensability rather than causal addition. Higher-level accounts are legitimate when they identify stable organizational forms, constraint-structures, and patterns of permissibility that matter for understanding the system and cannot be omitted without explanatory loss. Their legitimacy lies not in competing with lower-level processes to produce outcomes, but in answering different explanatory questions about stability, organization, and open form.

If these conclusions are compressed into a single judgment, it would be this:

Emergence should be understood not as the appearance of new causal powers, but as the appearance of new stable structural organizations within a realized world.

This reformulation does not eliminate the higher level, and it does not mystify it. It neither collapses emergence into lower-level description nor inflates it into a new ontological inventory. It instead relocates the problem of emergence from the axis of extra causation to the axis of structural organization. Once that relocation is made, many of the supposed metaphysical pressures surrounding emergence lose much of their force.

This also prepares the way for the next chapter. If higher-level phenomena can be real and explanatorily significant without possessing independent causal power, then the broader question is no longer whether higher-level explanation competes with lower-level causation, but how higher-level explanation itself should be understood. Chapter 5 turns directly to that issue.

Chapter 5

The Legitimacy of Higher-Level Explanation: Beyond Exclusion and Mechanistic Monopoly

The legitimacy of higher-level explanation has long seemed difficult to secure because it is typically caught between two powerful pressures. On one side stands the exclusion problem: if lower-level physical causes are already sufficient, then higher-level explanation appears redundant unless it claims an additional causal contribution. On the other side stands the prestige of mechanistic explanation: if genuine explanation is taken to consist in showing how organized lower-level components generate an outcome, then higher-level explanation appears secondary unless it can be folded back into a mechanism. Between these two pressures, higher-level accounts are repeatedly forced into a narrow choice: either claim more causally than they plausibly can, or accept a diminished explanatory standing.

The argument of this book so far suggests that this choice is itself badly posed. If explanatory roles are differentiated, then the legitimacy of higher-level explanation need not depend on whether it competes for the same productive position as lower-level causation, nor on whether it can always be redescribed as a finer-grained mechanism. The real question is not whether higher-level explanation adds a second causal chain, but what kind of explanatory work it performs. If higher-level accounts identify stable patterns, structural boundaries, organizational relevance, and ranges of permissible variation that matter for understanding the system, then their legitimacy may rest on those roles rather than on causal rivalry.

This chapter develops that claim. It will argue that higher-level explanation is legitimate not because it escapes realization-dependence, and not because it introduces mysterious surplus powers, but because it answers questions that are

not exhausted by lower-level productive description alone. To show this, the chapter will first revisit the pressure generated by exclusion-style reasoning and explain why that pressure depends on a compressed picture of explanation. It will then examine the tendency to treat mechanism as the default explanatory ideal and show why that tendency, though methodologically powerful, cannot by itself settle the status of higher-level accounts. Finally, drawing on the framework of causal production, structural constraint, and permissibility, the chapter will articulate a more positive account of what higher-level explanation contributes and why that contribution is theoretically legitimate.

The aim is not to deny the importance of lower-level causation or mechanistic analysis. It is to show that once we move beyond the assumption that all explanation must terminate in one productive tribunal, higher-level explanation can be understood as a genuine and noncompetitive part of scientific understanding.

5.1 From Higher-Level Phenomena to Higher-Level Explanation

The legitimacy of higher-level explanation remains under pressure not because higher-level accounts are absent from scientific practice, nor because their practical usefulness is seriously in doubt. In many domains, higher-level descriptions, models, and explanatory frameworks are indispensable. The pressure persists, rather, because once the question turns from practical utility to theoretical legitimacy, higher-level explanation is repeatedly asked to justify itself before standards that were not designed for the kind of work it performs. It is judged as though its task must be the same as that of lower-level causal production, and once that assumption is in place, its standing becomes difficult to defend.

This pressure has a familiar structure. If the lower level is taken already to provide a causally sufficient account of how a result occurs, then the higher level

seems to face an immediate challenge: what is left for it to explain? If explanation is understood as the revelation of a productive pathway, then a higher-level account appears legitimate only if it either contributes additional causal efficacy or identifies some productive element missed below. If it does neither, it is easily demoted to heuristic summary, descriptive convenience, or merely provisional language. In this way, the problem of higher-level explanation is not framed as a question about explanatory plurality, but as a question about whether one more claimant can justify occupying the same already-filled explanatory slot.

That is why the pressure remains even when philosophers explicitly acknowledge that higher-level science is real and methodologically successful. One may readily grant that biology, psychology, economics, or systems theory employs higher-level patterns with great effectiveness, and still insist, at the decisive moment, that the “real” explanation must lie elsewhere, namely, at the level where the relevant outcome is ultimately produced. The higher level is then tolerated as useful, but not fully recognized as explanatorily authoritative. Its problem is not that it cannot speak; it is that it is repeatedly treated as speaking only by permission of a more fundamental productive account.

This is also why the pressure on higher-level explanation often takes two different but convergent forms. In the first form, familiar from exclusion-style reasoning, the question is whether higher-level explanation is redundant once lower-level causes are in place. In the second form, familiar from the prestige of mechanism, the question is whether higher-level explanation is incomplete until it can be anchored in an organized lower-level production story. The former pressures higher-level explanation through competition; the latter through subordination. But both presume that the ultimate standard of explanatory legitimacy lies in successful productive disclosure.

From the standpoint of this book, what keeps the issue under pressure is therefore not simply a dispute about levels, but a deeper contraction of

explanatory space. So long as explanation is treated as if it had one final tribunal, higher-level accounts will always appear vulnerable. If they claim causal autonomy, they risk inflation. If they refuse that claim, they risk demotion. If they appeal to stable patterns, structural boundaries, or organizational relevance, they are asked how those things are “really” implemented. If they answer by pointing to implementation, the higher-level role is then treated as dispensable. In each case, the problem is generated not only by the content of the higher-level account, but by the prior assumption that only one kind of explanatory success ultimately counts.

This helps explain why appeals to usefulness alone never fully settle the matter. To say that higher-level explanations are useful is true, but insufficient. Many things are useful that do not thereby secure strong theoretical standing. The real issue is not whether higher-level accounts help inquiry, but whether they answer genuine explanatory questions not already exhausted by lower-level production stories. Unless that point is made clearly, higher-level explanation will remain trapped in an unstable position: too important to ignore, yet always vulnerable to being described as secondary.

The pressure is intensified by the fact that higher-level explanation often abstracts away from many lower-level details. This abstraction can make it seem as though the higher level is merely selective in a pragmatic sense, choosing what is convenient for limited knowers while leaving the deeper reality untouched. But abstraction is not the same as arbitrariness. The issue is whether what is selected tracks a real structural regime, a stable pattern, or a space of relevant variation. If it does, then the higher-level account is not merely omitting detail; it is identifying a real explanatory form. The pressure persists because this difference between principled abstraction and mere convenience is repeatedly blurred.

A further source of pressure is the tendency to conflate two different questions: how something is realized, and why it matters explanatorily at the level under

discussion. The first question concerns the bearing relation between levels. The second concerns the function of explanation. Once these two questions are run together, realization-dependence is too quickly taken to imply explanatory dependence in a much stronger sense. But a higher-level account may depend on lower-level realization while still answering questions about pattern, organization, and stability that are not discharged simply by specifying that realization. If this distinction is missed, then higher-level explanation is bound to appear either redundant or illicit.

What follows from all this is that the pressure on higher-level explanation is not accidental. It is the predictable result of evaluating higher-level accounts within a compressed explanatory framework. The task of the present chapter is therefore not merely to defend the higher level by appeal to practice, nor to deny the importance of lower-level explanation. It is to show that once explanatory roles are properly differentiated, the demand that higher-level explanation justify itself as if it were competing for the same productive role begins to lose its grip.

The next step is to examine that pressure in its sharpest form by revisiting exclusion-style reasoning itself.

5.2 The Standard Schema of Exclusion

The exclusion problem exerts such lasting pressure on higher-level explanation because it appears to offer a very simple and very powerful challenge. If lower-level physical causes are already sufficient for an outcome, then what explanatory work can remain for a higher-level account to do? This challenge seems difficult to evade precisely because it presents itself not as hostility to higher-level explanation, but as a demand for theoretical economy. Why multiply explanatory claims where one sufficient lower-level account already exists? Why not say that once the productive story is complete, anything further is at best repetition or convenience?

5.3 The Conflation of Causal Sufficiency and Explanatory Sufficiency

The force of this question depends, however, on more than the principle of physical sufficiency. It depends on the move from productive sufficiency to explanatory sufficiency. Once that move is made, the space of higher-level explanation is narrowed from the outset. If the lower level is already taken to have explained the result simply in virtue of having produced it, then any higher-level account can seem legitimate only by doing one of three things: repeating the same explanation in different words, competing by claiming additional causal force, or withdrawing into a weaker pragmatic role. That is exactly why exclusion reasoning has such a powerful contracting effect. It does not merely raise a problem for one kind of causation; it reshapes the terms under which higher-level explanation is allowed to count at all.

This becomes especially clear when we separate two questions that exclusion-style reasoning often fuses together. The first question is: what brought about the outcome in this case? The second is: what explanation is needed to understand the system at the relevant level? These are not always the same question. A lower-level causal account may answer the first very well by specifying how the outcome was realized. But that does not automatically answer the second, because understanding a system at a given level may also require us to know what pattern is instantiated, what organization is stable, what boundaries structure the relevant regime, and what variations are significant or insignificant there. Once these tasks are recognized, the alleged exhaustiveness of the lower-level account becomes less obvious.

None of this means that exclusion reasoning is simply confused or unimportant. Its pressure points to something real: we do need an account of why higher-level explanation is not just loose talk layered on top of lower-level causation. The point is rather that exclusion-style reasoning asks the right question in the wrong form. It senses that higher-level explanation needs justification, but it

assumes too quickly that such justification must come in the form of additional productive contribution. As a result, it turns the issue into a contest over causal territory, when the more appropriate issue is whether different explanatory roles are being answered.

Once the problem is framed in that broader way, the apparent inevitability of the exclusion challenge weakens. A higher-level account need not justify itself by claiming that it produces the outcome over and above the lower-level account. It may justify itself by answering a different explanatory question: why this outcome belongs to a stable higher-level pattern, why certain lower-level variations do not matter at this level, why the system permits this form of organization, or why this regime of behavior remains coherent across multiple realizations. None of these questions is answered simply by saying that the lower-level physical causes were sufficient to bring about the token outcome.

This also helps explain why higher-level explanation can remain indispensable even when no one denies lower-level causation. The issue is not whether there was enough lower-level causation for the event to occur. The issue is whether causal sufficiency at one level settles every explanatory question relevant to the phenomenon. Exclusion reasoning tacitly assumes that it does. But if explanation is multi-role, then lower-level causal sufficiency no longer functions as a universal veto. It secures one kind of answer, not every kind of answer.

At the same time, the multi-role framework avoids the opposite mistake of treating higher-level explanation as exempt from discipline. Not every higher-level description earns explanatory standing merely by being higher-level. The question remains whether it identifies a stable structural fact, a pattern of permissibility, or a form of organizational relevance that genuinely matters for understanding. If it does not, then exclusionary suspicion may be justified. But if it does, then the fact that lower-level causation is sufficient for realization does not by itself remove its

legitimacy. What matters is not whether the higher level adds another producer, but whether it answers another question.

In this light, the exclusion problem is best understood not as proving the illegitimacy of higher-level explanation, but as revealing the cost of evaluating all explanations by a single productive standard. Once that standard is relaxed, the higher level need not be defended as a rival cause. It can instead be defended as a distinct explanatory role within a layered account of the same reality. The lower level may explain how the outcome occurred; the higher level may explain what kind of organized phenomenon that outcome instantiates and why that organization matters.

The next source of pressure to examine is different in form but similar in effect: the tendency to treat mechanistic explanation as the default or privileged model of what explanation really is.

5.4 Mechanistic Explanation and the Expansion of Its Paradigm

If exclusion pressures higher-level explanation by treating lower-level causation as already sufficient, the prestige of mechanistic explanation pressures it in a different way. Here the higher level is not so much challenged as redundant, but subordinated as incomplete. The idea is not necessarily that higher-level accounts are false, but that they remain explanatorily unfinished until they can be anchored in an account of how organized lower-level parts and activities generate the phenomenon in question. Mechanistic explanation thus appears less hostile than exclusion, but it can still impose a deep narrowing effect: it encourages the thought that genuine explanation reaches completion only at the point where production is displayed through mechanism.

This pressure is powerful because mechanistic explanation genuinely captures something central in scientific practice. In many areas of inquiry, to explain a

phenomenon is indeed to show how components, activities, and their organization produce it. Mechanistic analysis often provides the detail, manipulability, and structural articulation that make an explanation satisfying. For that reason, the present argument cannot and does not proceed by minimizing the importance of mechanism. The issue is not whether mechanistic explanation matters, but whether it should be treated as the default model against which all explanation must ultimately be measured.

Once that default status is assumed, higher-level accounts are placed in an unstable position. They may be tolerated as heuristically useful, as ways of guiding inquiry, or as temporarily serviceable summaries. But their standing remains secondary so long as the real explanatory endpoint is held to lie in the disclosure of the underlying mechanism. On this picture, a higher-level explanation is always waiting either to be vindicated by mechanistic grounding or to be superseded by it. Even where the higher level is not rejected outright, it is not allowed to count as fully explanatory in its own right.

The difficulty with this picture is that it once again conflates different explanatory tasks. A mechanism may explain how a phenomenon is produced, but it does not automatically explain why a certain higher-level pattern is the relevant unit of understanding, why some variations within the mechanism matter and others do not, why a stable organizational regime persists across multiple realizations, or why certain forms of system-level behavior are open while others are excluded. These are not anti-mechanistic questions. They are questions to which mechanism alone may not suffice, because they concern structural relevance and permissible organization rather than merely productive unfolding.

This is especially important when the same phenomenon can be realized through different lower-level pathways. In such cases, the mechanistic story may vary while the higher-level explanatory target remains stable. If one insists that explanation is fully complete only where the mechanism is completely specified,

then one risks missing what the higher-level account is actually tracking: not one unique lower-level route, but a pattern that remains explanatorily significant across several routes. The higher level is then not standing in for an unknown mechanism; it is identifying a stable organizational form that no single mechanism-description exhausts.

A further difficulty is that the mechanistic ideal can absorb structural and organizational considerations without granting them independent standing. Boundary conditions, stability relations, and regime-level constraints may all be redescribed as parts of the mechanistic setup. In some cases this is entirely appropriate. But if this move becomes universal, then every form of explanation is forced into the role of either contributing directly to production or waiting to be translated into those terms. Structural constraint and permissibility are thus not denied; they are domesticated. They appear only as secondary features of mechanism rather than as explanatory roles in their own right.

The present argument therefore does not oppose mechanism as such. It opposes only the monopoly of mechanistic explanation. Mechanism is one crucial explanatory role, especially where the problem is how a result is produced through organized activity. But it is not the sole measure of explanatory legitimacy. Higher-level explanations may be legitimate not because they reveal a hidden mechanism, but because they clarify a stable pattern, a structural boundary, or a space of relevant variation that matters for understanding the system. In such cases, mechanism may remain indispensable at one level, yet incomplete at another.

Once this is acknowledged, the relation between higher-level explanation and mechanistic explanation can be redescribed more productively. They need not stand in a relation of rivalry, where one replaces the other, nor in a relation of strict hierarchy, where one is explanatory only by permission of the other. Rather, they may answer different questions about the same phenomenon. Mechanistic explanation may tell us how the phenomenon is produced. Higher-level

explanation may tell us what pattern is being produced, why it counts as the same pattern across varying implementations, and what organizational conditions define its explanatory relevance. The point is not to separate them absolutely, but to prevent one from absorbing the other's task.

This shift matters because it opens the way for a more positive account of higher-level legitimacy. If higher-level explanation need not justify itself either as an extra cause or as an incomplete mechanism, then its standing can be articulated in its own terms. The next step is to say what those terms are, and how a higher-level account earns explanatory authority without competing for the same productive role as the lower level.

5.5 The Noncompetitive Standing of Higher-Level Explanation

Up to this point, the chapter has shown that higher-level explanation has long appeared suspect not simply because some philosophers have an excessive preference for lower-level physics, but because explanatory space has been systematically organized as a single productive structure. Within that structure, once a lower-level causal account is taken to be sufficient, a higher-level account can occupy only one of three unfavorable positions: it is either redundant, because the lower level has already done the work; competitive, because it seems to claim additional efficacy with respect to the same result; or pragmatic, because it is merely a convenient expression for finite cognizers. If this picture is accepted, higher-level explanation will of course struggle to secure stable standing.

The central claim of this section is that this picture is so oppressive not because higher-level explanation really is always competing for the same position, but because it has been wrongly understood in that way. The legitimacy of higher-level explanation should not rest on its ability to compete with lower-level explanation for the same causal position. It should rest instead on the fact that it performs a different explanatory task. Higher-level explanation is nonredundant not because it

adds another hidden chain outside the lower-level causal chain, but because it answers questions that lower-level productive accounts do not automatically answer.

To understand this point, we must distinguish more clearly between explaining the same phenomenon and explaining the same phenomenon from the same explanatory position. Higher-level and lower-level explanations often do concern the same phenomenon, the same system, or even the same result. But that does not mean that they are competing for the same explanatory task. A lower-level explanation may mainly answer how a given result is brought about through physical processes, component activities, or mechanistic chains. A higher-level explanation may mainly answer why that result constitutes a stable pattern at a certain level, why some lower-level differences are irrelevant there, why the system is open only to certain organizational forms, or why some forms of behavior count as relevant at that level.

If this point is not seen, higher-level explanation will continue to be misread. Once we assume that there is only one genuine explanatory task, namely to show how a result is generated, every additional account can only appear to be a repetitive supplement to the same task. Higher-level explanation will then look either superfluous or as though it were smuggling in extra causal power. But if explanatory tasks are themselves multi-role, this exclusionary conclusion no longer follows. A lower-level causal account may be entirely sufficient with respect to the productive task while a higher-level account remains indispensable with respect to the structural task.

From this perspective, the noncompetitive standing of higher-level explanation is not a retreating claim, as though higher-level explanation were forced to withdraw to the margins because it cannot enter the true core of explanation. On the contrary, it means that higher-level explanation has its own legitimate domain of task. It is not noncompetitive because it is unable to compete; it is

noncompetitive because it should never have been placed in that competition in the first place. An account of organizational stability, level-relevance, and pattern-boundaries does not lose explanatory content merely because it does not additionally push some result.

A further point must also be emphasized. The noncompetitive status of higher-level explanation does not mean that it is unrelated to lower-level explanation, nor that it can stand entirely apart from the lower level. This book does not advocate a sealed pluralism of levels. Higher-level explanation usually depends on lower-level realization, and it often cooperates closely with lower-level accounts. The issue is simply that dependence and competition are not the same thing. An explanation may depend on facts of realization at another level while still answering different questions at its own level.

This also shows why exclusion arguments and mechanistic explanation fail at this point. The problem with exclusion is that it moves directly from 'the lower level is sufficient to bring about the result' to 'there is no explanatory space left for the higher level.' The problem with mechanistic explanation is that it moves from 'the generative mechanism has been displayed in sufficient detail' to 'mature explanation is complete.' What both overlook is that productive sufficiency for a result is not the same as the exhaustion of all explanatory tasks concerning that result.

In concrete scientific practice, this noncompetitive standing is not at all unfamiliar. Scientists often do not treat higher-level accounts as repetitions of lower-level causal chains. On the contrary, they frequently identify key facts in a system through functions, organization, behavioral patterns, and constraint-structures at higher levels. Such accounts are not used merely because the lower-level processes have not yet been found, nor do they automatically become invalid once those lower-level processes are discovered. They continue to operate because their concern is not how a single event occurs, but what counts as relevant

change at a given level, what counts as stable behavior, and what forms the system permits.

Accordingly, the legitimacy of higher-level explanation does not lie in adding another layer of the same productive content to a lower-level causal account. It lies instead in making visible structural facts that would not naturally appear within a lower-level productive narrative. It characterizes units of organization, boundaries of variation, constancy of pattern, and the selective structure of level-relevance. Without such an account, even if we knew how events occur at the lower level, we might still not know why the pattern at this level is theoretically important, why certain differences are ignored, why certain organizational forms can be maintained, or why the system can present itself only in these forms rather than others.

The conclusion of this section can therefore be stated quite clearly: higher-level explanation is legitimate not because it supplements lower-level explanation with causes that have not yet been discovered, but because it performs an explanatory role distinct from lower-level productive accounts. It may concern the same system as lower-level explanation, yet it need not compete for the same position. Only on this basis can we avoid both mystifying higher-level explanation as an additional causal power and demoting it to mere pragmatic convenience.

The next section builds on this point by explaining positively what higher-level explanation is doing, that is, through which structural functions it obtains its explanatory credentials.

5.6 The Structural Functions of Higher-Level Explanation

If the legitimacy of higher-level explanation does not rest on additional causal power, and does not rest on competing with lower-level accounts for the same productive position, then a more direct question immediately follows: what exactly is higher-level explanation doing? If that question is not answered positively,

'noncompetitive' will too easily be misunderstood as a merely defensive posture, as though higher-level explanation were assigned to a weaker place only because it cannot enter the true causal core. The present section argues the opposite. Higher-level explanation is legitimate not because it does less, but because it answers different questions. It undertakes a group of explanatory tasks centered on organization, stability, level-relevance, and structural boundaries.

The first structural function of higher-level explanation is the identification of organizational units. Lower-level description can often provide very detailed local chains of events, but it does not automatically tell us what should count, at a given level of inquiry, as a functional unit, what constitutes a relevant system, or which components should be treated as a jointly organized whole. Higher-level explanation does not merely rename lower-level elements here. It identifies which mode of organization has theoretical significance for the problem at hand.

The second function is the sorting of level-relevant differences. A system may contain an enormous number of differences at the lower level, but not all of them matter at the higher level. One of the key tasks of higher-level explanation is to determine which variations alter the pattern at that level and which are absorbed or ignored by that level's organization. Explanation is not simply a matter of collecting every fact; it is a matter of identifying which facts are relevant to the present question. Higher-level explanation establishes a level's own explanatory geometry: what counts as the same, what counts as different, which changes are noise, and which count as structural turning points.

The third function is the characterization of stable patterns. As the previous chapter argued, higher-level phenomena are worth treating as real not because they introduce additional causal power, but because they form stable patterns. Higher-level explanation does not primarily show how one particular event was produced. It explains why a certain kind of pattern can be maintained in the system. It is concerned not with one token realization, but with constancy across

change: why a behavioral form is preserved across many realizational conditions, why some structural relation continues to constrain later system-states once formed, and why some organizational arrangement repeatedly appears as a theoretically identifiable fact at that level.

The fourth function is the characterization of structural boundaries. Lower-level causal accounts can show how things happen, but they do not always automatically reveal which states, behaviors, or organizational forms are open at a given level and which do not belong to that level's space of possibility at all. Higher-level explanation is partly indispensable because it can characterize such boundaries. It tells us what forms of change the system allows at this level, which local perturbations do not destroy the overall pattern, and which changes mean that the structure at that level has broken down.

The fifth function is the grasp of cross-realization invariance. Many higher-level patterns have explanatory value precisely because they are not locked to a single lower-level path, but can be borne by multiple realizations. Higher-level explanation does not ignore lower-level realization here; it explains how different realizations can still be grouped under one and the same pattern at the higher level. It identifies the structural conditions under which distinct local compositions can still realize the same higher-level organizational form.

Taken together, these functions display the real structural content of higher-level explanation. It is not an 'as-yet-incomplete mechanistic account,' not a compressed summary of lower-level description, and not a substitute language used only when causal explanation fails. It answers a set of questions that lower-level productive accounts do not automatically answer: what defines the object-boundary at this level, which differences matter here, which patterns can be stably maintained, which forms of change belong to this level's open space, and how different realizations are unified at the same level.

It must also be emphasized that the structural functions of higher-level explanation are not disconnected from causal production. A higher-level pattern of course has to be realized in concrete processes, and a structural boundary must be connected with lower-level mechanisms. The point is not that the two are unrelated, but that productive explanation does not automatically complete structural explanation. Even if a mechanism is described in great detail, we may still not know why it counts as a functional unit at a certain level, why some micro-differences are irrelevant there, or why one organizational form remains constant across multiple realizations.

This also lets us see more clearly why higher-level explanation is not an optional 'second language.' A merely secondary language could in principle be wholly replaced by a more basic one without any loss of content. That is exactly what this book rejects. Higher-level explanation is not adopted because we are too lazy to inspect lower-level detail; it is adopted because it corresponds to a different type of content. Even if every lower-level detail were in hand, if we no longer distinguished functional units, structural boundaries, or level-relevant differences, we would not have a more complete explanation. We would likely lose what is theoretically most important.

The conclusion of this section can therefore be stated in a compact form: the legitimacy of higher-level explanation rests on its structural functions. By identifying organizational units, sorting level-relevant differences, characterizing stable patterns, drawing structural boundaries, and grasping cross-realization invariance, it provides an irreplaceable form of understanding of the system. This point also prepares the next section on downward causation, because once higher-level explanation is positively understood as structural explanation, we can ask whether many cases traditionally described as 'downward causation' are better understood as the operation of structural constraint and level-boundaries.

5.7 Rethinking Downward Causation

In discussions of higher-level explanation, few notions are at once as attractive and as dangerous as downward causation. On the one hand, the phrase seems to capture an important phenomenon: higher-level forms of organization, systemic structure, or global patterns do indeed seem to 'feed back' upon local processes. An institutional arrangement shapes individual behavior; the organized state of a living whole constrains cellular activity; a functional requirement filters the realizational paths available below. If we reject this language altogether, we may seem to lose a way of expressing the reality of higher-level efficacy. On the other hand, downward causation is very easily heard as meaning that the higher level adds an extra pushing force outside the lower-level causal chain. Once understood that way, it immediately falls back into the dilemma this book has repeatedly rejected.

For that reason, the basic judgment of this section is that the long-running controversy over downward causation arises not mainly because the phenomenon it tries to describe is wholly fictitious, but because it runs together two things that should be separated. One is the fact that higher-level structure, global organization, and level-patterns can indeed alter the feasible range, stability-conditions, and path-selection of lower-level processes. The other is the claim that they therefore constitute additional productive causes of lower-level events. The former may well be real; the latter need not be accepted.

Seen through the book's multi-role framework, this is not hard to understand. Causal production, structural constraint, and permissibility must be distinguished. Downward causation feels both necessary and dangerous precisely because it tries to express, in the language of causal production, something that is often closer to structural constraint. A higher-level pattern may indeed change the boundaries within which lower-level processes unfold: making some micro-changes compatible with the whole function and others incompatible, opening some local

paths while excluding others. But this does not automatically mean that the higher level has entered the lower-level causal chain as an independent additional cause.

From this perspective, what is often called downward causation is more accurately described as higher-level organization reshaping the constraint-environment and the space of permissibility within which lower-level processes proceed. The term 'causation' is what makes the idea so prone to misunderstanding, because it immediately suggests a top-down transmission of force, pushing, or difference-making, as though the whole had opened a second causal channel outside its parts. Yet in many real contexts, what researchers want to express is something more restrained: the whole's organization makes local processes meaningful only within certain boundaries, the system-level stable form filters which lower-level paths are realizable, and the functional requirements of the level render some local differences irrelevant while making others decisive.

The advantage of this rewrite is that it preserves the reality of higher-level efficacy while avoiding the theoretical burden of additional causal power. It allows us to say that the whole really does affect the parts, but that this effect is first of all structural limitation rather than productive extra force. The system-level does not open a new chain outside lower-level causal closure; it determines which lower-level changes are maintainable, coordinable, and integrable within that system. Without such determination, lower-level processes would still occur, but they would not necessarily compose the same whole, sustain the same function, or realize the same level-pattern.

This rewrite does not deny real consequences at the higher level. On the contrary, it is proposed precisely to describe those consequences more accurately. If we insist on speaking in the language of additional productive causation, higher-level efficacy becomes more likely to look suspicious or mysterious, because it seems to require a revision of physical closure. But if we shift the focus to structural constraint, boundary-conditions, and spaces of permissibility, the efficacy of the

higher level becomes both real and intelligible. It asks us not to posit a new force, but to acknowledge that forms of organization can reshape the relevance-structure of local processes.

This matters especially for the status of higher-level explanation itself. Many theorists reach for the language of downward causation because they are reluctant to surrender the efficacy of the higher level, yet lack another vocabulary for expressing it. They are then forced to oscillate between two unsatisfying options: either say that the higher level really exerts a downward force and accept a metaphysical burden, or abandon talk of downward efficacy altogether and seem to weaken the substance of higher-level explanation. The present proposal offers a third path. The higher level can be fully efficacious, but its efficacy is mainly structural rather than additionally productive. What it changes is not the number of causes in the world, but the boundaries and organization within which causal processes unfold.

Caution is still needed. Not every use of the term downward causation can simply be translated into structural constraint; some writers do intend the stronger thesis of an independent higher-level causal contribution. This book does not deny that such a stronger thesis can be formulated. It claims only that, in the great majority of contexts relevant to the legitimacy of higher-level explanation, we do not need to go that far. To preserve higher-level explanation, we do not need to win for the higher level an additional productive causal identity. Recognition of structural downward constraint is already enough to explain why holistic organization genuinely shapes partial processes.

The conclusion of this section can therefore be put in a straightforward way: if downward causation is understood as an additional productive force exerted by the higher level upon lower-level events, then it pushes higher-level explanation back into the strong-emergence picture this book has consistently rejected. But if it is rewritten as the shaping of lower-level processes by higher-level organization

through structural constraint and permissibility, then the phenomenon it was trying to capture can not only be preserved, but expressed more clearly.

The higher level's 'downward effect' is not a surplus cause. It is the fact that lower-level processes can realize the system only within certain organizational boundaries.

With exclusion pressure, mechanistic expansion, and the miswriting of downward causation all reinterpreted in this way, we are now in a position to state the chapter's main conclusion more directly: the legitimacy of higher-level explanation rests not on extra causal power, but on the distinct structural tasks it performs.

The next section gathers those tasks and pressures into a compact summary.

5.8 Responses to Three Typical Objections

Objection 1: Higher-Level Explanation Is Mere Compressed Description

If higher-level explanation were merely compressed description, then removing it would cost us no understanding of real structure. But in many cases, removing higher-level explanation means losing information about stable patterns, realizational invariance, and organizational boundaries. Higher-level explanation is therefore not an arbitrary compression. It is a grasp of real structure at a certain scale of organization.

Objection 2: Structural Constraint Is Ultimately Still Nothing More Than Lower-Level Causal Relation

Structural constraint certainly depends on lower-level realization, but dependence is not the same as explanatory reduction. Lower-level causal chains explain how some events occur; structural constraint explains which paths are reachable or unreachable within the system's organization. The two can both be true while still answering different questions.

Objection 3: This Is Only Pragmatic Pluralism

If the value of higher-level explanation came only from human cognitive convenience, then different higher-level partitions would in principle have no real superiority or inferiority. But in actual scientific practice, higher-level partitions are constrained by stability, realizability, counterfactual support, and cross-realization preservation. Higher-level explanation is therefore not a purely pragmatic choice. It is constrained by the structure of reality itself.

These replies do not show that every higher-level description is automatically legitimate. They show only that the most familiar objections lose their force once explanatory roles are differentiated and structural tasks are taken seriously.

What matters is not whether the higher level adds a second productive chain, but whether it captures a real organizational fact, a stable pattern, or a structurally constrained space of variation.

5.9 Chapter Summary: From the Exclusion Structure to the Differentiation of Explanatory Roles

This chapter began from the question left by the previous chapter: even if the reality of higher-level phenomena no longer needs to rest on additional causal power, why is higher-level explanation itself still legitimate? The question is harder than the question of whether higher-level phenomena exist, because it directly concerns the allocation of explanatory space. A level of phenomena may be fully real, yet still have its explanations treated as redundant, repetitive, or merely pragmatic so long as explanation is presupposed to be a single productive task.

The chapter first reconstructed the standard schema of exclusion and argued that its enduring pressure comes not only from physical closure and anti-overdetermination intuitions, but from a further slide: lower-level causal sufficiency is too quickly turned into a condition excluding higher-level explanation. That is the crucial conflation between causal sufficiency and explanatory sufficiency. Even if

lower-level physical processes are sufficient to produce some result, they do not thereby automatically exhaust every explanatory task surrounding that result.

The chapter then turned to mechanistic explanation and its paradigmatic expansion. Mechanistic explanation exerts pressure on higher-level accounts not because it openly denies their importance, but because it is so successful that it easily becomes the default image of mature explanation. Against both exclusion and mechanistic monopoly, the core response developed here is that higher-level explanation is legitimate not because it competes for the same causal position, but because it undertakes different explanatory roles. It identifies organizational units, sorts level-relevant differences, characterizes stable patterns, draws structural boundaries, and grasps cross-realization invariance.

On that basis, the chapter also reinterpreted downward causation. Many cases traditionally described in that idiom are better understood not as the higher level exerting an additional productive force upon lower-level events, but as higher-level organization shaping lower-level processes through structural constraint and permissibility. Once framed that way, higher-level explanation no longer needs either to declare war on lower-level explanation or to wait to be replaced by it. The differentiation of explanatory roles shows its force precisely here: once productive accounts no longer monopolize explanatory legitimacy, higher-level explanation no longer has to oscillate between competition, repetition, and weak convenience.

This also carries the discussion naturally into Chapter 6. Even if higher-level explanation, higher-level phenomena, and structural forms of account are all legitimate, a further challenge remains: if such frameworks do not always directly yield new empirical predictions, why should they still be taken seriously within philosophy of science? Chapter 6 turns to that broader methodological question by asking how non-predictive frameworks can still possess genuine scientific standing.

Chapter 6

The Scientific Standing of Non-Predictive Frameworks

If the arguments of the previous chapters are sound, then we already have reason to reject a fairly common contraction: it is not only productive causal accounts that count as genuine explanation; higher-level explanation, structural constraint, and permissibility can all bear independent and legitimate theoretical tasks. Even so, at this point the book still encounters a further challenge. Even if these forms of account are not empty, are they still merely philosophically interesting peripheral language, lacking strict scientific standing? In other words, if a theoretical framework does not directly yield new testable predictions, why should it be regarded as serious cognitive work rather than as temporary conceptual housekeeping, a methodological expedient, or an abstract posture waiting to be filled with empirical content?

This challenge has force because the modern scientific picture has long granted prediction an almost natural privileged status. A theory that accurately predicts phenomena not yet observed is easily treated as weighty and substantial; by contrast, a framework chiefly concerned with clarifying concepts, constraining explanatory space, reorganizing problem structure, or showing that certain theoretical demands are themselves misplaced is easily suspected of lacking hardness. Thus the structural and higher-level accounts defended in the previous chapters may perhaps secure a place at the level of explanation, only to be pushed back to the margins at the level of methodological appraisal. They may be said to be enlightening without being productive, philosophically significant without being scientifically substantial.

The central claim of this chapter is that this suspicion rests on another unnecessary contraction. The problem is not that prediction is unimportant, but that prediction, causation, and explanation are often bundled together as though they were one and the same kind of achievement. It is as though a theory counts as having genuinely contributed only when it yields new eventive outcomes, while all work concerning structure, constraint, level-boundaries, explanatory roles, or ontological restraint is treated as mere preparation for “real theory.” The chapter argues that this methodological picture is too narrow. A framework may fail to aim primarily at direct prediction and yet still perform irreplaceable functions in scientific understanding, such as clarifying the form of a problem, excluding illegitimate explanatory demands, constraining the space of possible theories, unifying apparently dispersed phenomena, or providing stricter conceptual boundaries for empirical research. Without such work, prediction itself often cannot be properly situated or interpreted.

The task of this chapter, then, is not to oppose prediction, but to limit prediction-centrism. It will address, in sequence, several questions: why prediction has acquired such a high evaluative status in the modern image of science; why that status is so easily conflated with causation and explanation; what non-predictive frameworks are actually doing, rather than merely “not making predictions”; by what standards a framework should be assessed if its chief function is structural rather than directly predictive; and what methodological standing the multi-role view of explanation developed in the earlier chapters should occupy in this wider context. Only once these questions are clarified can structural explanation and higher-level accounts avoid being preserved at the explanatory level only to lose standing again at the methodological level.

6.1 Why Non-Predictive Frameworks Are Always Suspected

Non-predictive frameworks often find themselves at a disadvantage in academic evaluation not merely because some scholars happen to prefer empiricism or positivism, but because a deeper scientific-cultural intuition has long been at work: a truly powerful theory ought to tell us what will happen next. If a framework cannot generate new empirical expectations, cannot display its grip in cases not yet observed, or cannot lock in some future result of the world in advance, it is easily judged not hard enough. Against this evaluative background, any work concerned chiefly with conceptual clarification, structural articulation, the drawing of explanatory boundaries, or ontological organization naturally comes to seem second-rate, as though it were at most setting the stage for genuine empirical theory rather than forming an important part of theoretical work itself.

Part of the reason this suspicion is so persistent is that prediction really does present a very direct image of success. Prediction seems to display the value of a theory in its clearest form: if a theory can say, before the result appears, what will happen, and reality later conforms to it, then the theory seems not merely to have reorganized already-given facts, but to have genuinely grasped some structure of the world. For that reason, prediction is easily taken to be the purest manifestation of theoretical power. By comparison, frameworks that chiefly redescribe a problem, limit acceptable explanation, or show that certain forms of question are themselves mistaken appear less dramatic and less able, at first glance, to exhibit their accomplishments.

And it is precisely here that non-predictive frameworks begin to look suspect. They do not usually aim directly at “telling you what new phenomenon will occur.” More often they undertake another sort of task: telling you which questions were wrongly posed, which accounts do not in fact compete with one another, which ontological expansions are unnecessary, which distinctions of level cannot be skipped in theory, or which structural relations must first be clarified lest empirical

results themselves be misread. Work of this sort seldom produces an immediately impressive “new prediction,” and so it is easily treated as merely theoretical hygiene on the periphery rather than as knowledge-growth itself.

Yet this suspicion is worth examining precisely because it presupposes an overly single image of theoretical achievement. It assumes that if a theory is to be valuable, it must display that value first of all by prediction; and that if a framework operates chiefly through organization, constraint, clarification, and limitation, then its contribution is thereby weaker. The problem is that this picture never adequately explains why prediction should serve as the unified end-point for all theoretical work. Still less does it explain why a framework whose main task is to answer questions such as what counts as a legitimate explanatory demand, what theoretical space is open, or what form of problem avoids generating pseudo-problems should not count as a substantive contribution.

Non-predictive frameworks are also often suspected because they are easily misread as “unfalsifiable.” If a theory offers no new prediction, critics may worry that it is thereby evading empirical test, or that it can always save itself afterward by redescribing the language of the issue. In other words, non-predictiveness is often directly associated with methodological looseness: as though, so long as one does not incur predictive risk, one may indefinitely extend explanatory language without facing genuine theoretical constraint. This concern is not wholly without basis, since there have indeed been historical frameworks that expanded themselves through grand terminology while offering almost no testable boundaries.

But it is too quick to infer from this that every non-predictive framework is methodologically suspect. A framework may fail to produce new predictions directly and yet still be constrained in other ways; it may fail to take eventive forecasting as its primary aim and yet still be unable to say just anything it pleases. It may be rigorously evaluated by asking whether it has clarified a previously disordered problem structure, reduced unnecessary ontological burden, unified

formerly disconnected phenomena, identified which explanatory demands are themselves misplaced, or drawn stricter conceptual boundaries for empirical inquiry. If it performs these tasks well enough, that too constitutes theoretical advance, even if the advance does not first appear in predictive form.

From the standpoint of this book, the persistent suspicion directed at non-predictive frameworks resonates deeply with the explanatory monism criticized in the previous chapters. Explanatory monism compresses genuine explanation into causal production; prediction-centrism further compresses genuine theoretical achievement into new prediction. What both share in the background is the picture that only those accounts that directly take over the occurrence of events count as the hardest and most central theoretical achievements. Under that picture, work concerning structure, constraint, permissibility, level-boundaries, and ontological restraint is easily demoted en bloc. Even when it has content at the explanatory level, it comes to seem methodologically not quite scientific enough.

A provisional conclusion may therefore be drawn here: non-predictive frameworks are always suspected not because they naturally lack theoretical value, but because a too narrow measure of achievement has long prevailed within the modern image of science. So long as theoretical value is presupposed to show itself first as new prediction, every framework whose main work lies in clarification, limitation, organization, and reconstruction will look like second-order labor. To secure a proper methodological place for such frameworks, the next step is to examine more carefully why prediction has acquired so special a status, and at what points that status has been wrongly generalized.

6.2 Why Prediction Acquires Methodological Privilege

If non-predictive frameworks so easily arouse suspicion, the next question is why prediction enjoys so stable a priority in modern scientific evaluation. This status was not formed by accident, nor is it merely the residue of some narrow

empiricism. On the contrary, it possesses a strong historical rationality and methodological appeal. Science has long treated prediction as a conspicuous sign of theoretical power because prediction seems most directly to show that a theory is not merely reorganizing facts after the event, but genuinely grasping some generalizable structure in the world. If a theory can indicate, before a result appears, how that result will appear, then it seems to be not just describing the past, but laying hold of the order of reality itself.

This point explains, first, why prediction is so easily associated with “real theory.” Retrospective explanation always leaves room for suspicion, because one can worry that it is nothing more than a re-description of already-known results in the form of a tidier-looking story. Prediction, by contrast, places theory under a stricter kind of risk: if the prediction fails, the credibility of the theory is visibly damaged; if it succeeds, the theory seems to have displayed its power in a more compelling way. For that reason, prediction is often regarded as a harder kind of theoretical performance than retrospective organization. Its appeal lies not only in the surprise or sense of accomplishment it can produce, but in the fact that it seems naturally tied to testability, fallibility, and a genuine grasp of the world’s structure.

Second, prediction acquires privilege because it is closely linked to the ideals of control and intervention. Modern science seeks not only to understand the world, but often to act effectively within it. If a theory can predict how a system will behave under different conditions, then it has not only cognitive value but practical value. It can help us design experiments, avoid risks, improve technologies, optimize policy, or achieve particular outcomes by manipulating variables. In this setting prediction naturally appears not merely as an additional virtue of theory, but almost as the best evidence that theory and effective practice have been successfully joined. When a theory cannot support prediction, people are easily led to doubt whether it is genuinely close enough to the way the world works.

Third, prediction appears methodologically privileged because it is often taken to be an outward sign of theoretical unity. A theory that explains only already-known isolated phenomena may still seem local and piecemeal; but a theory that can derive previously unobserved results from one and the same set of principles seems to exhibit greater internal organization. Prediction, in other words, is not merely a matter of “getting something right in advance”; it is also treated as a test of whether a theory rises above local patchwork and forms a more general structure. This is why major theoretical breakthroughs are so often narrated in the history of science as episodes of predictive success: prediction appears as the sign that theory has moved from the reorganization of experience to the disclosure of structure.

Precisely because these reasons are genuinely attractive, prediction’s privileged status cannot simply be dismissed. This book does not claim that prediction is some misleading idol, nor that science ought to abandon its concern with predictive strength. The problem lies elsewhere: the fact that prediction is extremely important for some tasks does not automatically show that it should serve as the unified evaluative standard for all theoretical work. A standard may be most appropriate in one class of task and yet begin, once generalized, to obscure other equally important theoretical functions. The problem with prediction is not that it matters, but that it too easily slides from being one important form of achievement into being the only form of genuine achievement.

This slide is especially easy in the context of theories of explanation and philosophy of science. Once prediction is treated as the primary measure of theoretical value, it becomes natural to assume further that any work not directly increasing predictive power is merely preparatory to real theory. Conceptual clarification, the differentiation of explanatory roles, the identification of boundaries between levels, ontological restraint, and the analysis of structural constraint all then come to seem secondary or even pre-theoretical. They may be

helpful, but they no longer appear to count as positive advances in knowledge. Prediction's privilege thus ceases to be only a criterion for whether a theory can handle future results and becomes instead a global criterion for whether theoretical work is worth doing at all.

That generalization is especially worth resisting because it quietly changes the very meaning of theoretical progress. If only frameworks that produce new predictions count as genuine progress, then all work aimed at limiting mistaken questions, excluding pseudo-explanatory demands, reducing ontological inflation, unifying relations among levels, or reconstructing explanatory space begins to look like peripheral circling. Yet many important advances in theory consist precisely in our no longer asking an old question in its old form, in our recognizing that certain theoretical demands were themselves wrong, or in our seeing that some dispute lacked not more facts but clearer structural division. Such advances do not necessarily appear first as predictions, but they may still be foundational for scientific understanding.

This already prepares the next issue. If prediction has privileged status for good reasons, but those reasons do not justify turning it into the universal measure of all theory, then we must ask how prediction became so tightly bound up with causation and explanation themselves. Only once that bundle is loosened can we see more clearly how a framework may contribute rigorously and substantially without first yielding new predictions.

6.3 The Conflation of Prediction, Causation, and Explanation

If prediction acquires methodological privilege not only because it is important in some contexts, but because it is further generalized into a comprehensive standard of evaluation, then the next question is how that generalization occurs. This book argues that a crucial part of the answer lies in the long-standing bundling together of prediction, causation, and explanation, as though they were merely different

surfaces of one and the same theoretical achievement. As a result, when a theory does not directly increase predictive power, it is easily taken not to possess genuine causal insight either; and if it does not increase causal insight, it is then easily taken not to provide real explanation. It is through this chain of slippage that non-predictive frameworks gradually lose methodological standing.

To begin with, prediction, causation, and explanation do indeed often appear together. Many powerful scientific theories both explain why certain results occur, identify the relevant causal structure, and generate testable expectations in new cases. Precisely because the three reinforce one another in many successful episodes, it becomes easy to form a stronger impression: not only do good theories usually possess all three features, but the three seem almost naturally to entail one another. The better a theory explains, the more it reveals causation; the more it reveals causation, the more it predicts; and conversely, if a framework does not visibly enhance predictive power, it appears not really to have reached causal structure, and thus not really to have provided substantive explanation.

The problem is that the inference from “often appearing together” to “essentially indistinguishable” does not hold. Prediction, causation, and explanation are related, but they are not identical. A theory may, without directly yielding new predictions, still clarify the boundaries of certain causal questions. An account may, without centering on causal production, still constitute a genuine explanation. And a model may possess considerable predictive power without thereby yielding especially deep explanatory understanding. Unless these three functions are disentangled, methodological evaluation will remain vulnerable to domination by a standard that seems plausible only because it is too strong.

The conflation of prediction and explanation is especially important here. Explanation does not always aim to tell us what will happen next. Often it aims instead to show why certain results are possible, why certain paths are excluded, why a certain pattern can remain stable, or why a given difference is theoretically

relevant at a certain level. Accounts of this sort may greatly increase understanding without automatically translating into new empirical predictions. If explanation is wrongly compressed into “being able to say the result in advance,” then any framework whose primary work lies in organizing problem-space, constraining explanatory demand, or clarifying the structure of levels will appear to be operating outside explanation. The issue is not that such frameworks fail to increase understanding, but that “understanding” itself has been contracted into another name for predictive success.

The conflation of causation and explanation connects directly with the explanatory monism discussed in the earlier chapters. So long as explanation is tacitly treated as a single productive task, any account that does not in some way reveal difference-making, causal chain, or mechanistic pathway is easily suspected of not being a “real explanation.” And once that assumption is accepted, prediction naturally becomes its methodological echo: if genuine explanation reveals causal production, then good theory should obviously be able to predict future results on the basis of that causal structure. Explanatory standing and predictive standing thus reinforce one another. A framework that does not clearly increase prediction is suspected of not having genuinely revealed causation; and if it has not genuinely revealed causation, it is not thought to have genuinely explained. This chain of reinforcement is exactly one of the main reasons non-predictive frameworks are excluded.

It is at this point that the connection between the previous chapters and the present one becomes especially visible. This book has argued all along that explanatory activity includes not only causal production, but also other roles such as structural constraint and permissibility. If that is right, then prediction and explanation can no longer maintain a one-to-one correspondence. A framework may well provide substantive explanation by clarifying structural boundaries, identifying level-relevance, and excluding mistaken explanatory demands, without

these contributions first appearing as new predictions. Once explanatory space is pluralized in this way, prediction-centrism loses its grip on explanatory value.

It should also be noted that prediction and causation themselves do not perfectly coincide. Many predictions can succeed without especially clear causal understanding, and many causal insights do not directly translate into new empirical forecasts. The present book does not attempt to rewrite causal theory in full, but only to make one limited point: even if a framework's primary task is not to generate new predictions, it may still advance theory by clarifying which causal questions are legitimate and which causal demands are themselves misplaced. A framework need not directly increase predictive power in order to be deeply connected with causal understanding; it may do so by clearing the boundaries within which genuine causal inquiry can proceed.

This point matters especially for the present project. The work of the earlier chapters has not consisted in proposing a set of claims immediately convertible into experimental predictions. Rather, it has repeatedly shown that many existing disputes compress different explanatory tasks into one productive framework and thereby generate pseudo-problems. From emergence to higher-level explanation and onward, the book's mode of advance has been not primarily to "predict some new event," but to rewrite a problem that had itself been miswritten. So long as prediction, causation, and explanation remain fused, work of this sort will appear to evade the real theoretical task. But once the three are disentangled, we can see that such work is not decorative theorizing on the periphery. It is a reconstruction of what counts as a legitimate problem, what counts as a legitimate explanation, and what kinds of causal demand ought not to be made in the first place.

The key conclusion of this section may therefore be stated as follows: non-predictive frameworks appear methodologically suspect not because they are inherently loose, but because prediction, causation, and explanation have long been measured as one unified achievement. So long as that conflation remains in

place, a framework that does not directly enhance predictive power will automatically be suspected of also lacking genuine causal insight and genuine explanatory force. Once, by contrast, we acknowledge that the three are related but not flattenable into one another, we can begin to evaluate more fairly those frameworks that advance knowledge chiefly by clarifying structure, limiting theoretical space, and reorganizing the form of problems.

This makes the next question sharper. If non-predictive frameworks are not automatically devoid of theoretical content merely because they do not make predictions, then what exactly are they doing? Put differently: if the value of such frameworks is not first expressed in predictive success, how should their positive function be described? The next section addresses that question directly.

6.4 Theoretical Functions of Non-Predictive Frameworks

If prediction, causation, and explanation should not be flattened into one and the same kind of achievement, then the next question is what exactly a framework is doing when its primary aim is not direct prediction. Unless this can be stated positively, non-predictive frameworks will continue to be misunderstood as a merely negative category, as though they were either half-finished theories that have not yet achieved predictive power or substitutes that retreat into the conceptual realm only because empirical means are lacking. It is precisely this understanding that the present book rejects. Non-predictive frameworks are distinctive not because they lack prediction, but because they bear another set of theoretical functions. These functions are not prologues or footnotes to predictive work. They are part of what makes scientific understanding possible at all.

First, a non-predictive framework may reconstruct the form of a problem. The reason theoretical disputes become entrenched is often not that there are too few answers, but that the problem itself has been posed in the wrong way. Some questions are framed too narrowly, some too strongly, and some bind together

demands that do not belong at the same level in the first place. In such cases, pursuing more empirical results along the original line may not yield genuine progress. A deeper advance may consist in seeing that the difficulty arises because the problem was miswritten from the start as a different kind of problem. The role of a non-predictive framework here is not to offer a new answer to the old question, but to show why the form of the old question was itself distorted. Without such work, empirical inquiry can become ever more detailed while the theoretical deadlock grows ever more entrenched.

Second, a non-predictive framework may constrain explanatory demand. Not every seemingly natural theoretical demand is in fact legitimate. Some demands already smuggle in unexamined assumptions at the moment they are posed, for example, the assumption that all explanation must be productive, or that all genuine novelty must appear as additional causal power. Framework-level work in this respect tells us which explanatory demands are themselves misplaced and which questions ought not to be posed in the way they have been posed. This is not an evasion of explanation. On the contrary, it is a cleaning of explanatory space, ensuring that explanation responds only to demands that are genuinely meaningful. A framework that reduces mistaken demand has already advanced theoretical understanding rather than retreating from it.

Third, a non-predictive framework may organize theoretical space. A field advances not only by accumulating more facts, but by learning to distinguish which paths are genuinely open, which inferences are unnecessary, which concepts need to be separated, and which apparently conflicting accounts are actually answering different questions. Part of theoretical progress is therefore the reorganization of space itself. By drawing boundaries, dismantling confusions, and identifying levels and roles, a non-predictive framework provides a clearer map of the field. Such a map may not immediately tell us some new empirical fact, but it can determine which theoretical paths are worth pursuing, which questions are worth continuing

to ask, and which conflicts are merely products of a misplaced competitive structure.

Fourth, a non-predictive framework may exclude category mistakes. Many persistent theoretical difficulties arise not because the facts are too complex, but because things of different kinds have been placed within the same comparative frame. Structural constraint is treated as if it were the same kind of explanation as causal production; higher-level autonomy is understood as if it had to mean causal independence; organizational boundary-questions are rewritten as problems of event-generation. Once category mistakes enter the theoretical core, the whole discussion begins rotating within the wrong coordinates. The importance of a non-predictive framework here lies precisely in its ability to show which comparisons do not hold, which exclusion-relations are artificially manufactured, and which demands appear forceful only because we are conceptually standing in the wrong place. Such work is not “mere philosophy on the side.” It is theoretical hygiene itself.

Fifth, a non-predictive framework may unify scattered phenomena. Many disputes from different areas appear unrelated on the surface while in fact sharing the same deep structure. A good framework may not first propose a new prediction, but it can show that apparently dispersed problems all arise from the same conceptual dislocation, the same explanatory schema, or the same methodological bias. That sort of unification is itself theoretical progress, because it allows us to stop treating each problem as an isolated case and begin seeing the recurrent organizing principle beneath them. Much of the work of the earlier chapters is exactly of this kind: from emergence to higher-level explanation and onward, many disputes are difficult not because each has separately failed, but because all are jointly constrained by explanatory monism and the priority of productive disclosure. To bring that shared structure to light is already a genuine framework-level achievement.

Sixth, a non-predictive framework may provide conceptual boundaries for empirical research. Empirical work does not proceed in a concept-free space. How we distinguish variables, define systems, determine the relevant level, interpret model-results, and decide when explanation is complete all depend on a prior conceptual framework. If the framework is unclear, empirical results are easily misread, and may even be recruited to support questions that were badly posed from the outset. A non-predictive framework does not replace experiment or modeling here; it ensures that the questions experiment and modeling are meant to serve are themselves clear. In that sense, it is not external commentary on empirical work, but part of the conditions under which empirical work can correctly locate itself.

From the standpoint of this book, these functions are not isolated from one another. Together they constitute a form of theoretical achievement distinct from direct prediction. A framework may fail to tell us first what we will observe tomorrow and yet still substantially deepen our understanding of a domain by reconstructing problems, limiting mistaken demands, organizing theoretical space, excluding category-confusions, unifying scattered disputes, and clarifying conceptual boundaries. Without such work, prediction itself may be pursued in the wrong way, interpreted wrongly, or placed where it does not belong.

For that reason, non-predictive frameworks should not be understood as “theories lacking prediction,” but as theories that advance understanding in another way. Their value does not lie in imitating the outward form of predictive theory, but in accomplishing tasks that predictive theory does not automatically accomplish. Such frameworks can of course fail: they may organize badly, clarify insufficiently, unify in a strained way, or fail genuinely to reduce theoretical confusion. But the way they fail should not be judged by whether they produce new predictions. It should be judged by whether they actually perform the functions just described.

This prepares the next step. If non-predictive frameworks do possess these positive functions, then clearly they cannot be evaluated solely in terms of predictive power. The question that follows is: if a framework's main role is structural, conceptual, and constraining rather than directly eventive, by what standards should we assess whether it has succeeded? In other words, the legitimacy of non-predictive frameworks cannot rest on abstract defense alone; it also requires a more concrete set of criteria of evaluation. The next section takes up that issue directly.

6.5 How Structural Frameworks Should Be Evaluated

If non-predictive frameworks are not naturally weaker simply because they do not generate predictions, but instead perform a different set of theoretical tasks, then the standards by which they are evaluated cannot simply remain the single measure of whether they produce new predictions. Otherwise we would merely acknowledge their existence conceptually while sending them back, methodologically, to the same disadvantaged position. The issue is no longer whether structural frameworks may sometimes happen also to yield empirical consequences. It is this: if a framework's main function is to reconstruct problems, constrain mistaken demands, organize theoretical space, and clarify explanatory roles, then by what criteria should we judge whether it is successful, unsuccessful, or merely empty redescription?

This book proposes at least five kinds of standard.

The first is problem-clarifying power. A structural framework should first be able to show where a given theoretical deadlock actually lies. It must not merely repeat the old problem in a new vocabulary; it must make visible how previously entangled demands can be separated, how certain conflicts are artificially compressed, and how certain questions have been miswritten from the start. If a framework claims to clarify while merely moving vagueness from one language to

another, then it has not genuinely advanced. If, by contrast, it makes more clearly distinguishable the different tasks, levels, and demands in a dispute, then even if it has not yet produced all the answers, it has already achieved methodological progress.

The second is the capacity to dissolve mistaken demands. Many frameworks matter not because they add more explanation, but because they show that some strongly posed explanatory demand is itself unnecessary. A good structural framework should be able to distinguish which theoretical pressures arise from genuine difficulty and which arise from misplaced questioning; which disputes are persistent not because the world is intrinsically contradictory, but because explanation is being asked to perform a task it ought not to perform. If a framework can reveal that some supposedly unavoidable question is in fact a pseudo-problem, an overstrong demand, or a category mistake, then it has already displayed a significant theoretical gain. That gain is not weaker than prediction; it simply operates differently.

The third is space-organizing power. The value of a structural framework lies not only in clarifying one local concept, but in reorganizing the map of a theoretical field. It should help us see which positions actually share a common presupposition, which disagreements are only surface-level conflicts, which research paths are genuinely independent, and which are in fact competing for a position that was itself wrongly defined. If a framework can identify a common structure across several issues that appear unrelated and thereby rearrange the theoretical landscape, then it is doing more than tidying language. It is restructuring the field's orientation.

The fourth is the capacity to exclude category-confusion and false competition. A framework succeeds when it can show that certain oppositions are artifacts of a mistaken classificatory scheme rather than real theoretical alternatives. If it can demonstrate that two accounts are not actually in conflict because they answer

different questions, or that some demand has been wrongly imposed because a difference of explanatory role has been mistaken for a difference of ontological commitment, then it has accomplished something substantial. Structural frameworks should therefore be judged not only by what new claims they add, but by what confusions they remove.

The fifth is the capacity to provide disciplined conceptual boundaries for empirical work. A structural framework should not float free of inquiry. It should help determine how variables are carved up, what counts as the relevant level of analysis, what kind of model-result matters, what counts as explanatory completion, and which empirical questions are worth pursuing. If a framework sharpens the boundary conditions under which empirical work can correctly situate itself, then it contributes to scientific progress even without directly generating a new prediction. Its value lies in helping empirical research ask the right questions under the right conceptual description.

These standards make clear that structural frameworks can fail in ways just as real as predictive theories can fail. They may fail to clarify; they may multiply distinctions without necessity; they may reorganize space badly; they may not really eliminate confusion; they may leave empirical inquiry no better oriented than before. In that sense, they are not protected from criticism merely because they are non-predictive. Their legitimacy depends on whether they actually accomplish the structural tasks they claim for themselves.

This also means that a structural framework need not apologize for not being primarily eventive. Its success is not measured by whether it says in advance what will happen tomorrow, but by whether it enables us to understand what kind of question is worth asking tomorrow, under what conceptual conditions an answer would count, and which theoretical demands should never have been posed in the first place. That is not a lesser form of rigor. It is rigor exercised at a different level.

Seen this way, the evaluation of structural frameworks becomes continuous with the multi-role conception of explanation developed earlier in the book. Just as causal production, structural constraint, and permissibility require different criteria appropriate to different explanatory tasks, predictive and non-predictive theoretical work require standards appropriate to different modes of contribution. The mistake in both cases is the same: treating one successful form of achievement as the universal model of all achievement.

This is the point at which the methodological significance of the book's earlier arguments becomes fully visible. If a framework shows that several disputes were misframed from the outset, that explanatory roles had been collapsed, that false causal demands were imposed, and that structural and higher-level accounts occupy legitimate positions within a clarified map of inquiry, then the framework has already done serious scientific work. It has not merely decorated the margins of theory. It has altered the conditions under which theory can proceed well.

The next step is to make this point explicit with respect to the book's own framework. If non-predictive frameworks can be evaluated in these ways, then what exact methodological status should be assigned to the multi-role view of explanation defended in this book?

6.6 The Methodological Status of the Multi-Role View of Explanation

If structural and non-predictive frameworks can be methodologically legitimate, then the next question is how the multi-role view of explanation defended in this book should itself be situated. The answer proposed here is that it should be understood neither as a replacement empirical theory nor as a merely external philosophical commentary on science. Its status is methodological and structural: it is a framework that seeks to clarify what kinds of explanatory demand are legitimate, how different explanatory roles should be distinguished, and under

what conditions theoretical disputes are being distorted by compression into a single productive model.

This means, first, that the multi-role view does not compete with empirical theories by attempting to generate a rival body of eventive predictions. It does not aim to displace physics, biology, neuroscience, or other empirical sciences with a superior predictive machinery. Its contribution lies elsewhere. It reconstructs the form of several recurrent disputes, shows that apparently different deadlocks share a common explanatory distortion, and provides a conceptual map by which those disputes can be re-asked in a less misleading way. In that respect, its role is comparable to that of a framework that reorganizes the problem-space rather than adding one more local hypothesis within it.

Second, the multi-role view should be regarded as methodologically disciplinary. It does not merely broaden explanatory vocabulary; it imposes stricter requirements on theoretical judgment. By distinguishing causal production, structural constraint, and permissibility, it forces us to specify what kind of question is being asked, what kind of answer is appropriate, and whether a demand for explanation has already smuggled in an illegitimate assumption. In this way, the framework does not loosen standards. It tightens them, but at the level of problem-formation and explanatory allocation rather than at the level of direct prediction.

Third, its contribution is negatively productive in an important sense. It produces progress not only by adding concepts, but by subtracting confusions. It shows why certain familiar demands are overstrong, why certain oppositions are falsely competitive, why certain metaphysical pressures are generated by explanatory miswriting, and why some theoretical burdens should never have been placed where they were. This kind of negative productivity is not absence of contribution. It is a substantive form of theoretical advancement: progress through the disciplined removal of pseudo-problems, false alternatives, and misplaced explanatory expectations.

Fourth, the framework is integrative across domains. One of its major methodological strengths is that it does not treat emergence, higher-level explanation, entanglement, ontological inflation, and related issues as isolated local puzzles. It identifies a common deep structure across them, namely, the contraction of explanation into a single productive role. That integrative function is itself a sign of serious theoretical standing. A framework that can show how multiple apparently disconnected debates are governed by a shared structural distortion is not merely inventing a vocabulary; it is reorganizing the architecture of inquiry.

Fifth, the multi-role view remains open to discipline and criticism. It is not protected from evaluation simply because it is non-predictive. It can fail if its distinctions are unclear, if it redescribes without illuminating, if it multiplies roles without necessity, if it cannot actually reduce theoretical confusion, or if it offers no real guidance for how disputes should be reframed. Its legitimacy therefore does not rest on immunity from test, but on its capacity to perform the methodological tasks appropriate to frameworks of its kind. That is precisely why it should be taken seriously as knowledge work rather than dismissed as decorative philosophy.

This also clarifies the sense in which the book's framework is scientific without being itself a predictive scientific theory. It is scientific in the sense that it aims to improve scientific understanding by clarifying explanatory roles, organizing legitimate theoretical space, limiting unwarranted ontological expansion, and giving empirical inquiry a more disciplined conceptual setting. It does not stand outside science as a detached observer; nor does it stand within science as one more causal model. It occupies an intermediate but rigorous position: a framework for the ordering of explanatory practice itself.

From this perspective, the multi-role view of explanation belongs to the class of theories or frameworks whose importance lies in making other theoretical and empirical work more intelligible, more disciplined, and less prone to pseudo-

problem formation. That is not a second-rate role. Indeed, in many cases such work is a precondition for empirical findings and predictive models to be interpreted correctly. A field that cannot distinguish its explanatory roles is often a field that does not yet know what would count as resolving its disputes.

The methodological standing of this book's framework can therefore be expressed in a compact way: it is a non-predictive but scientifically serious framework because it reconstructs explanatory space under conditions of conceptual compression. Its legitimacy rests on clarification, constraint, integration, and the reduction of theoretical mislocation, not on immediate novel prediction.

This makes possible a broader conclusion. Once non-predictive frameworks are granted disciplined methodological standing, the contrast between "scientific theory" and "mere conceptual work" becomes less crude. Some conceptual work is itself part of scientific advance, provided that it clarifies the right things, constrains the right demands, and reorganizes the field in ways that matter for understanding. The multi-role view defended here is proposed as exactly such a case.

The next step is to gather the chapter's argument into a more explicit conclusion about why prediction-centrism should be limited rather than abolished, and how scientific understanding becomes broader once that limit is recognized.

6.7 Chapter Summary: Limiting Prediction-Centrism Without Rejecting Prediction

The task of this chapter has been to defend the scientific standing of non-predictive frameworks. The argument has not been that prediction is unimportant, still less that science ought to abandon predictive ambition. The argument has been that prediction has too often been generalized from one powerful form of theoretical achievement into the sole final measure of theoretical legitimacy. Once that happens, frameworks whose main contribution lies in clarification, constraint,

reorganization, and conceptual boundary-setting are pushed to the margins, even when they play an indispensable role in scientific understanding.

Several conclusions follow.

First, non-predictive frameworks are persistently suspected not because they naturally lack value, but because modern scientific culture has long operated with an overly compressed image of achievement. So long as theoretical success is presumed to show itself first and foremost in new prediction, any framework whose primary role is structural, conceptual, or disciplinary will appear second-rate by default.

Second, the methodological privilege of prediction has real historical and practical justification, but those justifications do not support turning prediction into a universal criterion for all theory. Prediction is especially powerful because it is connected with testability, fallibility, control, intervention, and the appearance of theoretical unity. Yet what makes it central for some tasks does not make it the proper standard for every task.

Third, much of the pressure against non-predictive frameworks is sustained by the conflation of prediction, causation, and explanation. Once these three are bundled together as if they formed a single achievement, any framework that does not directly increase predictive power is automatically suspected of also lacking causal insight and explanatory force. The earlier chapters have already shown why explanation itself cannot be reduced in that way. This chapter extends the same lesson to methodology.

Fourth, non-predictive frameworks perform positive theoretical functions of their own. They reconstruct the form of problems, constrain illegitimate explanatory demands, organize theoretical space, exclude category-mistakes, unify dispersed disputes, and provide conceptual boundaries for empirical work. These are not merely preparatory or decorative tasks. They are part of the way scientific understanding advances.

Fifth, structural frameworks can be evaluated rigorously even when they do not produce new predictions directly. They should be judged by their power to clarify problems, dissolve mistaken demands, reorganize theoretical space, eliminate false competition, and supply disciplined conceptual orientation for inquiry. Their standards of success differ from those of eventive predictive theories, but they are not thereby weaker or less serious.

Sixth, the multi-role view of explanation defended in this book should be understood as precisely such a non-predictive but scientifically serious framework. Its contribution is methodological and structural: it reconstructs explanatory space where explanatory roles have been compressed, identifies common distortions across multiple debates, and provides a disciplined way of reallocating explanatory demands. Its legitimacy lies not in replacing empirical science with a rival predictive model, but in clarifying the conceptual conditions under which empirical science can ask the right questions.

If the main point of the chapter is compressed into a single judgment, it is this:

Prediction should remain a central mark of some scientific achievements, but it should not function as the sole tribunal for all forms of theoretical progress.

Once that limit is recognized, a broader image of scientific understanding comes into view. Scientific progress does not consist only in saying in advance what event will occur. It also consists in learning which questions are well-posed, which explanations are genuinely needed, which theoretical spaces are open, which demands are illegitimate, and how apparently separate problems are structurally related. A science that loses sight of these tasks narrows its own understanding even when its predictive machinery grows more powerful.

This broader picture also prepares the way for the next chapter. If explanation is multi-role and theory-assessment cannot be reduced to prediction alone, then cases long treated as requiring additional causal stories may have to be reconsidered. One of the clearest such cases is quantum entanglement. The next

chapter will argue that entanglement has often been burdened with a misplaced demand for nonlocal causation, when what may really be required is a different kind of structural and permissibility-based understanding.

Chapter 7

Quantum Entanglement and the Misplaced Demand for Nonlocal Causation

In contemporary philosophical discussions of quantum theory, few problems have done as much as entanglement both to provoke explanatory impulse and to generate conceptual anxiety. Strong correlations appear between spatially separated systems, and those correlations violate constraints supported by certain classical intuitions. It therefore feels natural to think that something must be happening here that exceeds the ordinary picture of local causation. For this reason, entanglement has long been repeatedly formulated as a problem of nonlocal causation, as if, once two distant events display correlations that cannot be accommodated by local hidden-variable models, we must go on to ask: how exactly does some influence propagate across space? What hidden connection binds the two local events together?

The central claim of this chapter is that, although this way of posing the issue is intuitive, it is not necessarily the right form of the problem. Entanglement continues to seem mysterious not merely because quantum theory presents us with anomalous correlations, but because those correlations are almost automatically placed within a productive explanatory framework. It is tacitly assumed that if results at two locations are highly correlated, then explanation must show how some influence establishes a link between them; and if that link cannot be expressed locally within ordinary spacetime, explanation is quickly pushed toward nonlocal causation, action at a distance, hidden connections, or stronger holistic ontological commitments. In other words, entanglement is read as a problem of nonlocal causation not only because the experimental results are

strange, but because correlation is presupposed to require a propagative explanation.

What this chapter argues is that this demand is, to a large extent, misplaced. Entanglement does not necessarily require us first to add a new causal story of propagation. It may instead be better understood as a problem of structural inseparability. That is to say, the key issue may not be how some influence travels from one place to another, but rather that the space of permissible joint outcomes in the combined system is not generated by simply multiplying two independent local spaces of permissibility. If that is right, then the explanatory center of gravity in entanglement shifts from “what was propagated?” to “what forms of joint outcome are open to this system?” At that point, the framework of structural constraint, permissibility, and differentiated explanatory roles developed in the previous chapters ceases to be merely a general philosophical proposal and is tested in a highly sensitive case.

The line of argument in this chapter follows the same pattern as in the earlier chapters. First, it must explain why entanglement is so quickly read as a problem of nonlocal causation, and how this reading acquires its pressure within the Bell picture (Bell, 1964). At least under one explanatory rewriting, the core pressure of quantum entanglement need not first be formulated as a problem of nonlocal causal propagation, but can instead be formulated as a problem of the structural inseparability of joint outcome-space. Second, it must identify what sort of leap occurs between correlation and the demand for a propagative explanation. Third, the chapter introduces structural inseparability as an alternative understanding, and explains why entanglement is better interpreted as a problem of constraint on joint outcome-space than as a problem of anomalous influence-propagation. Finally, the chapter explains that this rewriting does not attempt to add a new empirical theory to quantum mechanics, but rather reorganizes a long-misplaced philosophical demand for explanation.

This chapter does not attempt to offer a new interpretation of quantum mechanics, nor does it seek to render a final verdict among Bohmian theory, the Many-Worlds interpretation, Relational Quantum Mechanics, QBism, or other interpretive options (Bohm, 1952; Wallace, 2012; Rovelli, 1996; Fuchs, Mermin & Schack, 2014). Its aim is more limited: to analyze a common shift in problem-formulation, namely, why joint correlations in quantum entanglement are so often rapidly translated into a demand for nonlocal causal propagation. In other words, the chapter concerns the form of explanatory demand, not a physical model intended to replace quantum theory.

To avoid overstatement, the chapter also distinguishes several different technical claims: Bell nonlocality is not the same thing as controllable superluminal signaling; local causality, separability, parameter independence, and outcome independence should not be treated as interchangeable. No-signaling constrains the ways in which entanglement can be understood as an ordinary propagation mechanism, but it does not remove the pressure entanglement exerts on a locally separable picture. Accordingly, what this chapter opposes is not every nonlocal formulation in the foundations of quantum theory, but the compression of these different layers of pressure into the single philosophical demand that there must be nonlocal causal propagation.

7.1 Why Entanglement Is So Often Read as a Problem of Nonlocal Causation

Entanglement is almost naturally read as a problem of nonlocal causation, first of all because it directly offends a very deep explanatory intuition: if there is a stable and significant correlation between two events, then that correlation seems to require some mechanism of connection. This intuition becomes especially strong when the correlated outcomes occur under conditions of spatial separation. It seems almost irresistible to think that if the local states are separated from one

another and yet the outcomes are somehow locked together, then explanation must tell us what exactly links them. Without such a link, the correlation itself looks like a coincidence that has not really been explained.

In classical experience, this intuition is usually reliable. We are accustomed to taking correlation as a sign of some connection because many correlations do in fact arise from common causes, direct interaction, background mechanisms, or traceable chains of transmission. Precisely for that reason, when quantum correlations display features that cannot be handled by certain local causal models, the first reaction is usually not to revise explanatory demand, but to intensify it: if ordinary local mechanisms are not enough, then some stronger connection must be needed; if the usual pathways of transmission are not enough, then some sort of action at a distance must be needed; if classical structures of independence are not enough, then some deeper holistic binding must be needed. In other words, quantum entanglement is easily turned into a philosophical problem of “nonlocal causation” not because the experimental results directly say that there is action at a distance, but because we bring to the problem a pre-fixed explanatory template.

The Bell picture intensifies this pressure still further. The violation of Bell inequalities clearly shows that certain classical local hidden-variable structures cannot reproduce quantum predictions (Bell, 1964; Clauser et al., 1969; Aspect, Dalibard & Roger, 1982). For many readers, that result is almost immediately translated into the thought that the local causal worldview has failed, and that some nonlocal causal structure must therefore take its place. The difficulty is that this translation is not identical with Bell’s result itself. What Bell excludes is a certain class of joint conditional structures, not a direct proof that some propagative influence has crossed space. Yet in ordinary discussion the two steps are often compressed together: once the local hidden-variable picture fails, explanation rapidly turns to the question of how influence propagates farther than locality would allow.

This turn is so persistent in part because “nonlocal causation” can seem like the only way of taking the problem seriously. Anyone who says that entanglement does not require action at a distance is easily suspected of merely changing the words or evading the demand to explain the “real connection.” In that sense, the demand for nonlocal causation in the entanglement debate is not merely one theoretical option; it almost carries the force of an explanatory ethic, as though anyone who does not provide a propagation story has not genuinely faced the depth of quantum correlation. For that reason, any attempt to rethink the issue cannot be content merely to conclude that there is “no nonlocal causation.” It must go further and show why treating correlation automatically as a demand for propagation is itself an explanatory habit that deserves scrutiny.

From the perspective of the preceding chapters, this structure is familiar. Once again we see the same pattern: a phenomenon appears difficult to understand not simply because we lack more facts, but because the problem has already been written in advance as a productive explanatory task. So long as we assume that to “truly explain correlation” is to show how influence establishes a link between events, entanglement will almost inevitably be read as a problem of nonlocal causation. If, by contrast, we allow another possibility, namely, that correlation sometimes reveals not an anomaly of propagation but the inseparability of permissible structure itself, then the space of explanation changes significantly. The problem is no longer first of all “what travels faster than locality allows?” but rather “why does the joint system open only certain combinations of outcomes and not others?”

A provisional conclusion can therefore already be drawn: entanglement is so often read as a problem of nonlocal causation not because quantum theory directly commands such an interpretation, but because we are accustomed to handing significant correlations over to a productive explanatory template. Within that template, Bell pressure appears especially strong, because it excludes certain

local conditional structures and thereby naturally encourages a turn to stronger propagative connections. The next task is to take that leap apart: what, exactly, does the Bell picture rule out, and at what point is the demand for nonlocal propagation built too quickly out of mere correlation?

7.2 The Standard Pressure in the Bell Picture

If the previous section explained why entanglement is almost naturally read as a problem of nonlocal causation, then the next step must be to enter more specifically the theoretical background that gives this reading its greatest force, namely, the Bell picture. For many discussants, what makes entanglement not merely “a bit strange,” but a genuine challenge to locality and causation, is precisely the pressure generated by Bell inequalities and their violation. Unless the standard form of that pressure is first clarified, the later rewriting in terms of structural inseparability will easily be mistaken for an evasion of the hardest part of the problem.

The starting point of the Bell picture is well known. Consider two systems that have interacted in the past and are then measured at different locations once spatially separated. Quantum theory predicts that the measurement outcomes on the two sides will display a certain strong correlation, and that under suitable experimental arrangements this correlation will violate a class of inequality constraints satisfied by local hidden-variable models. Bell’s key contribution is not merely to note that “quantum correlation is strong,” but to show that if we simultaneously retain certain classical intuitions, for example, that measurement outcomes are determined before measurement by local states or common hidden variables, and that spatial separation means that the measurement setting on one side cannot immediately affect the result on the other, then those intuitions together cannot reproduce quantum predictions.

For that reason, the pressure in the Bell picture does not arise merely from the experimental values themselves, but from the fact that a set of assumptions that initially look quite natural are forced to become jointly unsustainable. It compels us to recognize that, at least under certain traditional interpretations, locality, predetermined value-structure, and quantum correlation cannot all be retained together. That is why Bell's argument is so philosophically powerful. It does not merely say that "quantum systems are somewhat counterintuitive"; it says that if one wishes to preserve a certain classicalized causal picture, one must pay a price. That price may take the form of abandoning local hidden variables, abandoning some independence condition, reinterpreting realism, or accepting a stronger holistic structure.

In ordinary discussion, this pressure is often compressed into a more direct formula: Bell's result shows that the local causal picture has failed. It is precisely here that the temptation toward nonlocal causation becomes especially strong. For once "the failure of local causation" is treated as the core issue, the most natural next step seems to be the search for some stronger causal connection to replace it. If local common causes are not enough, then perhaps there is direct influence across space; if classical local mechanisms are not enough, then perhaps there is some deeper connection. In this way, the standard pressure in the Bell picture is often reformulated as the claim that quantum theory forces us to accept some kind of nonlocal causal structure.

What must be treated with the greatest caution, however, is that this step is not Bell's result itself, but only one way of interpretively organizing it. Bell's theorem does indeed rule out certain combinations of conditions, but it does not logically entail that there must be superluminal causal propagation. More exactly, it tells us that quantum correlations cannot be reproduced by a certain class of model satisfying certain conditions of local causality and predetermined structure. It does not automatically tell us that the only remaining path is to understand correlation

as influence transmitted across space. The pressure generated by the Bell picture is real, but its direction is not uniquely fixed in the way ordinary summaries often suggest.

This matters because much of the philosophical anxiety surrounding entanglement arises precisely from an overhasty leap from “model restriction” to “propagative conclusion.” It is easy to move directly from “local hidden-variable structure does not work” to “there must be some nonlocal influence,” as if no route remained other than adding an anomalous link into spacetime. But the earlier chapters of this book have repeatedly argued that theoretical difficulty often arises not from the facts themselves, but from the explanatory roles being narrowed in advance. If Bell pressure is placed from the start into a productive explanatory template, then it will naturally appear to compel us toward some new mechanism of propagation. But if the core issue is not first of all propagation, but the structural conditions on joint outcome-space, then the meaning of Bell pressure is rearranged.

This helps us distinguish two layers more clearly. The first is the constraining fact revealed by Bell’s argument: certain joint conditional structures cannot accommodate quantum correlation. The second is the philosophical interpretation of that constraining fact: must we understand it as evidence for nonlocal causal propagation? The first is something this chapter must take with full seriousness. The second is precisely what the chapter aims to reopen. In other words, the chapter does not seek to weaken Bell’s force, still less to retreat into some easy consolation of locality restored. On the contrary, its purpose is to preserve the hardness of Bell pressure while refusing an additional explanatory demand that Bell itself does not directly impose.

From this perspective, the standard pressure in the Bell picture can be more accurately stated as follows: quantum theory forces us to abandon certain joint models based on local independence and predetermined structure, but it does not

automatically determine how the explanatory picture must be rewritten. We may of course choose nonlocal causation as the direction of rewriting, but another direction may also be available: to understand the issue as one of the joint system's space of permissible outcomes not being assembled out of two independent local spaces of permissibility. If so, then what Bell pressure points toward is not first of all an anomaly in propagation-paths, but structural inseparability.

The conclusion of this section is therefore twofold. First, the pressure in the Bell picture must be fully preserved, because it clearly shows that certain classical local conditional structures are insufficient for quantum correlation. Second, this pressure is not automatically equivalent to the conclusion of nonlocal causal propagation; it is commonly read that way only because the demand for explaining correlation has from the outset been taken over by a productive template. The next question, then, is to take that takeover itself apart: where exactly does the leap from correlation to propagative explanation occur, and why does that leap feel natural while remaining unnecessary?

7.3 The Miswriting from Correlation to Propagation

If the standard pressure in the Bell picture does not automatically yield nonlocal causal propagation, then what most needs to be clarified next is how this transition takes place. Why is it that, starting from the claim that certain local conditional structures are insufficient to reproduce quantum correlations, people so naturally go on to think that there must be some superluminal influence establishing the connection? The judgment of this book is that what happens here is not a logically unavoidable inference, but the rapid takeover of the issue by a familiar explanatory habit. More precisely, once correlation is taken to require explanation, it is almost immediately handed over to a productive explanatory template, and the question is rewritten as follows: how is some connection established between the two

outcomes? If ordinary local mechanisms fail, then a stronger story of propagation must be supplied.

This miswriting is so common because, in many non-quantum cases, it is in fact effective. We are accustomed to moving from correlation to propagation or connection because, in classical experience, many stable correlations really can be explained by common causes, local interactions, or some traceable mechanism. In other words, the thought that correlation calls for connection is not absurd; in many contexts it is a reasonable explanatory instinct. The problem is that this instinct easily slides from being a template that is often useful into the only legitimate template. Once that happens, the anomalous correlations exhibited in quantum entanglement are no longer allowed to be understood first as a problem of joint structure, but are automatically required to provide a propagation pathway.

In the terminology developed in the earlier chapters of this book, what occurs here is a typical miswriting of problem-type. What might originally need to be discussed is why the space of permissible joint outcomes is organized as it is, why some combinations are systematically allowed while others are excluded, and why this organization cannot be decomposed into two mutually independent local structures of permissibility. But once correlation enters the productive explanatory template, the question is rewritten as: from one side to the other, what exactly is at work? Who affects whom? And how does that effect propagate through space? In this way, joint structure is compressed into causal pathway, and what is strange about quantum correlation is redescribed as an anomaly of propagation.

This miswriting deserves particular attention because it does not merely add one more philosophical option; it quietly changes what counts as “really facing the problem.” Once the propagation template takes over the discussion, anyone who fails to provide some story of connection is easily suspected of evading explanatory responsibility. It is as though, unless correlation can be translated into a propagative connection, it has not yet truly been explained. For this reason,

rejecting nonlocal causation is not automatically treated as a legitimate reinterpretation, but more often as a retreat: one is simply refusing to acknowledge the depth of the connection while offering no genuine alternative. That is why, if we are to make room for another route, it is not enough merely to say that Bell did not directly prove propagation; we must go further and show why understanding “explaining correlation” as “showing how influence propagates” is itself a presupposition that ought to be examined.

This miswriting can be analyzed into three levels.

First, there is the leap from correlation to the demand for connection. Merely because two outcomes are stably correlated, we tend to assume that there must be some fact establishing a connection between them. This inference works in many classical cases, but it does not automatically show that the connection must take the form of productive propagation. Correlation may reveal not only some transmitted influence, but also a deeper structural restriction on the joint system.

Second, there is the leap from the demand for connection to the propagation template. Even if one grants that some connection must be explained, it does not follow that the connection must take the form of an influence flowing from one location to another. Yet in ordinary discussion, connection is almost immediately spatialized and eventized: explanation is required to tell us how one result “reaches” another by way of some path. In this way, the problem of structural relation is prematurely compressed into a dynamical picture.

Third, there is the leap from the propagation template to the conclusion of nonlocal causation. Once local models of propagation fail, the most natural substitute seems to be that the propagation must therefore be nonlocal. But if the first two steps were already unexamined, then the third step merely seeks a stronger version within an unsuitable template instead of reconsidering whether the template itself is appropriate.

This analysis also helps explain why the Bell picture is so easily read as a “metaphysical blow” to locality. In fact, what Bell’s result directly undermines is a certain class of joint conditional structures. The reason it is so often redescribed as undermining the very “local worldview” is precisely that the problem of joint structure has already been translated into the problem of propagation. If we think only within the propagation template, then once locality can no longer support the correlation, nonlocal propagation appears to be the only remaining option. But if we step back and ask whether what truly needs to be explained here is propagation or rather the organization of joint outcome-space, then the center of gravity of the problem shifts.

This does not mean that correlation is never related to causal propagation, nor that every discussion of propagation is illegitimate. The present book is not proposing an absolute principle to the effect that all correlations should never be understood causally. Its claim is more limited and more rigorous: in the case of quantum entanglement, the automatic handing over of correlation to the propagation template requires further argument, and cannot simply be treated as the default starting point of explanation. Without such further argument, the demand for nonlocal causation is likely to be a pressure imposed by the template rather than a requirement directly posed by the phenomenon itself.

The key conclusion of this section, then, is that the transition from Bell-type correlation to nonlocal causation is not a seamless logical chain, but is structured by a set of explanatory habits that have not been sufficiently examined. Correlation is taken to require propagation; the problem of joint structure is thereby miswritten as a problem of causal pathway; and the difficulty of quantum theory is correspondingly rearranged. If this miswriting can be dismantled, then another formulation of the problem becomes possible: what entanglement reveals may not first be an anomaly of propagation, but rather the structural inseparability of the joint system. The next section develops that alternative directly.

7.4 Structural Inseparability: An Alternative Understanding

If the transition from correlation to propagation-based explanation is not inevitable, then the next question is how quantum entanglement might instead be understood. The alternative route proposed in this chapter is to understand entanglement first as a case of structural inseparability rather than as a case of anomalous propagative connection. By structural inseparability I do not mean that spatially separated systems are physically and mysteriously glued back together, nor that there exists some additional ontic bond between them. What I mean is that, in certain quantum cases, the structure of permissible joint outcomes cannot be reduced to an independent combination of the permissible structures of two local systems taken separately. In other words, what needs to be explained is not “how some influence travels from one place to another,” but “why this joint system opens only some combinations of outcomes and not others.”

The key to this reinterpretation is that it shifts explanatory focus from propagation between events to the organization of outcome-space. Traditional readings in terms of nonlocal causation assume that if the measurement outcomes on the two sides are stably correlated, then explanation must tell us how one outcome acts upon the other. The idea of structural inseparability suggests that this demand itself may be inappropriate. What quantum entanglement reveals may not be some action at a distance, but the fact that the joint system cannot be regarded as a simple juxtaposition of two mutually independent local state-spaces. The issue, in other words, is not how local outcomes are later connected, but that those outcomes are defined from the outset within an inseparable joint structure of constraint.

It is important here not to let “inseparability” degenerate into a vague slogan of holism. If the notion is to have theoretical force, it must be understood as a judgment about permissible structure rather than as rhetoric about a mysterious totality. To say that a joint system is structurally inseparable means that the

characterization of its outcome-space cannot first be completed at the local level and then mechanically assembled into a whole by putting the local outcomes together. On the contrary, the constraining structure of the whole system determines from the outset which combinations of joint outcomes are open, and this determination does not need to be established by some subsequent process of propagation. In the language of the earlier chapters, entanglement here does not demand the addition of new causal production; it reveals instead a structure of permissibility that differs from the classical picture of local assembly.

This reinterpretation matters because it offers a different explanatory direction. It neither denies the force of Bell nor retreats to a facile re-classicization of quantum correlation. On the contrary, it preserves the central fact revealed by Bell, namely, that certain structures of local independence are insufficient to accommodate quantum correlation. What it rejects is the further translation of that fact into the claim that there must therefore be some superluminal propagation. The claim of structural inseparability is that the pressure generated by Bell can be understood as a restriction on joint-system structure, rather than as evidence of anomalous propagative pathways in spacetime. Entanglement is thus special not because influence travels farther, but because the organization of outcome-space is not of the kind presupposed by classical local models.

From the standpoint of the overall framework of this book, this should sound familiar. The previous chapters have repeatedly argued that many theoretical deadlocks arise because questions about structural boundaries, organizational relations, and permissibility are miswritten as questions about event-generation. Entanglement furnishes a particularly sensitive version of the same phenomenon. So long as we continue to understand "explaining correlation" as "telling the story of how connection propagates," quantum theory will keep appearing to force us toward some more exotic causal mechanism. But if we step back and ask whether what truly needs explaining is the organization of joint outcomes rather than the

transmission of influence, the philosophical picture of entanglement changes quite dramatically.

This does not mean that structural inseparability automatically solves everything. It is not a new interpretive system for quantum mechanics, nor does it claim to replace the mathematical formalism of the theory. Its role is more limited, but also more crucial: it re-specifies the direction of philosophical explanatory demand. It asks that we not begin by treating quantum correlation as an anomaly of propagation, but instead treat it first as an anomaly of joint structure, or more precisely as the refusal of joint permissible space to submit to classical local decomposition. In this way, much of the pressure that was previously thought to require nonlocal causation is reassigned to the analysis of structural constraint and permissibility.

This also helps us distinguish more carefully between “the failure of locality” and “the failure of the picture of local assembly.” Bell’s result is often quickly summarized as a failure of locality, but the more careful formulation is that a certain kind of joint model based on local independence and predetermined structure fails. If so, then what we are facing may not be “an extra nonlocal causal line added to the world,” but rather “the world’s refusal to allow us to organize outcome-space by first fixing the local pieces and only later assembling the whole.” Structural inseparability captures exactly this point: entanglement does not primarily undermine causal speed, but the way in which outcome-space is decomposed.

That is also why this understanding can preserve the force of entanglement without ontological inflation. Under a nonlocal-causal reading, we are easily pushed toward stronger ontological commitments, as though we had to posit some new connective mechanism, some action-at-a-distance relation, or some nonstandard channel of propagation. The framework of structural inseparability is more restrained. It adds no new entity; it instead reinterprets the modeling demands we impose on joint systems. It claims that some of the “gaps” we thought

had to be filled by additional causes are in fact generated by our having wrongly presupposed how permissible structure ought to be decomposed.

The central conclusion of this section can therefore be stated as follows: quantum entanglement does not necessarily first require a nonlocal causal explanation; it may instead be understood as structural inseparability in the joint system. What is genuinely anomalous here may not be how influence propagates, but how outcome-space is organized. The former demands that we search for some superluminal productive connection; the latter requires us to rethink the system's permissible structure. Once this distinction is recognized, the philosophical pressure of entanglement no longer has to discharge itself only in the direction of nonlocal causation, but can be relocated within the framework of structural constraint and permissibility.

7.5 Entanglement, Constraint, and Permissibility

If entanglement is better understood as structural inseparability rather than anomalous propagation, then the next key question is what this inseparability means in explanatory terms. Using the language developed in the previous chapters of this book, the most natural answer is that entanglement should first be understood as a problem of structural constraint and permissibility. That is to say, what quantum entanglement reveals is not primarily the extension of some causal chain through space, but the fact that the outcome-space of the joint system is constrained from the outset in a way that resists local decomposition, so that only certain forms of joint outcome are allowed while others are excluded.

The core of this rewriting lies in shifting explanatory focus from "how the results affect one another" to "why certain combinations of results are allowed." Traditional readings in terms of nonlocal causation ask: what happens on one side such that the other side must change correspondingly? A reading in terms of constraint and permissibility asks instead: for this joint system, which combinations of outcomes

belong to its open space from the start, and which do not belong to it at all? These two ways of posing the problem may look similar, but they are fundamentally different. The former demands a story of propagation; the latter demands an account of structural boundary. The former understands the issue as one of connection between events; the latter as one of the organization of joint outcome-space.

From the standpoint of structural constraint, what is special about an entangled system is not that it contains an extra connecting line, but that its joint state-structure imposes restrictions on combinations of local outcomes. These restrictions are not retrospective. It is not that two local outcomes first form independently and are then forcibly coordinated by some action at a distance. Rather, the restriction arises from the structure of the joint system itself, and specifies in advance which combinations of local outcomes are open to that system and which are not. In other words, constraint here does not operate as a causal correction added afterward, but as the structural boundary of the joint system as such.

Once the problem is rewritten in this way, the importance of permissibility becomes immediately salient. Entanglement does not first require us to explain “why the outcome happened to come out this way on this occasion.” It requires us instead to explain why, for this joint system, only these and not those combinations of outcomes are possible, stable, and compatible with the structure of the whole. This is exactly the kind of task that an analysis of permissibility is suited to handle. It does not ask how influence crosses space; it asks what sort of result-space is open to a system. Put differently, the deeper difficulty in entanglement need not be an anomaly of generative pathway, but an anomaly of permissible space: the result-structure allowed by the joint system cannot be given by a simple composition of two locally allowed spaces.

It is worth emphasizing that this is not using “constraint” as a vague substitute term to cover over our ignorance of what happened causally. On the contrary, it is a precise change in explanatory task. When we say that a system is constrained, we do not mean “we have not yet found the real cause.” We mean that the structure of the system itself determines its open space. In the entanglement case, this means that we do not need to search for an additional channel of influence from A to B in order to explain why certain combinations of results do not occur. What we need is an account of why the structure of the joint system does not permit those combinations from the outset. That is not a failed substitute for causal knowledge, but a different explanatory target altogether.

Looked at from this angle, Bell pressure can also be restated more clearly. The significance of the violation of Bell inequalities does not lie in the fact that it directly forces some action at a distance upon us, but in the fact that it shows that if we try to understand joint outcome-space by a mode of local independent assembly, then quantum correlation cannot be accommodated. What is first challenged, therefore, is not an intuitive picture of “the speed limit of local causation,” but the deeper classical presupposition that joint permissible structure can be generated by multiplying local structures together. In other words, the difficulty exposed by Bell can perfectly well be read as a refutation of a certain way of decomposing permissible space, without being read immediately as a refutation of some path of propagation.

This also helps us distinguish more precisely between two different kinds of impossibility. In the classical propagation template, impossibility usually means that there is no sufficient causal path for establishing a connection. In the permissibility template, impossibility means that some combinations of outcomes do not belong to the open structure of the system at all. The former concerns how one gets from one result to another; the latter concerns which combinations of results are not within the system’s permitted range from the outset. A large part of why quantum

entanglement feels so powerful is that we are accustomed to understanding what is in fact closer to the second kind of impossibility in terms of the first.

Moreover, understanding entanglement in terms of constraint and permissibility has a methodological advantage: it allows us to preserve the strong limitations delivered by quantum theory while avoiding premature ontological inflation. If correlation can only be explained through nonlocal causation, then we are easily pushed toward the picture that there must be some additional mechanism of connection in the world. But if correlation arises first from structural constraint on the space of joint outcomes, then the ontological pressure is significantly reduced. We must still admit that classical local decomposition is insufficient. But we no longer need, on that basis alone, to commit ourselves immediately to some new superluminal productive entity or anomalous causal channel. In this way, the theoretical strength is preserved while the inflationary explanatory impulse is restrained.

Of course, this rewriting does not mean that causal processes are entirely irrelevant in the quantum case. Measurement, evolution, experimental arrangement, and apparatus setting all remain parts of concrete physical process. This book does not propose to erase all process-language by means of structural language. More exactly, the claim of this section is that in the entanglement problem the real philosophical difficulty does not lie in how those processes unfold as such, but in why the joint structure of results they yield cannot be understood in terms of locally independent permissible spaces. Thus, even if process-analysis retains its place, the explanatory center of entanglement should not automatically be surrendered to a propagative causal template.

The conclusion of this section may therefore be summarized as follows: if entanglement is brought back into the framework of structural constraint and permissibility, then we can say more clearly that what quantum theory truly requires us to revise may not be 'how causation must propagate,' but rather 'how the

outcome-space of a joint system must be organized.' Entanglement is profound not because it compels us to search for a stranger pathway of influence, but because it refuses the classical explanatory habit of putting the local first and assembling the whole afterward. In that way, the philosophical interpretation of entanglement shifts from a demand for nonlocal causation to a problem of the structural inseparability of joint outcome-space.

The next issue that must be treated is the relation between this rewriting and the concepts of locality and no-signaling. For once "nonlocal causation" is downplayed, the natural question arises: what, then, does locality still preserve? And what does no-signaling mean at the explanatory level? The next section reassigns the roles of these concepts within the discussion of entanglement.

7.6 Locality, No-Signaling, and Explanatory Roles

If entanglement is to be understood first as the inseparability of joint structure rather than as a mechanism of action at a distance, then the next natural question is what role locality and no-signaling still play in explanation. In standard discussion, these two concepts are often pulled directly into the same structure of tension. On the one hand, people wish to preserve some sort of local picture of the world; on the other hand, they see that Bell correlations cannot be produced by certain models of local independence, and thus they easily conclude that locality itself has been "pierced through" by quantum theory, while no-signaling is often treated as a kind of merely technical remainder that holds mathematically but is explanatorily insignificant. But if the rewriting developed earlier in this chapter is correct, then this way of arranging the concepts itself needs to be reexamined.

First, locality and no-signaling must be distinguished more clearly in terms of the work they perform in discussion. No-signaling says that, within the structures allowed by standard quantum theory, the measurement setting on one side cannot be used in a controllable way to transmit information to the other side. It is a

constraint concerning communicability and operational dependence of results. It tells us that, however strong entanglement-correlation may be, it does not open a channel for superluminal message transmission. Locality, by contrast, is more complex. It is often treated as a broader metaphysical or model-building requirement, meant to express the thought that spatially separated local events should not directly determine one another by action at a distance, or that joint structure should be intelligible in terms of local states and their local causal relations. The trouble is that, in actual discussion, these two layers are often compressed together, as though no-signaling preserved locality in some residual form, while any correlation that exceeds what classical local models can accommodate thereby destroys locality as such.

The claim of this book is that this compression conceals what most needs to be distinguished: locality and no-signaling do not constrain the same type of problem. No-signaling constrains, first of all, what is physically realizable at the level of operation and communication. It tells us that quantum entanglement does not provide a propagative mechanism by which distant results can be controlled. From the perspective of causal production, this is extremely important, because it means that we have no reason to identify entanglement too quickly with some usable process of action at a distance. That is to say, no-signaling weakens the intuitive force of the thought that "there must be some anomalous propagation here." If there really were some action-at-a-distance process that was controllable and exploitable in a manner analogous to classical causation, we ought to expect more direct traces of it at the level of communication. Quantum theory does precisely not give us that.

At the same time, however, no-signaling does not automatically solve the explanatory problem of entanglement. It does not tell us why the space of joint outcomes has this form rather than another, nor why some combinations are allowed while others are excluded. In other words, no-signaling constrains whether

correlation can be used for controlled transmission; it does not answer why joint outcome-structure is organized as it is. If these two layers are not distinguished, discussion tends to collapse into one of two equally crude reactions: either one says that since no-signaling holds, there is no deep problem at all; or one says that no-signaling is merely a surface technical condition, because the real issue still lies in some hidden connection. Both reactions miss the point that no-signaling does not cancel the need for structural explanation; it only blocks one intuitive route for a propagative explanation.

The case of locality requires even finer subdivision. If locality is understood as a principle governing controllable propagation of influence, then no-signaling preserves an important part of it. But if locality is understood as the stronger model-requirement that the structure of joint outcomes must be generated by independently combining local state-spaces, then entanglement clearly places real pressure on that requirement. What quantum correlation unsettles, in that case, is not necessarily the narrower form of locality according to which no controllable superluminal influence exists, but the stronger structural presupposition that a joint system can be understood by putting the local first and the whole afterward. This is precisely why the structural inseparability proposed earlier in the chapter can be seen as a relocation of the focus of the locality-problem: the difficulty lies not necessarily in the distance traversed by influence, but in the impossibility of decomposing outcome-space into locally independent pieces.

From the standpoint of explanatory roles, this point is especially important. The standard reading is inclined to think that locality has been “causally breached” because it presupposes that the core explanatory task is to show how some connection propagates between events. Once a local picture of propagation cannot accommodate Bell correlation, locality appears to have failed absolutely. But if the explanatory center of gravity is shifted to structural constraint and permissibility, then the problem is reorganized. Locality no longer appears first as a

binary verdict on whether superluminal propagation exists; rather, the questions become: to what extent is the joint structure of outcomes still constrained by local operations, and to what extent does it refuse local independent decomposition? In this way, locality is broken down from a single propagative concept into several layers differently related to operation, communication, structural decomposition, and explanatory demand.

This also allows a fairer understanding of the status of no-signaling. It is not an insignificant technical patch. It is an important sign that the strangeness of quantum entanglement does not manifest itself as usable action at a distance. If we still insist on describing entanglement as “some influence propagating,” then we must explain why that influence is supposed to be so explanatorily central while leaving so little trace at the operational level. By contrast, if entanglement is first understood as an anomaly in joint permissible structure, this tension is significantly reduced. No-signaling then no longer looks like a limitation imposed on nonlocal causation; it looks more like a reminder that perhaps the propagation-template was not the best point of entry from the start.

From the standpoint of the book’s general framework, this rearrangement has a broader methodological significance. The earlier chapters repeatedly argued that many theoretical difficulties appear insoluble because different explanatory roles have been compressed into a single demand. The discussion of locality and no-signaling in entanglement provides a vivid example. No-signaling properly constrains propagative operability, while the dispute over locality is often directly treated as if it were simply a question of whether productive causation can still be maintained. The result is that structural problems are mistaken for propagation-problems, and propagation-problems are mistaken for the whole of the explanatory issue. To reassign these roles is not to diminish the difficulty, but to return it to the place where it properly belongs.

The conclusion of this section can therefore be stated as follows: no-signaling does not eliminate the explanatory pressure of entanglement, but it weakens the rationale for automatically translating that pressure into anomalous propagation; and what locality is really challenged by may not be the prohibition of controllable causal influence, but the requirement that joint outcome-space be decomposable into locally independent parts. If that is so, then what most needs to be rewritten in discussion of entanglement is not the slogan “has locality collapsed altogether?” but the allocation of roles among locality, no-signaling, and explanatory demand. Only once that allocation is made clear does structural inseparability emerge not merely as an alternative vocabulary, but as a more accurate reorganization of the problem-form itself.

The next section draws a further boundary: what exactly this rewriting resolves, and what it does not resolve. Without clarifying both its gain and its limit, the chapter is liable to be misunderstood either as overclaiming that the problem has disappeared, or as covertly introducing a new interpretation of quantum mechanics. In fact, its aim is neither.

7.7 The Gains and Limits of This Rewriting

At this stage, the chapter has put forward a fairly definite alternative route: quantum entanglement should not first be understood as a problem of nonlocal causal propagation, but rather as a problem of the structural inseparability of the joint system. More specifically, this inseparability is better characterized in terms of structural constraint and permissibility than in terms of additional propagative connection. If this rewriting is to be genuinely persuasive, it must explain two things at once: what exactly it does resolve, and what it does not claim to resolve. Otherwise, it will either be misunderstood as an overly easy “philosophical cooling down,” as though the difficulty vanished merely by changing the language, or as a new interpretive system that seeks to replace existing physical theory.

Let us begin with what it does resolve. First, this rewriting clearly weakens a long-too-hastily imposed explanatory demand, namely, the demand that “correlation must be explained by propagation.” The argument of this chapter is not that all propagative language is meaningless, but rather that there is no seamless logical path from Bell-type correlation to the demand for action at a distance. Once that is seen, an especially constraining pseudo-dilemma in the entanglement debate is significantly weakened. We are no longer forced to choose between “either accept nonlocal causation or admit explanatory failure.” What is resolved here first, then, is a coercive feature of problem-formulation itself.

Second, this rewriting helps us distinguish more accurately among different kinds of theoretical pressure. What the Bell picture reveals is not simply an “anomaly in causal speed,” but the inadequacy of certain forms of local independent decomposition for handling joint outcome-structure. By bringing the problem back into the framework of structural constraint and permissibility, the chapter redistributes several difficulties that had previously been run together: some concern operation and signaling, some concern the organization of outcome-space, and some concern the explanatory template itself. The gain here does not consist in yielding new experimental predictions, but in preventing quantum entanglement from being crushed at once under all these different layers of pressure, and instead assigning them back to their proper places. That is, what is resolved is a problem of theoretical pressure being wrongly concentrated.

Third, this rewriting yields an ontological advantage: it reduces unnecessary inflationary impulse. If entanglement can only be understood through nonlocal causation, then we are easily driven toward stronger ontological commitments, as though the world must contain some extra line of connection, some action-at-a-distance relation, or some anomalous mechanism lying outside ordinary causal structure. The chapter does not prove that such commitments are impossible in every circumstance. It argues instead that, in many cases, our sense that we are

forced to go that far arises only because the propagation-template has been presupposed as the sole legitimate entry-point of explanation. Once that presupposition is itself lifted, ontological pressure no longer escalates automatically. The gain here is therefore first of all a gain in restraint rather than in expansion.

Fourth, the chapter also provides a concrete test of the general framework developed earlier in the book. If, in a case as sensitive as entanglement, the language of structural constraint, permissibility, and differentiated explanatory roles can still help us disentangle the problem rather than functioning only at an abstract level, then the earlier claims about the multi-role view of explanation receive stronger support. In other words, the chapter is not only about entanglement; it is also a test of whether the overall framework of the book can actually work in a high-pressure case. The gain in this respect is a gain in framework-applicability.

At the same time, the limits of the chapter must be stated equally clearly. First, the chapter does not provide a new empirical theory of quantum mechanics. It does not alter the formalism, propose new dynamical equations, or produce a set of experimental predictions to compete with existing theories. In other words, the problem it addresses is not “how quantum theory should be physically rewritten,” but “how we should philosophically organize its explanatory demands.” To require it to answer every dispute in physical interpretation by the standards of empirical theory would simply be another miswriting of problem-type.

Second, the chapter does not claim that all explanatory difficulties concerning entanglement are thereby dissolved. Even if structural inseparability is accepted as a legitimate rewriting, quantum theory may still present deep problems at other levels, for example the measurement problem, the interpretation of state-reality, comparisons among rival interpretations, and the relation between the mathematical form of joint systems and their ontological picture. The chapter does

not aim to provide a master-answer to all of these. What it addresses is one especially persistent pressure: the tendency to organize entanglement automatically as a problem of nonlocal causation. Once that pressure is removed, the remaining problems still exist, but they no longer stand under the same distortion.

Third, the chapter does not deny that some interpretations of quantum theory may continue to employ some form of nonlocal language. Its claim is not that every nonlocal formulation is mistaken. Its claim is that nonlocal causation should not be regarded as the single philosophical conclusion directly forced upon us by entanglement itself. What the chapter seeks is not to close every interpretive route, but to break an overly early monopoly of interpretation. Once that is granted, different interpretive approaches can be compared on a clearer terrain of problems instead of all being dragged into the default frame of propagation.

Fourth, the rewriting offered here does not mean that structural language itself needs no further clarification. To say that entanglement reflects a joint space of permissibility that cannot be locally decomposed is not yet to provide a complete, residue-free explanation. The framework still needs to be coordinated more precisely, in further theoretical work, with quantum formalism, interpretive schemes, and experimental structure. What this chapter does is to clear the problem-space for that further work, not to substitute for it altogether. It provides a better point of entry, not a self-executing final destination.

Seen in this light, the significance of the chapter is closer to a kind of philosophical decompression and relocation. It neither abolishes the difficulty nor manufactures a false ease. What it does is to take apart some pressures that had been wrongly piled on top of one another, and to show that not every quantum strangeness must bear its explanatory weight through an anomaly of propagation. Some of that weight properly belongs to joint structure itself, some to limitations on local independent models, and some to the misuse of explanatory template.

Only once those weights are redistributed can quantum entanglement be discussed more clearly.

The conclusion of this section may therefore be summarized as follows: what this rewriting genuinely resolves is a problem of misplaced explanatory demand, not the total problem of quantum interpretation. It weakens the pressure to slide automatically from Bell-type correlation to nonlocal causation, reorganizes the relation among locality, no-signaling, and outcome-space structure, and reduces unnecessary ontological inflation; but it does not provide a new physical theory, nor does it claim that all the philosophical difficulties of quantum theory are thereby finished. Its value lies not in “ending the debate,” but in making it possible for the debate to proceed from a less confused position.

This prepares the way for the chapter’s final summary. We can now state more clearly that the reason quantum entanglement has so long been taken to force some picture of action at a distance is not that the theory itself directly yields that conclusion, but that correlation, locality, and explanatory demand have long been organized within a single propagation-template. The next section gathers the main claims of the chapter together and carries the discussion forward to Chapter 8 on unknown phenomena and ontological stability.

7.8 Chapter Summary: From Nonlocal Causation to Structural Inseparability

The point of departure of this chapter was a challenge to the long-standing tendency to organize quantum entanglement in philosophical discussion automatically as a “problem of nonlocal causation.” On the surface, this organization appears extremely natural: spatially separated systems display strong correlation, local hidden-variable pictures cannot accommodate that correlation, and explanation is therefore rapidly pushed toward some superluminal connection, anomalous propagation, or stronger holistic ontological commitment.

What the chapter has argued, however, is that this turn is not directly forced upon us by quantum theory itself. It arises from a deeper explanatory habit: once a correlation is sufficiently striking, it is automatically handed over to a productive propagation-template. It is at that moment that entanglement is too quickly rewritten as the problem of “how influence crosses space.”

In order to dismantle this miswriting, the chapter first emphasized that the real pressure in the Bell picture should not be weakened. Bell’s result does indeed show that certain joint models based on local independence and predetermined structure cannot reproduce quantum correlation. That fact must be preserved in full. The issue is that the move from this fact to the conclusion that “there must therefore be some nonlocal causal propagation” is not logically seamless. What Bell directly excludes is a class of joint structural models, not the absence of every propagative connection. The hardness of the theoretical result lies in the restriction placed on classical local decomposition, not necessarily in a direct demand for action at a distance.

On that basis, the chapter proposed that quantum entanglement is better understood as a problem of structural inseparability. Structural inseparability does not mean that an extra line of connection has been added to space, nor that the whole mysteriously overrides its parts. It means that the permissible structure of joint outcomes cannot be obtained by independently assembling two local outcome-spaces. What needs to be explained first, therefore, is not how some influence travels from one side to the other, but why only certain combinations of outcomes are open to the joint system while others are excluded. In this way, the explanatory center of gravity shifts from anomaly of propagation to organization of outcome-space.

To keep this proposal from remaining merely verbal, the chapter then situated it within the framework developed in earlier chapters, arguing that entanglement is fundamentally a problem of structural constraint and permissibility. The

philosophical tension of quantum correlation need not be understood as some action at a distance exerting force between events; it may instead be understood as the fact that the joint system possesses from the outset a constraining structure that resists local decomposition. If understood in this way, Bell pressure is directed not at whether propagation reaches far enough, but at whether we have wrongly insisted on a picture in which the local comes first and the whole is assembled afterward.

The chapter also used this perspective to reorganize the roles of locality and no-signaling. No-signaling does not eliminate the explanatory difficulty posed by quantum correlation, but it does show clearly that entanglement does not provide a controllable mechanism of superluminal communication. The deeper challenge to locality may therefore lie not in “the world permitting action at a distance” so much as in “the structure of joint outcomes refusing classical local independent decomposition.” Once these roles are re-sorted, many sweeping slogans in the entanglement debate lose some of their earlier force.

At the same time, the chapter marked its own limits. It did not seek to choose among competing interpretations of quantum mechanics, nor to provide a new empirical theory, nor to solve every broader philosophical difficulty concerning measurement, state-reality, or interpretation. What it did was more limited but nonetheless important: it dismantled an explanatory demand that had been imposed too quickly, namely, the automatic translation of Bell-type correlation into a demand for nonlocal causal propagation. What was rewritten was the form of the problem, not the whole of the physical content. What was reduced was unnecessary ontological pressure, not the genuine difficulty of empirical theory itself.

Chapter 8

Unknown Phenomena, Ontological Robustness, and Explanatory Discipline

The preceding chapters have repeatedly shown that many apparently profound theoretical difficulties expand so quickly not simply because empirical facts themselves force us to do so, but because explanatory demands have from the outset been overly compressed. So long as explanation is tacitly treated as equivalent to causal production, and theoretical achievement as equivalent to predictive success, then whenever higher-level patterns, quantum correlations, or anomalous structures appear, it comes to seem almost natural that if the existing framework cannot accommodate them, what is needed must be a new causal power, a new connecting mechanism, or a new ontological level. The discussion of entanglement in Chapter 7 has already shown that this inflationary impulse is not always dictated by the phenomena themselves, but is often triggered by the way the problem has been miswritten.

Chapter 8 addresses the operation of this same mechanism in a more general setting. Quantum entanglement is not the only thing that provokes theoretical anxiety. Any unknown phenomenon, anomalous phenomenon, or empirical novelty may quickly provoke a similar reaction: perhaps the existing ontological inventory is no longer enough; perhaps we must add new entities, new forces, or even new levels of reality. Faced with phenomena that resist smooth integration, people easily form an almost instinctive judgment: explanatory difficulty shows that the world itself exceeds the current framework. In this way, empirical unfamiliarity is rapidly translated into ontological insufficiency, and temporary theoretical mismatch is elevated into evidence that reality itself must be structurally enlarged.

The central claim of this chapter is that this translation is usually too quick. Unknown phenomena do not first and foremost challenge our judgment concerning what there is in the world; more often, they challenge the way we assign explanatory tasks. What empirical novelty really tests first, in other words, is explanatory discipline rather than ontological inventory. A mature theoretical framework should not move immediately to ontological expansion whenever an anomalous phenomenon appears. It should first be able to distinguish whether what is needed is a finer empirical description, a revision of the model, a redistribution of explanatory roles, or only after these paths have been exhausted, a stronger ontological commitment. Without this distinction, theory remains in a condition of being led around by anomalies and cannot develop genuine robustness.

This chapter therefore proceeds as follows. First, it explains why unknown phenomena so reliably trigger ontological anxiety. Second, it reviews the recurring historical linkage between empirical novelty and ontological inflation. Third, it advances a more restrained alternative judgment: what unknown phenomena first test is the robustness of an explanatory framework, not the completeness of ontological inventory. Finally, it argues that ontological robustness does not mean conservatively rejecting novelty, but rather a more mature theoretical discipline: in the face of the unknown, neither rashly mystifying nor arbitrarily dismissing it, but more strictly distinguishing the order-relations among empirical pressure, explanatory reorganization, and ontological commitment.

8.1 Why Unknown Phenomena So Readily Trigger Ontological Anxiety

Unknown phenomena so often trigger ontological anxiety because they directly disturb a deep sense of theoretical security. So long as a phenomenon can be smoothly incorporated into existing concepts, models, and explanatory habits,

people rarely stop to ask whether there might be “more things in the world.” But once some object, behavioral pattern, or experiential result appears difficult to place, the theoretical reaction escalates rapidly. What began merely as the problem “we do not yet know what this is” is easily rewritten into the thought that “the current picture of reality is no longer enough.” Unknown phenomena are unsettling, then, not only because information is missing, but because they disturb the usually unspoken fit between framework and world.

This anxiety so often turns in an ontological direction because a “new phenomenon” is very easily understood, intuitively, as a sign of a “new existent.” People naturally feel that if something exceeds current explanation, the most direct reason must be that the existing map of the world has omitted some crucial component. This inference is not wholly irrational, since the history of science does indeed contain cases in which anomalous phenomena ultimately forced us to acknowledge new objects, new processes, or new theoretical levels. The problem is that this historical possibility is easily misused as though it were a universal template. As a result, every instance of empirical unfamiliarity is treated too quickly as evidence of an ontological gap, while the more fine-grained explanatory work in between is passed over.

There is also a more concealed source of ontological anxiety: people often confuse “not yet explainable for the moment” with “in principle not explainable within the current framework.” The former may reflect nothing more than model insufficiency, limited data, unclear conceptual boundaries, or a wrongly posed question; the latter would mean that the framework itself is deeply incapable of handling the phenomenon. Yet in actual discussion, these two are rarely distinguished with patience. The collapse becomes even more likely when the phenomena in question are highly unfamiliar, widely circulated in public discourse, or emotionally charged. The result is that a feeling of empirical mismatch rapidly becomes a feeling of ontological crisis, as though whenever we cannot explain

something at once, we must immediately decide anew what else there is in the world.

This mechanism is especially common in discussions of anomalous phenomena. Whether the case concerns poorly understood physical behavior, difficult-to-classify observational data, anomalous patterns in highly complex systems, or more controversial objects such as UAP and signs of unknown intelligence, such phenomena are often placed in a high-pressure discursive setting: if one does not immediately acknowledge a new ontology, one seems to be suppressing the phenomenon; if one remains restrained, one is suspected of merely defending the old framework. Theoretical judgment then no longer unfolds according to explanatory order, but is driven by a false binary: either admit revolution or deny the facts. Ontological anxiety here is not merely a theoretical reaction; it becomes an atmosphere of discussion.

From the standpoint of the earlier chapters, this anxiety once again reveals the same underlying structure. It closely resembles the thought, in debates over emergence, that if something is genuinely new, then it must be causally new; and it resembles the thought, in discussions of entanglement, that if there is strong correlation, then there must be some transmissive connection. Unknown phenomena trigger ontological anxiety not because empirical unfamiliarity naturally equals ontological insufficiency, but because theoretical culture has long harbored a compression impulse: whenever existing explanation ceases to run smoothly, the pressure is automatically discharged at the level of “there must be more in reality.” That is to say, empirical novelty is first felt as an ontological problem rather than first analyzed as a problem of how explanatory roles are assigned.

But this step is not necessary. A more mature theoretical response should first ask: where exactly is the problem here? Is it instability in observation and classification? Insufficient model resolution? Faulty variable-selection? Unclear

level-distinction? Has a structural problem been miswritten as a causal one? Or is it only after such possibilities have been gradually excluded that we truly have reason to suspect the current ontological picture of being too impoverished? What the rest of this chapter argues for is precisely this discipline of sequence. Unknown phenomena may indeed eventually motivate ontological revision, but what they first test should not be our courage in expanding the inventory of entities, but our ability to preserve explanatory calm and layered judgment under empirical pressure.

The provisional conclusion of this section is therefore this: unknown phenomena trigger ontological anxiety not because that anxiety is always justified, but because empirical unfamiliarity is easily translated too quickly into structural insufficiency in reality itself. To free ourselves from this automatic reaction, the first step is not simply to declare that no new ontology is needed, but to look back historically at how this anxiety has repeatedly formed and how it has often later been corrected. The next section follows this path by discussing the historical pattern linking empirical novelty and ontological inflation.

8.2 The Historical Pattern of Empirical Novelty and Ontological Inflation

If unknown phenomena so often trigger ontological anxiety, then the next obvious question is: what pattern has this anxiety displayed in the history of science and theory? From a broad historical perspective, there does indeed seem to be a recurring linkage between empirical novelty and ontological inflation. Whenever an existing theoretical framework is temporarily unable to integrate certain phenomena smoothly, people often quickly incline toward the view that the problem is not merely one of model-detail, but that reality's own inventory is not rich enough. New entities, new forces, new media, and new levels then emerge as the direct outlet for explanatory anxiety.

It is not difficult to see why this pattern recurs. Faced with unfamiliar phenomena, one of the most immediate and psychologically satisfying reactions is to convert the difficulty into an existential judgment: if the current picture cannot explain it, then there must be something in the world not yet entered into the ledger. The attraction of this move lies in the fact that it seems both to take the anomaly seriously and to give theory a clear target. By contrast, admitting that “we do not yet know at what level the problem lies” appears cognitively and narratively far less stable, and far less satisfying. Ontological inflation is thus not only a theoretical choice, but also a rapid way of settling uncertainty.

Yet the history of science also shows that although this reaction is common, it is not always reliable. In many cases, empirical difficulty is not ultimately resolved by adding new existents, but by reunderstanding the structure of the problem, revising the form of the model, clarifying variable-relations, redrawing level-boundaries, or abandoning certain explanatory demands that once appeared natural. The lesson of history is therefore not that new ontology is never needed, but that there is no direct route from empirical novelty to ontological expansion. In many cases, what truly requires change is not what exists in the world, but how we describe, organize, and explain what is already before us.

What deserves special emphasis here is a recurring absence of mediation. Empirical novelty does not usually jump directly to new ontology. It ought instead to pass through a series of intermediate judgments: is the observation stable? Is the classification reliable? Is the model adequate? Have distinct problems been miswritten into one another? Is this a case of local theoretical failure rather than collapse of the overall ontological picture? But in many historical cases these mediating judgments were compressed. As soon as the anomaly appeared, ontology was mobilized first. For that reason, ontological inflation often looks less like a truly forced ontological conclusion than like impatience with explanation.

Historically, this pattern usually takes three common forms.

The first is treating a theoretical blank as an existential blank. When certain phenomena cannot, for the moment, be explained, people easily conclude that since current theory cannot make sense of them, the current map of the world must have omitted the object itself. The problem with this line of thought is that it translates “we do not yet have a suitable account” too quickly into “there must be something additional in the world.” Yet what later happens is often not an increase in the list of existents, but a change in how theory understands relational structure, constraining conditions, or level-organization.

The second is treating descriptive difficulty as ontological division. Complex phenomena especially invite this reaction. Once a system exhibits stable patterns, coordinated behavior, or macroscopic novelty at a higher level, people readily feel that if lower-level language cannot directly and losslessly cover that level, then reality itself must split into two parts and a new ontological level must be acknowledged. But as the earlier chapters have repeatedly argued, the non-flattenability of descriptive levels does not automatically imply further ontological multiplication. Much of the temptation toward so-called ontological division actually comes from misjudging structural novelty and explanatory role.

The third is treating anomalous correlation as a new mechanism of connection. Chapter 7 on entanglement already illustrated this: once striking correlations can no longer be absorbed by familiar models, people readily turn to the thought that there must be some additional connection. This pattern is not limited to quantum discussions, but appears more generally whenever theory confronts anomalous coordination, distant correlation, or complex coupling. Once again, the problem is that the intermediate steps are skipped: before we have asked clearly whether what is exposed is a problem of transmission, structure, boundary, or the model’s own mode of decomposition, ontology is already being expanded.

What these historical patterns show is that ontological inflation is not always the mistake of being “too bold.” More often it is an error in the order of judgment.

The problem does not necessarily lie in whether new ontology is eventually introduced, but in whether it is introduced too early, before the empirical, model-theoretic, and explanatory levels have been adequately distinguished. A framework that lacks this sense of sequence will be forced, whenever novelty appears, to choose between “expand immediately” and “deny stubbornly,” and will find it difficult to form a genuinely robust middle ground.

This is precisely why the chapter introduces the concepts of explanatory discipline and ontological robustness. The historical pattern shows that a mature theoretical response does not consist either in never revising ontology or in declaring revolution at the first sign of pressure, but in being able to distinguish which empirical novelties genuinely end up requiring ontological revision and which primarily require model revision, problem-rewriting, redistribution of explanatory roles, or clarification of structural boundaries. Without this distinction, theory will readily mistake every difficulty for a shortage in the world’s inventory.

In terms of the book’s overall framework, this historical observation is fully homologous with the arguments of the earlier chapters. In the case of emergence, structural novelty was translated too quickly into new causal power. In the case of higher-level explanation, explanatory surplus was translated too quickly into additional causes. In the case of entanglement, pressure arising from joint structure was translated too quickly into nonlocal transmission. Now, in the more general context of unknown phenomena, empirical novelty is translated too quickly into new ontology. The common structure across all four cases is that before explanation has been treated in a properly layered way, the pressure is discharged directly at the strongest ontological level.

The conclusion of this section can therefore be summarized as follows: there is indeed a stable historical linkage between empirical novelty and ontological inflation, but this linkage itself cannot serve as a justification for ontological expansion. It reveals, rather, a recurrent theoretical reaction-pattern: when an

explanatory framework is not yet ready to distinguish at what level the problem lies, ontology is mobilized too early. The real lesson to be learned from history is not “never add new ontology,” but “do not let empirical unfamiliarity automatically leap into an ontological conclusion.”

This prepares the way for the next section. If the historical pattern shows that unknown phenomena are often first misjudged as ontological crises, then we must now ask more directly: what do they really challenge first? The next section presents the chapter’s central diagnosis: what unknown phenomena first challenge is not “what there is in the world,” but “how we allocate explanatory tasks.”

8.3 What Unknown Phenomena First Challenge

If the historical linkage between empirical novelty and ontological inflation does not show that every anomaly directly demands new existents, then we must now answer more explicitly: what is it that unknown phenomena really challenge first? The central judgment of this chapter is that what they first challenge is not our final ontological list of what there is in the world, but the way we allocate explanatory tasks when confronted with new phenomena. What they first expose, in other words, is usually not that the structure of reality itself exceeds the existing inventory, but that we do not yet clearly know what level the problem belongs to, what sort of account it requires, and which explanatory demands are themselves being made too quickly.

This point is easy to miss because, intuitively, an “ontological challenge” seems more direct than a challenge to the allocation of explanation. If some phenomenon exceeds the current framework, we naturally feel that perhaps the framework is missing some existent. By contrast, to admit that the problem first lies in our not yet knowing what question to ask can seem overly cautious, even evasive. But this caution is not retreat. It is a sign of theoretical maturity. When empirical unfamiliarity first appears, we often do not yet know whether the difficulty arises

from observational stability, model-fit, variable-selection, level-confusion, or genuinely from a deeper ontological insufficiency. If, before distinguishing these layers, we jump directly to an ontological conclusion, then the apparent boldness is really only an inversion of the proper order of judgment.

From the perspective of the earlier chapters, the allocation of explanatory tasks has a very concrete meaning. It requires us first to ask: what is lacking here? Is it a productive account of how the result occurs? A structural account of why a pattern remains stable? A permissibility-account of which possibilities are allowed and which excluded? Or a meta-level diagnosis of why accounts at different levels have been wrongly organized into a competitive relation? Only once these questions have been at least provisionally distinguished can we begin to judge whether the tension in the existing framework reflects explanatory insufficiency or whether the ontological inventory itself is genuinely too narrow and must be expanded.

The reason unknown phenomena first test explanatory allocation is that what they usually expose first is a blank place in the existing map of explanation, not immediately a new existent as such. What an anomalous phenomenon most often presents, at the outset, is the fact that we do not yet know how to place it, not the fact that it straightforwardly points to some new entity waiting to be acknowledged. Empirical novelty therefore creates, first of all, a problem of localization, and a problem of localization is not automatically a problem of ontological shortage. It may indicate that our model-classification is too coarse, that we have mistaken a structural problem for a causal problem, or that we have overextended an explanatory role previously thought sufficient.

Especially important here is the distinction between two different kinds of framework-insufficiency. One kind of insufficiency occurs at the level of explanatory organization: the objects, variables, and level-relations within the existing framework have not yet been properly arranged, and so the phenomenon appears difficult to incorporate. The other kind occurs at the ontological level: even after

explanatory roles, model-boundaries, and level-distinctions have been made as clear as possible, the phenomenon still demands a place that the current inventory cannot provide. Many historical misjudgments arise from declaring the second form of insufficiency before the first has been properly examined. What this chapter calls explanatory discipline is, first of all, a guard against precisely that leap.

From this standpoint, the importance of unknown phenomena lies not in the fact that they always announce that there is something new, but in the fact that they force a framework to ask itself: is your division of explanatory labor mature enough? Are your distinctions among problem-types too coarse? Have your level-distinctions forced accounts into competition where none was required? Has some default explanatory template caused you to translate the phenomenon too quickly into one specific type of problem? Unknown phenomena are, in other words, first of all a pressure-test for the framework. What they test is whether a framework can preserve order of judgment, maintain conceptual boundaries, and identify where it is actually getting stuck when empirical scope expands.

This shift in emphasis has an important consequence: it means that taking a phenomenon seriously and immediately expanding ontology are not the same thing. Discussion becomes polarized very quickly when those two are tied together. It then appears that if one does not at once acknowledge a new ontology, one is not taking the new phenomenon seriously; and conversely, that taking the phenomenon seriously requires willingness to rewrite the inventory of reality. It is precisely this false binary that the present book rejects. To take an unknown phenomenon seriously is not to propose new existents as quickly as possible, but to distinguish as rigorously as possible what exactly the phenomenon requires us to revise. If what it first requires is a reorganization of explanatory structure, then premature ontological expansion is actually a form of theoretical rashness.

For that reason, when this chapter claims that unknown phenomena first challenge the allocation of explanation, it is not minimizing empirical pressure. It is locating that pressure more accurately. A phenomenon may indeed eventually lead to ontological revision, but it rarely appears at the outset in such a pure and transparent form. Much more often it first manifests itself as instability in our classifications, pressure among our model-demands, misalignment among explanatory roles, and the reopening of questions once thought closed. If these intermediate signals are not taken seriously, ontological judgment loses the layered support it ought to have.

The conclusion of this section can therefore be stated as follows: what unknown phenomena first really challenge is not what else exists in the world, but how we identify and organize the explanatory tasks we already face. What they first test is a framework's power to discriminate levels of problem, its capacity to allocate explanatory roles, and its discipline in locating the source of theoretical pressure. Only after work at those levels has been progressively carried out can ontological revision cease to be an anxiety-driven rapid reaction and become a genuinely supported theoretical conclusion.

This provides the direct basis for the next section. If unknown phenomena first test the robustness of an explanatory framework rather than automatically requiring ontological expansion, then we need to distinguish more carefully two things that are often conflated: explanatory robustness and ontological stability. How a mature framework remains open without becoming uncontrolled, and restrained without becoming rigid, is exactly what the next section addresses.

8.4 Explanatory Robustness and Ontological Stability

If what unknown phenomena first test is not ontological inventory but the capacity to allocate explanatory tasks, then the next key question is this: in the face of empirical novelty, in what respects should a mature framework remain stable, and

in what respects should it allow revision? Unless this is clarified, the argument of the present chapter is easily misread as a form of conservatism, as though not rushing to add new ontology were merely a strategy for buying time for the old framework. To avoid this misunderstanding, we must distinguish two notions that are often run together: explanatory robustness and ontological stability.

By explanatory robustness I mean a framework's ability, when confronted with new phenomena, to maintain a clear division of explanatory roles and a clear sense of problem-levels. It does not mean that existing models need never be adjusted, nor that all conceptual boundaries have already been fixed once and for all. On the contrary, it means that even under empirical pressure, the framework can still distinguish which problems belong to productive explanation and which to structural explanation; which difficulties arise at the model-level and which at the classificatory level; which are merely noise caused by unstable observation and which genuinely press against deeper theoretical boundaries. A framework that, in the face of novelty, immediately loses this layered power of discrimination and compresses every difficulty into "there must be something more in the world" lacks explanatory robustness. Explanatory robustness is therefore, first of all, a stability of discernment.

By ontological stability, by contrast, I mean that a framework does not immediately expand its inventory of reality in response to every empirical unfamiliarity. This does not mean that ontology never changes. It means, rather, that revision of ontological commitment should occur with order, threshold, and layered support. A framework with ontological stability does not, when faced with unknown phenomena, discharge pressure at once into conclusions about new entities, new forces, or new levels. It first asks whether the difficulty arises from conceptual misplacement, model-insufficiency, level-confusion, improper carving of variables, or an overstrong explanatory demand. Only after such possibilities have been progressively examined does ontological revision become an

appropriately grounded next step rather than an initial reflex. Ontological stability, in short, is not refusal of change, but refusal of unlayered change.

These two notions must be distinguished because they are so easily set in false opposition. It often seems as though a framework that is ontologically stable must be explanatorily rigid, and that a framework open enough at the level of explanation must therefore be ontologically volatile. The present book argues for the opposite. Genuine ontological stability presupposes explanatory robustness. A framework is able to avoid immediate ontological inflation in the face of novelty precisely because it possesses enough explanatory discrimination to place the difficulty first at the right level. By contrast, the frameworks most likely to expand ontology rapidly in response to anomalous phenomena are often those that are least robust in their explanatory stratification: they translate every difficulty too quickly into the strongest existential question.

From this angle, ontological stability is not a passive defense of the status quo. It is not the view that the old framework ought to be protected by default, nor that the addition of new ontology should always be treated with suspicion. It is closer to a disciplined postponement of judgment: ontology should not be the first layer to bear pressure; rather, model revision, redistribution of explanatory roles, clarification of structural boundaries, and correction of classificatory schemes should do their work first. Only then does ontological change appear as something forced rather than induced. Stability of this sort is not a synonym for conservatism, but a sign of theoretical maturity.

Explanatory robustness also performs another function here: it prevents the seriousness of a phenomenon from being falsely tied to ontological radicalism. Discussion often polarizes quickly because these two are treated as the same stance. If one does not immediately acknowledge new entities, one seems not to have taken the new phenomenon seriously; if one remains ontologically restrained, one appears instinctively to be defending the old world-picture. The importance of

explanatory robustness lies precisely in breaking this association. A framework can take anomalous phenomena extremely seriously while still maintaining strong discipline with respect to ontological commitment. Indeed, only if it can do so does it show itself capable of distinguishing empirical pressure from ontological revolution.

It should also be emphasized that ontological stability does not mean that a framework remains motionless at every level. On the contrary, what often changes first in a mature framework under empirical novelty is its explanatory structure. It may revise problem-classification, redefine relations among levels, abandon certain overstrong explanatory demands, admit that previous boundary-distinctions were too coarse, or introduce clearer structural distinctions. Such changes are themselves theoretically significant even if they do not yet take the form of ontological expansion. Ontological stability is therefore not equivalent to theoretical stasis; it often presupposes active conceptual and explanatory work. A genuinely rigid framework is not one that is ontologically stable, but one that is neither willing to revise ontology nor capable of reorganizing explanation.

Seen in relation to the book as a whole, this point continues the earlier logic. In the case of emergence, the proposal was to rewrite the meaning of novelty before adding new causal powers. In the case of higher-level explanation, the proposal was to separate causal sufficiency from explanatory sufficiency before granting the higher level an extra ontological place. In the case of entanglement, the proposal was to reorganize structural inseparability and the demand for transmission before accepting nonlocal causation. Now, in the general setting of unknown phenomena, the same order appears once again: what should first be tested is the robustness of the explanatory framework, not the elasticity of the world's inventory.

The conclusion of this section can therefore be stated simply: explanatory robustness and ontological stability are not opposed; the former supports the latter. A mature framework avoids rushing into ontological inflation in the face of

empirical novelty precisely because it possesses sufficient explanatory stratification to place difficulties first in finer-grained theoretical locations. What is genuinely suspect is not ontological restraint, but allowing ontology to bear all the pressure before explanation has been properly layered.

This leads directly to the next section. If ontological stability does not mean refusing revision, but refusing ungraduated revision, then we must now say more explicitly why new phenomena do not automatically entail new ontology.

8.5 Why New Phenomena Do Not Automatically Entail New Ontology

If both explanatory robustness and ontological stability require us to preserve layered judgment in the face of unknown phenomena, then what must now be stated most explicitly is a conclusion that seems simple yet is often ignored: new phenomena do not automatically entail new ontology. This point matters not because ontological revision is never necessary, but because there is no direct and unmediated inferential chain from empirical novelty to an increase in what exists. Between the two lies an entire body of explanatory work, and it is precisely this work that is so often skipped too quickly under the pressure of theoretical anxiety.

The first thing to recognize is that a phenomenon may appear “new” for many different reasons. It may simply never have been observed before. It may become visible only because observational conditions have changed. It may have long been placed in the wrong category because the existing classificatory scheme was too coarse. Or an existing model may cease to be sufficiently fine-grained at a certain scale, level, or boundary-condition. All of these possibilities can give a phenomenon the character of empirical novelty, but none of them by itself means that the inventory of reality must be expanded. “New phenomenon” means, first of all, only this: within our current descriptive and explanatory framework, something has appeared that has not yet been securely placed. It does not yet directly tell us

that the phenomenon must correspond to some new existent, new force, or new ontological level.

This matters because many theoretical discussions misread “not yet placed” as “a missing item in the ontological map.” In fact, failure of placement can occur at many nearer levels. The difficulty may lie in the carving of variables, in the way system-boundaries have been drawn, in a conflation of explanatory roles, in an overextension of the domain of a model, or in the fact that we are asking a question that ought not to have been asked in that form. If all these intermediate levels remain unexamined and the pressure is pushed directly to the ontological level, then the conclusion of “new ontology” is not forced by the phenomenon itself, but prematurely manufactured by a disordered sequence of judgment.

Logically speaking, what matters here is the distinction between two different inferences. The first is:

“A phenomenon has appeared that our present framework does not explain well; therefore the framework requires revision somewhere.”

This is a reasonable and necessary inference.

The second is:

“A phenomenon has appeared that our present framework does not explain well; therefore there must exist in the world something not previously included in our ontological map.”

This second step is not impossible, but it does not follow directly from the first. It requires further argument, and that argument cannot legitimately bypass the intervening work of explanatory stratification.

This also shows why the thought that new phenomena automatically imply new ontology is, at bottom, an explanatory shortcut. It allows us to skip the more difficult intermediate judgments and to package theoretical difficulty as existential discovery. The shortcut is attractive because it looks at once radical and clear: if the problem is this large, then the answer too must be large. But theoretical clarity is

not the same as theoretical legitimacy. In many cases, the genuinely difficult task is not to acknowledge a new ontology, but to show why all weaker levels of revision have already proved insufficient. If that demonstration has not been completed, ontological expansion remains an anxiety-driven premature reaction.

It should also be noted that ontological expansion carries a real cost. To add an entity, a force, or a level is not merely to add another noun; it is to rewrite the load-bearing structure of the entire theoretical picture. It affects how we individuate objects, organize explanation, understand level-relations, and coordinate the phenomenon with other domains. Ontological revision should therefore bear a heavier burden of proof than model-revision, classificatory adjustment, or redistribution of explanatory roles. A framework that places the heaviest revision first whenever empirical novelty appears is not merely being “open-minded”; it lacks methodological weight.

The anomaly of Mercury’s perihelion illustrates this discipline especially well. Its significance lies precisely in showing that empirical novelty does not itself determine the level at which revision must occur. Faced with anomaly, one may attempt to add a new object, or revise a background model, reorganize theoretical structure, or reinterpret constraining conditions. Historically, the postulation of Vulcan represented one route of object-expansion, whereas the success of General Relativity showed that some anomalies require not the addition of a new member to reality’s inventory, but a transformation in explanatory structure itself. This case makes clear that ontological discipline does not mean refusing change, but insisting on order among the levels at which change is considered.

The conclusion of this section can therefore be stated plainly: a new phenomenon is not yet a new existent. What it gives us immediately is pressure, not placement; difficulty, not decision; a demand for reorganization, not yet a warrant for inflation. If ontology is to remain a serious theoretical commitment

rather than a receptacle for unresolved explanatory anxiety, then the path from anomaly to ontology must always pass through disciplined explanatory work.

This also prepares the way for the next section. Once we admit that new phenomena are better understood as tests of explanatory discipline than as automatic evidence for new ontology, the most appropriate next step is to place this analysis in precisely those settings most likely to provoke public and theoretical anxiety: anomalous phenomena, UAP, and discussions of non-extraterrestrial intelligence.

8.6 Anomalous Phenomena as a Stress Test for Explanatory Discipline

If the argument of the previous sections has remained somewhat general, then the sharpest test of this framework's resilience lies precisely in those objects most likely to provoke collective anxiety and theoretical impulsiveness: anomalous phenomena, UAP, and various empirical clues and conjectures concerning non-extraterrestrial intelligence. These cases are chosen not because they are inherently more important than other scientific problems, but because they compress with particular clarity the very mechanism discussed throughout this chapter. They function almost as amplifiers of ontological anxiety. As soon as some observation, behavior, or pattern resists rapid classification, people easily slide from "the explanation is not yet complete" to "our picture of reality is no longer sufficient," and then further to "there must exist new entities, new technological agents, or new levels." For precisely that reason, such cases provide an especially apt test of whether explanatory discipline is strong enough.

What makes these cases philosophically interesting is not only that the evidence is often unstable, incomplete, or disputed, but that the pressure they exert is layered all at once. There is pressure at the observational level: was something really seen, measured, or recorded? There is pressure at the

classificatory level: if something was indeed observed, under what kind of object or process should it be placed? There is pressure at the model level: do existing explanatory models possess enough resolution to absorb the data? There is pressure at the structural level: have our inherited divisions among object, agency, intelligence, technology, and environment become too coarse? Only after these layers have been worked through does the ontological question properly arise: must we really admit a new kind of existent or a new domain of reality? In public discussion, however, these layers are typically compressed into one another and translated almost immediately into the strongest existential conclusion.

This is exactly why anomalous phenomena are so useful as a methodological stress test. They reveal whether a framework can preserve the order of judgment under high pressure. Can it separate uncertainty in observation from insufficiency in theory? Can it distinguish classificatory instability from ontological deficit? Can it resist the temptation to turn every unresolved pattern into proof of a new object? And at the same time, can it avoid the opposite reflex of dismissing all novelty simply because the evidence is uncomfortable or difficult to place? A framework that cannot maintain this balance under pressure does not merely have a view about anomalies; it reveals something deeper about its own explanatory maturity.

In this setting, the distinction made earlier between explanatory robustness and ontological stability becomes especially important. Explanatory robustness means that even when confronted with highly charged and unfamiliar cases, a framework does not immediately lose the ability to distinguish what kind of problem it faces. Ontological stability means that it does not discharge all of that pressure at once into new commitments about what exists. If a framework lacks the former, it will collapse every anomaly into the same type of crisis. If it lacks the latter, it will rush to inflate reality's inventory whenever empirical novelty appears. These cases therefore expose, in unusually stark form, whether a theory's restraint

is genuine discipline or merely defensive reflex, and whether its openness is genuine sensitivity or merely conceptual volatility.

This also helps explain why debates over UAP and related phenomena so quickly become polarized. One side often behaves as though taking the phenomena seriously requires immediate willingness to revise ontology; the other behaves as though resisting ontological revision requires treating the phenomena as negligible, delusional, or epistemically beneath notice. The present framework rejects both moves. Seriousness and inflation are not the same thing, just as restraint and dismissal are not the same thing. A disciplined theory must be able to acknowledge pressure without surrendering sequence, and admit the possibility of revision without declaring its object in advance. If it cannot do that, then what appears as “courage” is often just haste, while what appears as “skepticism” is often only conceptual rigidity.

The discussion of anomalous phenomena also makes something else clear: what public imagination often wants from such cases is not merely explanation but drama. The attraction of saying “this proves there is something radically new in the world” lies partly in the fact that it converts uncertainty into a striking verdict. Yet theoretical maturity rarely proceeds through the most dramatic inference available. It proceeds by preserving distinctions precisely where excitement encourages their collapse. In that sense, anomalous cases are not only tests of what a framework thinks the world contains; they are tests of whether it can keep thought proportionate to evidence.

There is another reason these cases matter. Because they are so often discussed in the borderland between science, media, political secrecy, and speculative culture, they make it especially difficult to preserve the distinction between evidential pressure and narrative overloading. This makes explanatory discipline all the more important. A mature framework must be able to say, at once: yes, there may be something here that places real pressure on our existing

conceptual and explanatory resources; and no, that does not yet authorize unrestricted escalation to exotic ontology. To sustain both sides of that judgment at once is not weakness but strength.

The conclusion of this section can therefore be summarized as follows: the real methodological significance of anomalous phenomena, UAP, and related objects does not lie in the fact that they automatically license some new ontology, but in the fact that they most clearly reveal whether a framework possesses explanatory discipline at all. They are not shortcuts to ontological conclusion, but stress tests for explanatory maturity. If a framework can preserve layered judgment, defer cross-level conclusions, and still acknowledge empirical pressure in these high-pressure cases, then its ontological robustness is genuinely worthy of trust.

This also leads naturally to the next section. If ontological robustness in such cases does not mean conservative refusal, but rather a higher-order form of explanatory maturity, then we must now say more positively what that robustness itself is good for.

8.7 The Positive Significance of Ontological Robustness

Even at this point, one lingering misunderstanding may remain: if this chapter repeatedly emphasizes ontological stability, explanatory discipline, and layered judgment, does it ultimately reduce theoretical posture to a form of cautious conservatism? In other words, is the claim that new phenomena should not automatically entail new ontology really just a way of buying more life for the old framework? If the chapter cannot answer this directly, then the preceding argument may too easily be read as a merely defensive strategy: commit less, judge more slowly, revise later. To avoid that misunderstanding, it must be stated clearly that ontological robustness is not the opposite of theoretical courage. It is a higher-order theoretical capacity. Its value lies not in doing less, but in deciding more accurately what ought to be done under empirical pressure.

First, the positive significance of ontological robustness lies in the fact that it allows a framework to endure novelty without immediately falling into disorder. What most easily happens in the face of empirical unfamiliarity is not total neglect, but overreaction. Once a phenomenon cannot be smoothly absorbed into familiar categories, thought tends to oscillate between suppression and inflation: either it is dismissed too quickly as noise, error, illusion, or nonsense, or it is elevated too quickly into evidence that the existing map of reality has collapsed. Ontological robustness interrupts this oscillation. It allows a framework to hold novelty in view without either shrinking it away or turning it at once into revolution.

Second, ontological robustness preserves proportionality. Not every anomaly contains the same amount of information, and not every theoretical difficulty has the same depth. A robust framework does not merely “stay calm”; it preserves a sense of scale. It can distinguish between local model-failure and global framework-failure, between classificatory friction and ontological rupture, between pressure that demands more careful explanation and pressure that genuinely warrants revision of what exists. This capacity for proportion is itself a major theoretical virtue, because without it theory loses the ability to allocate its efforts with any precision.

Third, ontological robustness helps preserve the informational content of phenomena. If one rushes too quickly either to dismiss or to inflate, one often overwrites the phenomenon before understanding it. Premature dismissal treats the unfamiliar as though it contained nothing worth reorganizing theory around; premature inflation treats it as though its meaning were already decided in ontological terms. In both cases, the phenomenon is effectively silenced. A robust framework, by contrast, leaves room for phenomena to continue informing inquiry before their final theoretical placement has been fixed. That makes it more, not less, responsive to what empirical novelty may eventually teach.

Fourth, ontological robustness functions as a restraint on pseudo-revolution. One of the recurring dangers in theory is not simply conservatism, but false radicalism: the temptation to dramatize a local or intermediate difficulty as though it already demanded a total rewriting of reality. Such gestures may feel bold, but they are often methodologically thin. Genuine theoretical transformation occurs when weaker and nearer revisions have been seriously tested and found insufficient, not when the strongest conclusion is simply the most exciting one available. Ontological robustness therefore protects theory from mistaking escalation for depth.

Fifth, ontological robustness gives a framework control over its own inferential tempo. A mature theory is not merely one that reaches important conclusions, but one that knows when a conclusion has actually become warranted. Timing matters in judgment. If a theory revises too early, it destabilizes itself without sufficient reason; if it revises too late, it becomes inert and evasive. Ontological robustness is the ability to keep that timing under control. In this sense, it is a form of reflective self-governance within theory itself.

This also clarifies why ontological robustness should be understood as an active virtue rather than a passive lack of adventurousness. It is an achievement of disciplined discrimination under pressure. It requires maintaining several distinctions at once, preserving conceptual order where public or theoretical excitement would rather collapse it, and allowing empirical novelty to remain serious without permitting it to dictate its own strongest interpretation prematurely. None of this is easy. It asks more of a framework than immediate escalation does, not less.

Seen in this way, ontological robustness is also what allows a framework to remain genuinely open. A theory that reacts to every unfamiliarity by immediate ontological expansion is not truly open; it is unstable. Yet a theory that refuses all pressure in advance is not open either; it is rigid. Robustness lies between these

failures, though not as a compromise. It is the capacity to remain revisable without becoming volatile, and restrained without becoming closed. That is a positive theoretical power, not a defensive posture.

The broader lesson, then, is that ontological robustness is itself a significant theoretical virtue. It manifests as resilience under pressure, a sense of proportion, preservation of informational content, resistance to pseudo-revolution, and control over the tempo of inference. A framework is mature not because it never changes, but because it knows how change ought to occur. In the face of unknown phenomena, what makes it trustworthy is not that it declares a new world most quickly, but that it best preserves the order of judgment under pressure.

This prepares the way for the chapter's final summary. By now, we have shown not only why unknown phenomena do not automatically entail new ontology, but also why maintaining ontological robustness is itself a theoretical achievement.

8.8 Chapter Summary: From Unknown Phenomena to Explanatory Maturity

This chapter began from a tendency that repeatedly appears in scientific practice, philosophical debate, and public culture alike: whenever anomalous phenomena, unknown objects, or empirical novelties arise that cannot be rapidly classified, people easily slide from "our present explanation is not yet enough" to "our present ontology is not yet enough." This slippage deserves special attention not because it is always mistaken, but because it happens so readily, and so prematurely. The central task of this chapter has been to dismantle that rapid passage from empirical pressure to ontological anxiety and to show why the explanatory work that ought to lie in between cannot be skipped.

First, the chapter argued that unknown phenomena so often trigger ontological anxiety because empirical unfamiliarity easily unsettles the usually tacit sense of fit between theoretical framework and world. As soon as a phenomenon

can no longer be stably placed within our existing concepts, models, or classifications, it begins to feel as though the world itself were exceeding its current map. Yet this feeling, the chapter insisted, should not be treated as a direct ontological verdict. What it reveals first is the fragility of our explanatory organization.

Second, the chapter distinguished what unknown phenomena first challenge. They do not first challenge our final list of what exists, but our ability to identify and organize explanatory tasks. What they first test is whether we can still distinguish observational problems from classificatory problems, model-insufficiency from structural insufficiency, and explanatory difficulty from ontological deficit. In that sense, unknown phenomena expose not simply what reality may contain, but whether our framework retains enough discipline to locate where the real pressure lies.

Third, the chapter distinguished explanatory robustness from ontological stability. Explanatory robustness is a framework's capacity to preserve clear role-distinctions and problem-levels under empirical pressure; ontological stability is its capacity to avoid immediate and unlayered expansion of what it takes to exist. These two are not opposed. On the contrary, explanatory robustness is what supports ontological stability. Only a framework capable of fine-grained explanatory discrimination can refrain from discharging every anomaly directly into ontological inflation.

Fourth, the chapter argued explicitly that new phenomena do not automatically entail new ontology. There is no direct inferential path from empirical novelty to ontological increase. Between the two lies an entire field of explanatory work: improved observation, revision of classification, refinement of models, clarification of structural boundaries, and reorganization of explanatory demands. If this intervening work is skipped, then ontological expansion is not forced by the

phenomenon itself, but prematurely manufactured by a disordered sequence of judgment.

Fifth, the chapter used anomalous phenomena, UAP, and related cases as high-pressure examples that make these issues especially vivid. Such cases matter philosophically not because they automatically prove the existence of exotic realities, but because they most clearly test whether a framework can preserve explanatory discipline under pressure. They reveal whether it can acknowledge empirical weight without collapsing into ontological haste, and whether it can remain open to revision without becoming conceptually unstable.

Sixth, the chapter argued that ontological robustness has a positive significance of its own. It is not merely a conservative hesitation to revise, but a theoretical capacity: the ability to bear novelty without disorder, preserve proportionality, retain the informational content of phenomena, resist pseudo-revolution, and govern the timing of inference. A mature framework does not deserve trust because it changes least, but because it knows how change should proceed.

The overall conclusion of the chapter can therefore be stated in compressed form: unknown phenomena are not automatic evidence that reality contains more entities than we had thought; they are, first of all, tests of explanatory maturity. They pressure us to determine whether we have properly distinguished roles, levels, models, structures, and ontological commitments. Only when that work has been carried out step by step can ontological revision become a genuinely warranted conclusion rather than an anxious reaction to empirical unfamiliarity.

This also prepares the transition to Chapter 9. If the earlier chapters have shown how many philosophical disputes are generated by the misallocation of explanatory roles, and if the present chapter has shown why unknown phenomena should not immediately trigger ontological inflation, then the next step is to ask what kind of ontology can preserve the gains of this explanatory reorganization

without falling back into either flattening reductionism or unrestricted expansion. Chapter 9 takes up precisely that task.

Chapter 9

Restrained Layered Structural Realism: From Explanatory Roles to Ontological Levels (with Discussion of “Radiance-Will Ontology”)

If the work of the preceding chapters has shown that many philosophical disputes do not first arise from a lack of empirical data, but from the systematic misallocation of explanatory roles, then a further question can no longer be avoided. If causal production, structural constraint, and permissibility are indeed irreducible explanatory roles, and if higher-level facts, stable patterns, and structural boundaries therefore retain legitimate explanatory standing, then what follows for ontology? Are these distinctions merely different ways of speaking, useful only at the level of method and description, or do they also require us to say something about how reality itself is organized?

This question cannot simply be set aside. If one only argues that higher-level explanation is legitimate, that emergence need not mean new causal powers, that entanglement need not force mysterious transmission, and that unknown phenomena should not automatically trigger ontological inflation, then one may indeed have achieved an important deflationary result. But such a result remains incomplete. For once explanatory legitimacy has been secured, the further pressure naturally arises: what, then, is the ontological standing of those structures, patterns, levels, and boundaries that explanation cannot do without? If they are not to be reduced away, and if they are not to be inflated into new substances or forces, then some more disciplined ontological account is needed.

What this chapter attempts, therefore, is not a return to the old reflex of ontological multiplication, but a restrained step forward. Its aim is to articulate an ontology that can acknowledge the reality of layered structure without collapsing

into either flattening reductionism or unrestrained proliferation. If the earlier chapters were chiefly occupied with freeing explanation from the monopoly of a single task, the present chapter asks what picture of reality becomes possible once that explanatory space has been reopened.

The chapter proceeds by arguing that explanatory differentiation does not by itself yield a bloated ontology, but it does create pressure for a more refined one. The point is not that every explanatory role corresponds to a separate realm of being. Nor is the point that every useful distinction deserves ontological elevation. Rather, the task is to identify under what conditions explanatory indispensability supports a restrained form of ontological recognition. To do that, the chapter develops what may be called a layered structural realism: a view according to which reality is not exhausted by a flat inventory of basic items, yet also does not require the postulation of freely proliferating higher entities. What deserves recognition, on this view, are stable structural features, organization-patterns, and level-sensitive facts that are both dependent on underlying realization and nevertheless not thereby eliminated.

In this way, the transition from explanation to ontology must be handled carefully. If it is too weak, then the earlier defense of higher-level facts and structural explanation risks collapsing back into instrumentalism. If it is too strong, then the framework simply reproduces the inflationary tendencies it has been warning against all along. The challenge, then, is to move from explanatory plurality to ontological articulation without turning every explanatory distinction into a new ontological addition.

The present chapter argues that this can be done by treating ontology not as a mere inventory of independent existents, but as a structured account of what kinds of organization, stability, and dependence deserve ontological acknowledgment. Under such a view, the reality of higher-level facts does not consist in their being extra substances layered on top of the physical, but in their being stable structural

realities whose explanatory role cannot be replaced without theoretical loss. Ontological recognition, in other words, should track structured indispensability rather than mere numerical addition to the stock of beings.

This shift also has an important methodological consequence. It means that the question “what is real?” should not always be posed in the crude form “what extra things must be added?” Sometimes the more appropriate question is: what forms of organization, constraint, and patterned stability does reality itself sustain? Once the question is reframed in that way, the opposition between reductionist austerity and metaphysical inflation begins to loosen. One can be ontologically restrained without being flattening; one can acknowledge higher-level reality without treating it as an autonomous second world.

The goal of this chapter, then, is to give conceptual shape to that middle path. It seeks to show how explanatory achievements concerning higher-level facts, structural relations, permissibility conditions, and stable organization can be carried into ontology in a disciplined way. If successful, this will provide a way of understanding reality as layered, structured, and differentially organized, while remaining resistant to the temptation to convert every theoretical pressure into a demand for new entities, new forces, or new domains of being.

9.1 Why This Chapter Must Turn to Ontology

At this point, one might still ask whether the book really needs to make an ontological turn at all. After all, if the earlier chapters have already shown that many philosophical pressures arise from the compression of explanation into a single production-oriented framework, why not stop there? Why not simply remain at the level of explanatory clarification and methodological discipline, leaving ontology aside as an unnecessary further risk? This is a serious question, because the entire argument of the book has, up to now, been marked by suspicion toward premature

ontological escalation. If the book now turns toward ontology, does it not risk undoing its own achievement?

The answer is that the turn is necessary precisely because the earlier chapters have already done more than merely criticize bad inferences. They have preserved, across several domains, the legitimacy of distinctions that cannot be treated as merely verbal or pragmatic conveniences. If higher-level facts are explanatorily irreducible in certain respects, if structural boundaries genuinely matter, if permissibility conditions help determine what patterns can stably exist, and if stable forms of organization repeatedly do work that cannot be replaced by lower-level production stories alone, then these are not merely features of our descriptions. They begin to look like features of reality's own organization. To stop entirely at methodology would leave this achievement suspended in an unstable way.

That instability matters because without some ontological follow-through, the earlier arguments remain vulnerable to collapse back into weaker interpretations. One could always say: yes, these distinctions are useful; yes, they help us organize inquiry; yes, they clarify explanatory labor; but none of that shows that anything corresponding to them is real. If that move is left unanswered, then the structural facts and higher-level realities defended throughout the book risk being reduced once more to epistemic convenience. The book would then have protected them methodologically while failing to secure them ontologically.

Seen in this light, the move toward ontology is not a betrayal of earlier caution, but a continuation of it under more demanding conditions. The earlier chapters warned against ontological inflation because they showed how too many theoretical pressures are prematurely rewritten as demands for new entities, new forces, or new levels. But caution of that kind cannot remain merely negative. If one rejects inflation without offering any more adequate picture of reality, then caution itself becomes fragile. It risks turning into a purely deflationary habit, unable to

explain why some explanatory achievements deserve more than instrumental status.

For that reason, this chapter turns to ontology not because the book now wishes to encourage rapid promotion of higher-level facts into new ontological items, but because after resisting those rapid promotions, it still needs a language strong enough to say what kind of reality structural facts possess. If the previous chapters secured a place for them in explanation, the present chapter must secure a stable place for them in our picture of reality.

9.2 From Explanatory Roles to Ontological Levels

To move from explanatory roles to ontological levels, two symmetrical errors must first be avoided. The first error is to translate explanatory plurality directly into ontological multiplication, as though once we distinguish causal production, structural constraint, and permissibility, reality itself must therefore contain three parallel classes of entities or three separate domains of being. That inference is too quick and too crude. Explanatory roles are first of all distinctions among kinds of question and explanatory function, whereas ontological levels concern judgments about which organizational differences in reality possess stable standing. The former answers: what are we explaining? The latter answers: what sorts of structural differences in the world deserve to be acknowledged as real?

The second error is the opposite one: to treat the differentiation of explanatory roles as nothing more than a cognitive arrangement, as though these were merely several descriptive strategies adopted by observers for convenience. If that were so, then the higher-level facts, stable patterns, and structural boundaries preserved in the earlier chapters would once again slide back into observer-relative utility. But the whole point of defending structural explanation has been that it tracks not arbitrary perspectives, but recurring differences in how reality is organized. If a form of explanation remains effective across multiple contexts, repeatedly

determines which changes are excluded, which patterns are sustained, and which differences count as relevant at a given level, then we have reason to say that what it grasps is not merely a habit of language, but some structural fact of reality itself.

For that reason, the transition this chapter undertakes can only be a restrained form of ontologization. What is being ontologized here is not some extra thing, but additional modes of organization; not a new inventory, but stable structures; not a second set of beings detached from lower-level realization, but different structural resolutions of the same reality at different organizational scales. The move from explanatory roles to ontological levels is not a leap from language to mysterious objects, but a transition from explanatory function to the acknowledgment of structural fact. Only if this is preserved can the chapter avoid both dissolving the earlier results and betraying the theoretical restraint maintained throughout the book.

9.3 Basic Principles of a Layered Ontology

To infer ontological recognition from explanatory differentiation is not an automatic step. The fact that an explanatory role is useful within theory does not by itself mean that it corresponds to an independent structure in reality. Otherwise any effective model, convenient classification, or cognitive compression could be over-ontologized. What the book needs instead is a more restricted inference: only when some explanatory role satisfies conditions such as stability, cross-context indispensability, counterfactual constraint, and realizational support does it become eligible to underwrite ontological recognition.

If this restrained ontologization is legitimate, then the principles it depends on must now be stated. The first is the principle of dependence without disappearance. Higher-level facts, structural boundaries, and permissibility relations all depend upon lower-level realization; without underlying materials, local processes, and concrete implementation paths, they could not appear. But

this dependence does not amount to cancelability. A structural form may depend ontologically on lower-level reality while still remaining explanatorily indispensable and possessing stable boundaries and pattern-identity in the world. Dependence without disappearance is the core of the book's resistance to flattening reductionism.

The second is the principle of structure before inflation. When theory encounters higher-level patterns, stable organization, or strongly connected structures, the first question should be whether there is some structural fact that needs to be acknowledged, not whether we must immediately add new entities, new levels, or new causal powers. The first step of ontology should not be addition, but reorganization. In many cases, what theoretical pressure calls for is not expansion of inventory, but recognition that reality is already constituted across multiple scales of organization.

The third is the principle of non-competitive coexistence. The fact that facts at different levels may be equally real does not mean that they compete over which is "more fundamental" or "more genuinely real." Lower-level realization facts and higher-level structural facts can both hold because they grasp different organizational aspects of the same reality. Layering is not territorial division; difference is not hostility; levels are not a redistribution of causal sovereignty. Unless this is acknowledged, higher-level reality will always be forced back into a demand for extra causal power.

The fourth is the principle of threshold ontologization. Not every theoretically useful distinction deserves ontological standing. What merits recognition as a level-fact must satisfy at least several conditions: it must appear stably across multiple contexts; it must have relatively clear boundaries and criteria of relevant difference; it must be sustainably realized across multiple implementations; and it must not be a merely optional rhetorical aid in explanation. By imposing this threshold, the chapter rejects both flattening and indiscriminate multiplication,

preventing layered ontology from becoming a mechanism for unchecked proliferation of terms.

9.4 Possibility, Manifestation, and Structural Boundaries

To give this layered ontology a clearer internal framework, we must also recast three concepts that have repeatedly appeared earlier in the book but have not yet been systematically ontologized: possibility, manifestation, and structural boundaries.

First, possibility in this book is not some empty modal repository floating free of reality. It refers to which states, trajectories, patterns, and outcomes are open to a given system under its concrete organization, boundary conditions, and realizational constraints. In other words, possibility is not an abstract set of alternatives hovering prior to reality; it is a dimension of reality's own structure. A system has a certain space of possibilities not because every logically imaginable form is waiting to be realized, but because the way that system is organized opens paths for some forms while closing off others.

Second, manifestation is not merely the neutral claim that "something has occurred." More precisely, it refers to the stable actualization of certain forms within the space of possibilities organized by given structural boundaries, such that they become identifiable, maintainable, and explainable objects. Manifestation depends both on causal production and on structural boundaries. Without concrete processes, forms do not enter reality; without boundary-organization, forms cannot stably become patterns of a system. It is at this point that manifestation links causal production, structural constraint, and permissibility together.

Finally, structural boundaries are not merely negative limitations. They do exclude certain impossible states, but more importantly, they positively shape what can stably manifest. Without boundaries, there are no stable forms; without

constraints, the space of possibilities cannot be organized; without organization, manifestation cannot be sustained. Boundaries are therefore not signs of reality's poverty, but conditions of its having form at all. Once this is clearly seen, it becomes possible to understand why higher-level facts are not accidental appearances: they are continuously carried by some structural boundary. Likewise, systems exhibit layeredness because boundaries at different scales of organization allow different forms to manifest.

9.5 Restrained Layered Structural Realism and "Radiance-Will Ontology": Proposal and Core Definition

On this basis, the book names its position 'restrained layered structural realism'; in a more synthetic formulation, it may also be called 'Radiance-Will Ontology.' If the earlier chapters have shown that reality cannot be understood solely in terms of causal production among objects, but must also be grasped through structural boundaries, spaces of permissibility, and stable forms of manifestation, then we need an ontological language capable of expressing these dimensions together. This secondary name is not meant to add a mystifying label beyond existing layered ontology, but to mark a more specific picture of reality: reality is not a passive accumulation of objects, but a layered process that organizes possibilities through boundaries, maintains forms through structure, and presents stable manifestation at different scales of organization.

What is meant here by "radiance" is not physical light, nor an ornamental poetic flourish. It refers to the stable manifestation that reality's structure acquires at a given scale of organization. A form has "radiance" not because it is endowed with some mysterious aura, but because it becomes an identifiable, maintainable, and explainable object within structural boundaries. Radiance emphasizes the layeredness, stability, and structural visibility of manifestation. It indicates that

reality exists not only as an itemized arrangement of lower-level objects, but also as the appearance of patterns at different scales of organization.

What is meant by “will,” meanwhile, does not imply that reality possesses subject-like psychology, nor that the world makes choices in some personified sense. It refers to the impersonal directedness exhibited by structure under given boundary conditions: some forms are opened, some paths excluded, some patterns maintained, and some changes systematically filtered out. Will here is not subjective desire, but the directedness of structural organization, the selectivity of boundary-maintenance, and the orientation through which forms persist. What it expresses is not what the world “wants,” but how the world, through its own conditions of organization, allows certain forms to manifest, endure, and retain reality-status.

Accordingly, so-called 'Radiance-Will Ontology' is not a new substance ontology, but a name for an ontology of how reality manifests, organizes, and stratifies itself. It holds that reality's basic character should not be understood solely as causal production among objects, but also as the stable manifestation of structural forms and the impersonal directedness of organizing boundaries. The former constitutes the way reality appears at different scales; the latter constitutes how reality opens, limits, and sustains possible forms. If layered ontology tells us how reality may be acknowledged, Radiance-Will Ontology goes further and tells us how reality organizes its own manifestation.

9.6 The Layered Structure of Restrained Layered Structural Realism

If Radiance-Will Ontology is to become a genuine layered ontology rather than a loose label for structural manifestation and organizational directionality, then it must further explain at what levels reality is organized, how these levels connect, and why such layering does not slip back into ontological inflation. In response to

that need, the book introduces a six-level architecture. This architecture is not an added surcharge on the world, nor a mechanical translation of explanatory distinctions into six classes of entity. More accurately, it is a layered expression of several ontological functions in reality that cannot be reduced into one another. Its purpose is to prevent level-confusion, and to keep manifestation, directedness, possibility-structure, empirical structure, and the logic of structural evolution from being compressed into one and the same type of fact.

The first level may be called L0: the unmanifest base layer. This is not some domain of objects within the empirical world, nor a directly observable “primordial state.” It marks the ontological presupposition prior to all differentiation and manifestation. Put in the most restrained way, it corresponds to an undifferentiated background in which structural difference, directedness, and boundary-specification have not yet taken shape. Its significance is not to offer a story about the origin of the cosmos, but to indicate that manifestation does not hang unsupported in midair. Layered structures in reality always arise against some pre-differentiated unmanifest background. Introducing this layer is meant to reserve an ontological place for the possibility of differentiation itself, not to add a mysterious starting point.

The second level is L1: the will layer. “Will” here does not mean personalized subjectivity or teleological consciousness; it means that once differentiation occurs, its difference does not immediately vanish, but acquires persistence. In other words, the will layer expresses the persistence of differentiation. If the unmanifest base layer lacks identifiable structure, then the will layer indicates that structural differences are no longer fleeting disturbances, but begin to hold themselves in place. It is in this sense that the book treats “will” as the earliest form of impersonal directedness: not a power that actively creates difference, but the structural condition under which difference, once arisen, can continue to be borne and maintained.

The third level is L2: the radiance layer. If the will layer marks the persistence of differentiation, then the radiance layer marks the fact that this persistent differentiation begins to acquire directional manifestation. "Radiance" here is not physical light and not an æsthetic metaphor, but the level at which enduring difference appears in reality as recognizable direction, traceable development, and visible structural contour. If will answers the question "why did difference not disappear?", radiance answers "how does difference begin to manifest as some directed form of reality?" The radiance layer can therefore be understood as the level at which manifestness genuinely begins to form: it advances mere preserved difference into a process of manifestation already bearing direction and contour.

The fourth level is L3: the topological layer. What this layer handles is not already manifested empirical structure, but the possibility-structure through which such manifestations can unfold along certain paths, be permitted within certain boundaries, and be sustained through certain continuities. In terms of the book's earlier vocabulary, the topological layer comes closest to the ontological convergence-point of structural constraint and permissibility. What it determines is not any specific result, but the continuity, boundedness, and traversability of a result-space; not how a particular event occurs, but which paths are open to a system and which are systematically excluded. In this sense, the topological layer is neither the mathematical model itself nor the empirical fact itself, but the level at which reality's space of possibilities is structurally organized.

The fifth level is L4: the manifested-structure layer. This is the level of reality most familiar in the ordinary sense: physical structures, biological organization, psychic order, social forms, and every other manifested reality that can enter experience, observation, experiment, and theoretical description. Its importance lies not only in the fact that it is the principal field of operation for empirical science, but in the fact that the functions of the preceding layers achieve recognizable structural results here. Causal production is chiefly tracked at this

level in the form of concrete processes and realization-paths; higher-level facts also chiefly appear here in the form of stable patterns, functional forms, and organizational structures. The manifested-structure layer is therefore not simply the “last layer,” but the site at which the ontological functions of the earlier layers make their integrated appearance in reality.

The sixth level is L5: the layer of structural logic and convergence. If the first five layers explain how differentiation is preserved, acquires direction, is topologically organized, and manifests as empirical structure, then one further point remains: how structure, amid instability, deviation, tension, and reorganization, tends toward certain locally stable forms. The layer of structural logic and convergence expresses precisely this ontological dimension of evolutionary relation. “Logic” here does not mean formal logic, but the way structure shifts under constraint, collapses under tension, reorganizes after instability, and tends toward new modes of local stability. It shows that reality does not simply become manifest and then remain motionless, but preserves its layered organization through ongoing structural rearrangement.

Once this is in place, the relations among the six layers can be stated more clearly. L0 provides the unmanifest presupposition; L1 provides the earliest condition under which differentiation can persist; L2 gives that persistence a directional contour of manifestation; L3 organizes the paths along which such manifestation may unfold; L4 turns those paths into concrete structures within empirical reality; and L5 explains how those structures, under instability and limitation, move toward new local stabilities. These are not stacked parallel entities, nor a linear sequence of temporal origins, but six irreducible ontological functions through which reality becomes layered being.

More importantly, this six-level structure can be stably coordinated with the explanatory framework developed earlier in the book. Causal production corresponds chiefly to the processes of actualization at L4; structural constraint and

permissibility receive ontological expression chiefly through the organizational and convergent relations found in L3 and L5; and the stable manifestation of higher-level facts falls chiefly at L4, though its conditions are prepared in advance by L1, L2, and L3. In other words, this six-level architecture is not a fresh departure, but a further ontological consolidation of the differentiated functions the book has already established at the explanatory level. If Chapter 3 provided a differentiation of explanatory roles, what is provided here is the layered definition of Radiance-Will Ontology.

9.7 The Unity of Radiance, Will, and the Three Explanatory Roles

Radiance-Will Ontology is not an arbitrary piece of naming. Its importance lies in the fact that it reunifies the three explanatory roles developed since Chapter 3. Causal production explains how a result becomes actualized, how events and states enter reality through concrete processes. Structural constraint explains that such actualization does not occur in a vacuum, but is delimited within certain boundaries, organizational relations, and system-structures. Permissibility, in turn, explains which forms are open to the system under those constraints, and which forms are systematically excluded. Originally these were distinctions of explanatory function; Radiance-Will Ontology elevates them into three aspects of one and the same process of reality's organization.

Seen from this angle, radiance corresponds to how manifested results become stable objects at different scales; will corresponds to why these manifestations are opened, maintained, or excluded along certain directions of boundary-organization; and causal production corresponds to the concrete processes by which such forms are actualized. Thus, the earlier distinctions no longer amount merely to the claim that "there are different kinds of explanation," but become the claim that reality itself organizes itself through different functional dimensions.

Explanatory roles here are no longer merely analytical tools, but different graspable aspects of how reality is organized.

This unification brings another important gain as well: it prevents us from sliding back into monism. If causal production alone is emphasized, radiance gets pushed back into a mere after-image of what has already occurred, and will is rewritten as an unnecessary extra motive force. If structural constraint alone is emphasized, the process of actualization is overly weakened. If permissibility alone is emphasized, manifestation risks being abstracted into modal language without realizational grounding. Radiance-Will Ontology must retain all three together precisely because reality must be actualized, organized within boundaries, and manifested in stable forms. Remove any one of these, and the whole picture loses its balance again.

9.8 The Ontological Status of Higher-Level Facts

Within Radiance-Will Ontology, higher-level facts acquire a clearer status. Their reality does not consist in their being a second set of entities, nor in their existing independently of lower-level realization, nor again in their possessing some extra causal sovereignty. They are real because they take the form of stably manifested structures, because they appear as boundaries, patterns, and selections of relevant difference at a given scale of organization. In other words, the reality of higher-level facts is reality in the sense of radiance: they genuinely appear at a particular level and are continuously sustained by the structure of reality itself.

At the same time, higher-level facts also possess reality in the sense of will. They are not arbitrary cuts imposed on lower-level materials, but genuine outcomes of a system's organization: within given boundaries, certain differences are selected as relevant, certain functional relations are maintained, and certain patterns are opened as stable forms. The fact that higher-level facts can remain consistent across multiple realizations is precisely because reality has already, at

the organizational level, structurally selected what kinds of forms are to be maintained. They do not add another layer of stuff on top of the lower level; rather, they form a non-cancelable layer of organized reality emerging through it.

Accordingly, higher-level facts both depend on lower-level realization and do not thereby disappear; they involve neither additional inventory nor extra causal force. Their standing arises from the conjunction of stable manifestation and impersonal directedness. Put in the most concise form: higher-level facts are real as structured forms of existence, not as inventory-style existents. They are forms of reality at a higher resolution of organization within one and the same world.

9.9 The Non-Competitive Relation Between Whole and Part

If higher-level facts have this sort of structured existence, then the relation between whole and part must also be rethought. Traditional zero-sum pictures always suggest that if the parts already constitute reality, then the whole can only be redundant; and conversely, that if the whole is treated as real, then the determinacy of the parts must somehow be threatened. The earlier chapters have repeatedly shown that this picture does not fit explanatory practice. At the ontological level too, it must be abandoned. Whole and part are not competing for who is more real; they articulate different facts at different scales of organization within the same reality.

The parts give us materials, local processes, and paths of realization. The whole gives us boundaries, pattern-stability, and organizational form. The former explains how something is brought into reality; the latter explains how that actualization becomes, within a certain structure, the manifestation of a system. The reality of the parts does not cancel the reality of the whole, and the reality of the whole does not require swallowing up the parts into some further substance. If the parts determine in the sense of realization, then the whole determines in the sense of organization. Only by distinguishing these two kinds of determination can we understand why

dependence does not imply elimination, and why realization does not exhaust reality.

Within Radiance-Will Ontology, the whole is real because it constitutes the stable supporting boundary of some manifested form, while the parts are real because they constitute the concrete conditions under which those forms are actualized. The whole is not a second world beyond the parts, and the parts do not automatically exhaust the whole once they are listed out. Together they organize reality across different functional dimensions. For that reason, the whole-part relation should no longer be understood in terms of “who occupies the true place of being,” but in terms of “who grasps what sort of structural fact at what scale.”

9.10 Revision of Existing Ontological Pictures

From this point, the chapter can also state more clearly how its position differs from several familiar ontological pictures.

First, it revises flattening reductionism. That view does not merely insist on the importance of lower-level realization; it goes further and rewrites realizational fundamentality as ontological exclusivity, as though once the lower-level facts are fully listed, higher-level facts must be reduced to descriptive convenience. What Radiance-Will Ontology opposes is not dependence, but the leap from dependence to cancelability. Higher-level facts depend on the lower level, but that does not mean they lack their own status of structural manifestation and organization.

Second, it also revises strong entity-based holism. That position often tries to preserve the reality of wholes by introducing new entities, extra causal powers, or independently existing higher-level units. This may look generous to the reality of the whole, but in fact it simply reintroduces the same inflationary impulse that the earlier chapters have been trying to dismantle. Radiance-Will Ontology insists that the reality of wholes does not require inventory-style expansion. Wholes are real

because they correspond to stable manifestation and organizational boundaries, not because they occupy an extra slot on the list of objects.

Third, it differs as well from crude pluralism. Not every useful analytical angle, convenient label, or macro-level tag automatically constitutes a genuine level of reality. Only when a form is stable in manifestation, bounded in organization, sustainably supported in realization, and explanatorily indispensable do we have reason to acknowledge it as a real level. In this sense, Radiance-Will Ontology is neither a single-level worldview nor an anything-goes multilayer view, but a threshold-based realism about levels.

9.11 Chapter Summary: From Explanatory Plurality to Restrained Layered Structural Realism

This chapter began by responding to a question that naturally arises after the earlier chapters: if causal production, structural constraint, and permissibility really are distinct explanatory roles, and if higher-level facts, stable patterns, and structural boundaries continue to prove explanatorily indispensable, must these achievements still remain confined to the methodological or semantic level? The chapter's answer has been no. If these distinctions repeatedly play decisive roles across multiple theoretical contexts, then we eventually have to say something about how reality itself is organized. Otherwise, the structural facts and higher-level reality preserved throughout the book would slide back into mere pragmatic convenience.

At the same time, the chapter has insisted that this ontological articulation cannot take the inflationary form of introducing new entities, new causal powers, or new world-layers. For that reason, what it proposes is not an expansive entity ontology, but a restrained layered ontology, rhetorically named Radiance-Will Ontology. "Radiance" refers to the stable manifestation of reality's structure at different scales of organization; "will" refers to the impersonal directedness shown

by reality's boundaries toward possible forms. Together they express a picture of reality in which the world is not a flat stockpile of passive items, but a structured process that opens possibilities through boundaries, maintains forms through constraint, and manifests itself stably at different scales.

In this way, the chapter gathers the explanatory reconstruction of the earlier chapters into a more complete image of reality. Causal production explains how forms become actualized, structural constraint explains how actualization is delimited within boundaries, permissibility explains which forms are open to a system, and Radiance-Will Ontology explains how these three together constitute the layered organization of reality. Higher-level facts are thus no longer understood as extra entities, but as structured facts of stable manifestation; and whole and part are no longer understood as rivals in a zero-sum contest, but as non-competitive coexisting realities at different scales of one and the same world. If the earlier chapters were mainly occupied with dismantling mistaken forms of questioning, the task of this chapter has been to offer an ontological articulation capable of carrying the results of that dismantling.

Chapter 10

Conclusion: From Causal Centrism to the Multi-Role View of Explanation

By this point, the book has followed a continuously unfolding path in recasting a set of disputes that appear scattered on the surface but are in fact deeply connected. From emergence to higher-level explanation, from quantum entanglement to unknown phenomena, from mechanistic explanation to the methodological standing of non-predictive frameworks, and finally to the expression of a layered ontology, what the book has repeatedly confronted is the same deeper difficulty: in contemporary philosophy of science and related theoretical discussions, explanation is too stably compressed into a single task, namely, showing how some result is causally produced. Precisely for that reason, many theoretical demands that do not belong to the same level are forced into a single structure, and once those demands cannot all be satisfied at once, philosophical difficulty quickly expands into ontological crisis, causal conflict, or methodological suspicion.

The central work of this book has therefore not been to append one more position to existing disputes, but to rewrite the very form of the disputes themselves. Again and again it has argued that many persistent theoretical difficulties arise not first from a lack of empirical facts, but from a systematic misassignment of explanatory roles. Structural novelty is rewritten as causal novelty; higher-level explanation is rewritten as competition with lower-level causation; quantum correlation is rewritten as anomalous transmission; unknown phenomena are rewritten as automatic gaps in reality's inventory. So long as these issues are understood within one and the same production-centered framework, they will almost inevitably reproduce the same dilemmas, the same exclusion

structures, and the same inflationary impulses. The reason the book has been able to make relatively unified progress across multiple topics is that it keeps returning to this deeper level of organization and rearranging the space of explanation itself.

For that reason, the task of this concluding chapter is not merely to repeat what came before, but to answer a more concentrated question: what, exactly, has this book changed? What it changes is not only our conclusions about several individual issues, but our understanding of theoretical maturity itself. If the arguments of the book are sound, then many difficulties long treated as “forced on us by the world” are in fact difficulties first produced by wrongly posed explanatory questions; many pressures thought to require new causal powers or new ontological levels can instead be reassigned through a stricter differentiation of explanatory roles and a recognition of structural facts; and many non-predictive frameworks often treated as insufficiently “hard” are in fact carrying out indispensable work of clarification, limitation, and organization in theoretical progress.

10.1 What Has This Book Changed?

If the theoretical contribution of the entire book had to be compressed into a single sentence, it might be this: what the book changes is the way we understand explanation, levels of reality, and the sources of theoretical pressure. It does not merely claim that “there are other kinds of explanation,” nor does it vaguely defend higher-level phenomena. More specifically, it argues that many recurring philosophical deadlocks arise because explanation has been over-centralized into causal production. Once that centralization is loosened, our ways of understanding higher-level facts, novelty, relevance, unknown phenomena, and even theoretical achievement itself all undergo systematic change.

First, the book changes how we understand the source of explanatory conflict. Standard pictures often treat many difficulties as if they were competitions within

reality itself: whether higher and lower levels are fighting over causal position, whether whole and part are fighting over reality-status, whether locality and correlation are fighting over explanatory dominance, or whether existing ontology and unknown phenomena are fighting over the right to define the world. The present book argues instead that many of these conflicts are not first clashes within reality, but the result of explanatory tasks having been wrongly organized. In other words, it shifts theoretical attention from “which side wins?” to “why were these things placed within the same adjudicative structure in the first place?”

Second, the book changes how we understand the concept of the new. In standard disputes, novelty is too easily read straightaway as a new causal power, a new entity, or a new level. By contrast, through its analysis of emergence, higher-level facts, and unknown phenomena, the book argues that novelty may first be structural. A pattern, a boundary, a stable relation, or a way of organizing a joint outcome space can be genuinely new without having to prove itself by expanding the productive inventory of the world. The significance of this shift is not merely that it gives milder answers to certain disputes, but that it provides a more accurate criterion for what counts as real novelty.

Third, the book changes how we understand the status of higher-level facts and higher-level explanation. Higher-level phenomena no longer need extra causal powers in order to preserve their reality, and higher-level explanation no longer needs to compete with lower-level accounts for the same explanatory slot in order to preserve its legitimacy. Their reality and explanatory force are grounded instead in structural constraint, permissibility, stable patterning, level-relative relevance, and cross-realizational consistency. In short, the book tries to free higher-level reality from the old standard of “independent efficacy” and relocate it within a new standard of “stable organized fact.”

The book also changes how we understand the form of theoretical achievement. If explanatory roles are themselves multiple, then theoretical

progress cannot be defined only by predictive success. The earlier chapters' defense of non-predictive frameworks proceeds from exactly this point. The contribution here is not merely to secure room for a certain kind of philosophical work, but to argue that restructuring problems, limiting misplaced demands, organizing theoretical space, and eliminating category confusion are themselves serious theoretical accomplishments. Without such accomplishments, predictive and empirical work alike can easily be guided by the wrong questions.

Finally, the book changes how we understand ontological restraint. In the standard picture, restraint often means flattening: the fewer levels, boundaries, and structural facts one acknowledges, the more economical and mature one appears. The present book tries to show instead that genuine mature restraint does not consist in rejecting levels, but in refusing to rewrite levels as inventory expansion. A framework can acknowledge higher-level facts, structural boundaries, and organized permissibility while still remaining ontologically restrained. In other words, what the book seeks is not "less reality," but "a less inflationary acknowledgment of reality."

More generally, then, what this book changes is not that it adds something extra to the world, but that it rearranges the order of theoretical judgment. It asks us first to distinguish explanatory roles, then assess theoretical pressure; first to analyze whether the problem has been wrongly posed, then ask whether stronger ontological commitments are needed; first to acknowledge the reality of structural facts, then decide whether new entities or causal powers are truly required. This reordering matters because it reopens a clearer, more restrained, and more stable space for dealing with disputes that have long seemed to admit only radical options.

10.2 Rewriting the Major Disputes

If the previous section answers the question “what has the book changed overall?”, this section asks more concretely how that change appears in the recasting of the central disputes. For the method of the book has not been to begin from one abstract master principle and then mechanically apply it to different problems. Rather, it has repeatedly identified the same deep structure across several difficulties that initially seem independent, and loosened the apparent necessity of their dilemmas by rewriting the form of the problem. It is precisely here that the unity of the book becomes clearest.

First, the book rewrites the debate over emergence. Standard discussions usually organize emergence around the dilemma of strong versus weak emergence: either one preserves the reality of higher-level phenomena at the cost of introducing new causal powers or irreducible laws, or one preserves physical closure and realization-dependence at the cost of reducing higher-level novelty to something merely epistemic, methodological, or computational. The book argues that this dilemma is so persistent not simply because each side holds one horn of it, but because novelty has from the beginning been rewritten as causal novelty. By introducing structural novelty, the book shifts the issue from “are there new causal powers?” to “have new forms of organization, stable patterns, and permissibility-boundaries emerged?” Emergence is thus no longer forced to oscillate between mysterious causation and explanatory nihilism, but can be re-understood as a question about structural fact and level-manifestation.

Second, the book rewrites the debate over higher-level explanation. Under the combined pressure of exclusion arguments and mechanistic explanation, higher-level accounts have long been forced into one of three unfavorable positions: either they are redundant, because lower-level explanation is already sufficient; or they are competitive, because they seem to claim extra causal efficacy with respect to the same result; or they are merely cognitively convenient, because they no

longer count as genuine explanation. The key to the book's recasting of this dispute is to sever the automatic link between causal sufficiency and explanatory sufficiency, and to show that higher-level explanation is legitimate not because it adds new productive content, but because it undertakes different explanatory tasks concerning units of organization, level-relative relevance, stable patterns, and structural boundaries. In this way, higher-level explanation no longer needs extra causal power in order to secure itself, but can regain theoretical legitimacy in a non-competitive form.

Third, the book rewrites the debate over quantum entanglement. Standard philosophical readings often translate Bell-type pressure too quickly into a demand for nonlocal causation, as though once strong correlations cannot be handled by locally independent models, they must be explained by some form of action at a distance, hidden linkage, or stronger holistic entity. The book argues that the crucial miswriting here occurs in the leap from correlation to a transmission template. What the Bell framework directly places pressure on is a certain class of joint-structure models, not a direct requirement for anomalous transmission mechanisms. By introducing structural inseparability, structural constraint, and permissibility analysis, the book rewrites entanglement as a question of how joint outcome spaces are organized, rather than a question of how influence travels across space. This preserves the strength of quantum correlation while significantly weakening the impulse toward premature ontological inflation.

Fourth, the book rewrites the debate over unknown and anomalous phenomena. In this context, empirical novelty is often rapidly elevated into ontological anxiety: as soon as existing theory cannot smoothly absorb some object, people readily conclude that reality's catalog must itself be incomplete. By distinguishing explanatory robustness from ontological stability, the book argues that unknown phenomena first challenge not "what else exists in the world," but "whether we have properly assigned levels of problem, variable structure, and

explanatory roles." Anomalous phenomena therefore function more like stress tests for explanatory discipline than shortcuts to new existents. What the book accomplishes here is to delay ontological impulse, allowing model revision, structural clarification, problem reconstruction, and level-analysis to do their work first.

Fifth, the book rewrites a broader methodological dispute, namely the question of the standing of non-predictive frameworks. Standard pictures tie theoretical value too steadily to predictive success, with the result that work aimed primarily at problem-reconstruction, conceptual clarification, structural organization, and boundary-limitation is easily treated as second-rate. By disentangling the long-standing conflation among prediction, causation, and explanation, the book argues that non-predictive frameworks can produce rigorous theoretical progress by other means. In this way, the book itself is repositioned as a methodologically legitimate form of structural work, rather than something that must apologize for being "not yet at the predictive stage."

Taken together, these rewritings reveal a very clear common structure. In almost every case, the original dispute is governed by some overly strong problem-form: emergence must involve new causal powers; higher-level explanation must compete for the same causal position; entanglement must involve anomalous transmission; unknown phenomena must signal new ontology; theoretical achievement must first appear as prediction. What the book does is not seek milder compromises within those strong demands, but show that many of the difficulties arise first from the misplacement of the demands themselves. That is to say, the book does not simply offer new answers to old questions; across multiple domains, it repeatedly demonstrates that the old questions were often already wrongly written before they became problems.

This also shows that the book's "rewriting" is not mere verbal substitution. If it were only a change of wording without a change in problem structure, then the

dilemmas of strong and weak emergence, the exclusion-structure of higher-level explanation, the nonlocal pressure of entanglement, and the ontological anxiety around unknown phenomena would all remain exactly as before. Genuine rewriting means changing what counts as the core difficulty, changing what is taken to require explanation, and changing when it is appropriate to escalate ontological commitment. It is here that the substantive force of the book's multi-role conception of explanation becomes visible.

The conclusion of this section can therefore be put simply: the book has not separately handled several accidentally similar problems, but has repeatedly performed the same operation across several major disputes, namely decomposing problems that had been compressed into single productive, ontological, or predictive demands into distinct explanatory tasks at different levels. It is this sustained rewriting that gives the book its unified theoretical outline.

10.3 From Explanatory Monism to the Differentiation of Explanatory Roles

If the book's central turn had to be reduced to its most concise and comprehensive formula, it would be this: from explanatory monism to the differentiation of explanatory roles. The disputes discussed in earlier chapters – emergence, higher-level explanation, entanglement, unknown phenomena, and the status of non-predictive frameworks – belong on the surface to different domains. But the reason they can be treated together in one systematic argument is not that they all happen to concern "complexity" or "higher-level phenomena." It is that they all expose the same deep structure: explanation has been too stably contracted into a single activity, especially the activity of showing how some result is causally produced. Once that contraction becomes the default background, the rest of the theoretical pressures begin to proliferate.

Back in Chapter 2, the book analyzed this explanatory monism into evaluative monism, role monism, and competitive monism. Looking back across the whole book, it becomes even clearer that this threefold diagnosis is not a local analytical device, but nearly the fulcrum of the entire project. Evaluative monism inclines people to regard the most weighty explanations as those closest to causal sufficiency. Role monism goes further and reduces explanation itself to the single task of answering how something was produced. Competitive monism then concludes that once some level's productive account is deemed sufficient, any other account must be redundant, competitive, or merely secondary convenience. It is precisely the combination of these three that places so many disputes into the wrong arena before they even begin.

That is why the book can sometimes look as though it is simply "making room for higher-level, structural, or constraining explanations," when in fact its deeper work is not to grant special exemptions to a few marginalized explanatory forms, but to redraw the role-structure of explanation itself. The book argues that explanatory activity cannot be compressed into one productive task alone. It also includes the articulation of structural constraint, the clarification of permissibility-spaces, the identification of pattern-stability, the selection of level-relevant differences, and even the rewriting of problem-boundaries themselves. Explanation, in other words, is not a single linear track, but a set of interrelated yet non-absorptive task-domains. It is precisely because of this that the earlier chapters' rewritings become possible: higher-level phenomena need not prove themselves by extra causal powers, higher-level explanation need not compete with lower-level causes, entanglement need not be forced into a transmission template, and unknown phenomena need not demonstrate their seriousness by demanding new ontology.

This turn from unity to differentiation changes not only our verdicts on individual cases, but also our diagnosis of theoretical difficulty itself. Under

explanatory monism, a problem tends to be interpreted as showing that some level's account is not yet causal enough, predictive enough, or productive enough, and all remaining pressure is then pushed upward to the strongest level: either add new causal powers, add new ontology, or admit explanatory failure. The differentiation of explanatory roles offers a different diagnostic path. When a theoretical problem appears stubborn, we no longer ask first "which segment of the causal chain is still missing?" but "which explanatory task has been rewritten into the wrong form here?" Sometimes what is missing is a productive account; sometimes a structural boundary; sometimes an analysis of permissibility-space; and sometimes what is missing is only a reconstruction of the problem itself. Once these roles are distinguished, the location of theoretical pressure becomes visible.

This also means that the book is not merely adding the slogan that "explanation is diverse." If one only says there are many kinds of explanation, without giving principles of differentiation, that plurality easily collapses into vague permissiveness: everyone keeps their preferred account, but the dispute itself makes no progress. The plurality the book pursues is not merely descriptive but structural. It demands distinctions among explanatory task-boundaries, level-relevant differences, the legitimacy of problem-forms, and the placement of different theoretical pressures. Explanatory-role differentiation is therefore not a celebration of plurality, but its disciplined organization.

From this angle, many of the book's central concepts can be seen as serving this single turn. Structural novelty matters because it partly frees "the new" from the causal role. The non-competitive standing of higher-level explanation matters because it frees explanatory legitimacy from a single productive slot. Structural inseparability matters because it shifts Bell-type pressure away from transmission and back toward the organization of result-space. Ontological robustness matters because it reallocates the pressure generated by empirical novelty from automatic expansion of ontology to the levels of models, boundaries, and explanatory

stratification. Without differentiation of explanatory roles, these ideas would appear as a series of local patches. Once placed back within this overarching turn, they form a clear theoretical line.

This also shows that what the book truly opposes is not causal explanation, mechanistic account, predictive success, or ontological revision as such. It opposes their overgeneralization into the only legitimate roles. Causal production is indeed important; mechanistic explanation is often powerful; prediction is unquestionably one central form of theoretical achievement; and ontological revision is sometimes unavoidable. The trouble begins only when any one of these is made exclusive. The significance of explanatory-role differentiation lies in restoring each of these achievements to its proper place rather than turning any one of them into a universal tribunal. If the book has a higher-order philosophical commitment, it is this: theoretical maturity does not require flattening the world into one explanatory template, but requires a more accurate grasp of when, where, and why different templates are legitimate.

The conclusion of this section can therefore be stated quite clearly: the book's deepest transformation is not a shift from one specific school to another, but a shift from explanatory monism to the differentiation of explanatory roles. It is this shift that allows higher-level facts, structural constraint, permissibility, non-predictive frameworks, and ontological robustness to appear not as marginal adjustments, but as different components within one and the same theoretical picture. Only in this sense do the several parts of the book truly belong to one chain of argument rather than a series of loosely resonant case studies.

10.4 The Unity of Explanation, Methodology, and Ontology

At this point, the book has advanced one and the same task across three distinct but continuously interdependent levels. First, at the explanatory level, it has moved from explanatory monism to the differentiation of explanatory roles. Second, at the

methodological level, it has moved from predictive centrism to the legitimation of non-predictive structural frameworks. Third, at the ontological level, it has moved away from both flattening reduction and inflationary level-entity metaphysics toward a restrained layered ontology.

These three levels should not be understood as three independent conclusions merely placed side by side. Their deeper significance lies in the fact that they mutually support one another and form a unified theoretical reorientation. If explanatory roles are not differentiated, then methodological space contracts as well: theoretical value gets tied too narrowly to prediction and production, and non-predictive frameworks are reduced to something preparatory or secondary. If methodology is not reconstructed, then ontology too easily swings between two extremes: either only what predicts and produces counts as real, or every unresolved pressure gets converted into a justification for ontological expansion. And if ontology is left unaddressed, the earlier explanatory and methodological achievements remain suspended, always at risk of collapsing back into instrumentalism or semantic convenience. For that reason, the book's contribution is not a set of disconnected corrections, but a linked reordering of explanation, methodology, and ontology.

From this perspective, the differentiation of explanatory roles is the starting point, but not the endpoint. It first makes visible that theoretical pressure does not arise from one source alone, and that not every difficulty is a deficit in causal production. Methodological reconstruction then shows that once this point is granted, theory must also recognize forms of progress other than prediction, including the reorganization of problem-space, the clarification of explanatory tasks, the restriction of illegitimate inferences, and the exposure of category mistakes. Ontological reconstruction finally shows that once these explanatory and methodological achievements are taken seriously, reality itself must be conceived in a way that can acknowledge structured, level-sensitive facts without collapsing

either into flattening reduction or into ontological excess. In this way, the three levels form a continuous progression rather than a mere sequence of topics.

This unity also clarifies the book's overall philosophical posture. It is not simply anti-reductionist, because it does not reject dependence, realization, or lower-level causal explanation. Nor is it straightforwardly pluralist, because it does not treat every useful distinction as equally real or equally legitimate. Nor again is it a form of metaphysical inflation, because it repeatedly refuses the automatic move from explanatory difficulty to new entities, new forces, or new levels. What the book proposes instead is a disciplined enlargement of theoretical space: one that preserves the importance of causal production without granting it exclusive authority, preserves the reality of higher-level and structural facts without detaching them from realization, and preserves ontological seriousness without letting it become a reflex of theoretical anxiety.

It is precisely here that the notion of maturity, which has hovered in the background throughout the book, can finally be stated more directly. Theoretical maturity does not consist in always choosing the most austere picture, nor in always embracing the most radical revision. It consists in preserving order among kinds of question, levels of analysis, and forms of commitment. A mature framework knows that explanatory adequacy is not identical with causal sufficiency, that methodological value is not exhausted by prediction, and that ontological recognition need not take the form of numerical addition to a stock of beings. Maturity, in this sense, is not hesitation but proportion.

This also reveals why the book's recasting of so many major disputes should not be mistaken for a merely local intervention. What is at stake is not only a better treatment of emergence, higher-level explanation, entanglement, unknown phenomena, or non-predictive frameworks taken one by one. What is at stake is a more general transformation in how theoretical pressure is identified, distributed, and answered. Once the unity of explanation, methodology, and ontology is seen,

the deeper aim of the book becomes clearer: not to defend a special class of phenomena or a favored metaphysical position, but to reconstruct the conditions under which theory itself remains disciplined, open, and adequately responsive to the structure of reality.

The conclusion of this section can therefore be put in the following way: the book's main achievement is not simply to defend explanatory plurality, methodological flexibility, or restrained ontology separately, but to show that these three belong together. Explanation requires differentiation, methodology requires more than prediction, and ontology requires structure without inflation. Only when these three insights are held together does the full outline of the book's position come into view.

10.5 Final Conclusion

The path traced by this book can now be stated in its most compressed form. Many of the philosophical pressures that dominate contemporary debates do not first arise because the world itself forces us into crisis, but because explanation has been prematurely narrowed to one task and then used as the universal measure of theory. Once explanation is treated as if it were exhausted by causal production, the rest follows almost automatically: higher-level facts look causally suspicious, emergence looks metaphysically explosive, entanglement looks like anomalous transmission, unknown phenomena look like evidence of ontological deficit, and non-predictive work looks theoretically inferior. What the book has tried to show is that these pressures are often generated less by reality itself than by the way explanatory labor has been wrongly organized.

Against that background, the book has advanced one central reorientation: from causal centrism to a multi-role conception of explanation. This does not mean replacing causal explanation with something else, nor denying the central importance of mechanisms, prediction, or empirical rigor. It means refusing to let

any one explanatory role become the tribunal for all the others. Causal production, structural constraint, and permissibility answer different but interconnected questions. Once that differentiation is restored, many disputes that previously seemed locked into impossible dilemmas become more tractable, more proportionate, and less vulnerable to inflationary conclusions.

The implications of this reorientation are wide-ranging. It changes how novelty is understood, how higher-level explanations are justified, how quantum correlations are interpreted, how anomalous phenomena are handled, how non-predictive theoretical work is valued, and how ontology itself is articulated. But these are not isolated consequences. Together they express one underlying shift: theoretical pressure should be distributed according to the kind of explanatory work actually required, rather than being automatically pushed upward into the strongest available categories of causation, prediction, or ontology.

This is why the book's final position is neither reductionist nor metaphysically exuberant. It does not flatten reality into one level of causally productive facts, but neither does it proliferate entities whenever explanatory friction appears. What it seeks instead is a disciplined realism: one that can acknowledge stable patterns, structural boundaries, higher-level facts, and organized spaces of possibility as real without turning them into a second inventory of beings. Reality, on this view, is neither exhausted by its lowest-level components nor enlarged by every theoretical unease. It is structured, layered, and differentially organized, and theory matures insofar as it learns to respond to that structure in a proportionate way.

In the end, then, the deepest claim of the book is not merely about any single dispute. It is about what counts as a mature relation between theory and world. A mature theory does not rush to compress every question into one form, nor does it treat every difficulty as a mandate for metaphysical escalation. It distinguishes tasks, preserves levels, accepts structure, and revises in order. If this book succeeds, its lasting contribution will not be that it settled every controversy it

touched, but that it offered a clearer way to see why so many controversies have been misshaped by the same underlying habit of explanatory compression.

Put most simply: the book argues that we understand the world better when we stop asking explanation to do only one thing.

Note to the Reader:

On the Book's Mode of Composition and Responsibility

This book is a theoretical work authored under human direction and completed with the assistance of artificial intelligence tools. In the course of its conception, writing, and revision, such tools were used in a number of auxiliary capacities, including clarifying the problem structure, checking continuity between chapters, improving expression, reviewing logical coherence, and offering draft passages and revision suggestions for selected sections.

It should be made clear that artificial intelligence tools do not count as the author of this book, nor do they bear co-authorship or named authorship. The book's core sense of the problem, theoretical direction, conceptual choices, chapter arrangement, final judgments, and overall responsibility all belong to the human author. Any material generated, organized, or suggested by artificial intelligence tools was reviewed, selected, revised, and integrated by the human author, and its adoption was determined solely by the human author.

This statement is retained out of concern for transparency in composition and academic honesty. Artificial intelligence tools can assist with linguistic organization, structural checking, and repeated revision, but they cannot replace the author's judgment regarding the content of thought, nor can they replace responsibility for concepts, arguments, or facts. Especially in a theoretical work such as this one, dealing with explanation, philosophy of science, emergence, and ontological commitment, what matters is not whether the text made use of tools, but whether the author is able to take responsibility for every concept, argument, and conclusion contained in it.

Accordingly, all views, arguments, errors, and shortcomings in this book are the responsibility of the human author. Artificial intelligence tools are merely auxiliary media within the process of composition, not subjects of theoretical responsibility. Readers, therefore, should understand this book as a philosophical work directed, approved, and owned by a human author, rather than as a body of text automatically generated by artificial intelligence.

This statement is both a gesture of respect toward the reader and an extension of the book's own theoretical position. For one of the book's central concerns is precisely how to distinguish different roles, different levels, and different sites of responsibility. The use of artificial intelligence in the writing process must likewise be assigned its proper place: it may participate in expression, organization, and feedback, but it cannot replace the author's judgment, responsibility, or intellectual accountability.

References

The references collected in this book are drawn mainly from core works relevant to theories of explanation, causal explanation, mechanistic explanation, higher-level causation, non-causal explanation, emergence, the foundations of quantum theory, and structural ontology. Since the aim of this book is not to provide a historical survey of philosophy of science, but to propose a theoretical framework centered on the differentiation of explanatory roles, the works listed here are limited to those directly engaged with the problem-space of the book's argument. Related research not discussed explicitly in the main text is generally not included unless it bears a particularly close relation to the book's core concepts.

I. Scientific Explanation, Causal Explanation, and Scientific Understanding

[1](#) Achinstein, P. *The Nature of Explanation*. Oxford: Oxford University Press, 1983.

[2](#) Bokulich, A. "How Scientific Models Can Explain." *Synthese* 180, no. 1 (2011): 33-45.

[3](#) Cartwright, N. *How the Laws of Physics Lie*. Oxford: Oxford University Press, 1983.

[4](#) Cartwright, N. *The Dappled World: A Study of the Boundaries of Science*. Cambridge: Cambridge University Press, 1999.

[5](#) De Regt, H. W. *Understanding Scientific Understanding*. Oxford: Oxford University Press, 2017.

[6](#) Duhem, P. *The Aim and Structure of Physical Theory*. Princeton: Princeton University Press, 1954.

[7](#) Friedman, M. "Explanation and Scientific Understanding." *The Journal of Philosophy* 71, no. 1 (1974): 5-19.

[8](#) Hacking, I. *Representing and Intervening: Introductory Topics in the Philosophy of Natural Science*. Cambridge: Cambridge University Press, 1983.

[9](#) Hempel, C. G. *Aspects of Scientific Explanation and Other Essays in the Philosophy of Science*. New York: Free Press, 1965.

[10](#) Hempel, C. G., and P. Oppenheim. "Studies in the Logic of Explanation." *Philosophy of Science* 15, no. 2 (1948): 135-175.

[11](#) Hitchcock, C. "The Intransitivity of Causation Revealed in Equations and Graphs." *The Journal of Philosophy* 98, no. 6 (2001): 273-299.

[12](#) Hitchcock, C., and J. Woodward. "Explanatory Generalizations, Part I: A Counterfactual Account." *Noûs* 37, no. 1 (2003): 1-24.

[13](#) Khalifa, K. *Understanding, Explanation, and Scientific Knowledge*. Cambridge: Cambridge University Press, 2017.

[14](#) Kitcher, P. "Explanatory Unification." *Philosophy of Science* 48, no. 4 (1981): 507-531.

[15](#) Kuhn, T. S. *The Structure of Scientific Revolutions*. 4th ed. Chicago: University of Chicago Press, 2012.

[16](#) Lakatos, I. "Falsification and the Methodology of Scientific Research Programmes." In *Criticism and the Growth of Knowledge*, edited by I. Lakatos and A. Musgrave, 91-196. Cambridge: Cambridge University Press, 1970.

[17](#) Lewis, D. "Causation." *The Journal of Philosophy* 70, no. 17 (1973): 556-567.

[18](#) Lipton, P. *Inference to the Best Explanation*. 2nd ed. London: Routledge, 2004.

[19](#) Pearl, J. *Causality: Models, Reasoning, and Inference*. 2nd ed. Cambridge: Cambridge University Press, 2009.

[20](#) Popper, K. R. *The Logic of Scientific Discovery*. London: Hutchinson, 1959.

[21](#) Potochnik, A. *Idealization and the Aims of Science*. Chicago: University of Chicago Press, 2017.

[22](#) Quine, W. V. O. "Two Dogmas of Empiricism." *The Philosophical Review* 60, no. 1 (1951): 20-43.

[23](#) Salmon, W. C. *Scientific Explanation and the Causal Structure of the World*. Princeton: Princeton University Press, 1984.

[24](#) Salmon, W. C. *Four Decades of Scientific Explanation*. Minneapolis: University of Minnesota Press, 1989.

[25](#) Spirtes, P., C. Glymour, and R. Scheines. *Causation, Prediction, and Search*. 2nd ed. Cambridge, MA: MIT Press, 2000.

[26](#) Stanford, P. K. *Exceeding Our Grasp: Science, History, and the Problem of Unconceived Alternatives*. Oxford: Oxford University Press, 2006.

[27](#) Strevens, M. *Depth: An Account of Scientific Explanation*. Cambridge, MA: Harvard University Press, 2008.

[28](#) van Fraassen, B. C. *The Scientific Image*. Oxford: Oxford University Press, 1980.

[29](#) Woodward, J. *Making Things Happen: A Theory of Causal Explanation*. Oxford: Oxford University Press, 2003.

II. Mechanistic Explanation, Levels, and Scientific Practice

[30](#) Bechtel, W. *Mental Mechanisms: Philosophical Perspectives on Cognitive Neuroscience*. London: Routledge, 2008.

[31](#) Bechtel, W., and A. Abrahamsen. "Explanation: A Mechanist Alternative." *Studies in History and Philosophy of Biological and Biomedical Sciences* 36, no. 2 (2005): 421-441.

[32](#) Bickle, J. *Philosophy and Neuroscience: A Ruthlessly Reductive Account*. Dordrecht: Kluwer Academic Publishers, 2003.

[33](#) Craver, C. F. *Explaining the Brain: Mechanisms and the Mosaic Unity of Neuroscience*. Oxford: Oxford University Press, 2007.

[34](#) Craver, C. F., and W. Bechtel. "Top-Down Causation without Top-Down Causes." *Biology and Philosophy* 22, no. 4 (2007): 547-563.

[35](#) Craver, C. F., and L. Darden. *In Search of Mechanisms: Discoveries across the Life Sciences*. Chicago: University of Chicago Press, 2013.

[36](#) Glennan, S. "Mechanisms and the Nature of Causation." *Erkenntnis* 44, no. 1 (1996): 49-71.

[37](#) Glennan, S. "Rethinking Mechanistic Explanation." *Philosophy of Science* 69, no. S3 (2002): S342-S353.

[38](#) Glennan, S. *The New Mechanical Philosophy*. Oxford: Oxford University Press, 2017.

[39](#) Illari, P. M., and J. Williamson. "What Is a Mechanism? Thinking about Mechanisms across the Sciences." *European Journal for Philosophy of Science* 2, no. 1 (2012): 119-135.

[40](#) Machamer, P., L. Darden, and C. F. Craver. "Thinking about Mechanisms." *Philosophy of Science* 67, no. 1 (2000): 1-25.

[41](#) Tabery, J. "Synthesizing Activities and Interactions in the Concept of a Mechanism." *Philosophy of Science* 71, no. 1 (2004): 1-15.

III. Higher-Level Causation, Realization, Exclusion Arguments, and Emergence

[42](#) Anderson, P. W. "More Is Different." *Science* 177, no. 4047 (1972): 393-396.

[43](#) Bedau, M. A. "Weak Emergence." *Philosophical Perspectives* 11 (1997): 375-399.

[44](#) Bennett, K. "Why the Exclusion Problem Seems Intractable, and How, Just Maybe, to Tract It." *Noûs* 37, no. 3 (2003): 471-497.

[45](#) Bennett, K. *Making Things Up*. Oxford: Oxford University Press, 2017.

[46](#) Chalmers, D. J. "Strong and Weak Emergence." In *The Re-Emergence of Emergence: The Emergentist Hypothesis from Science to Religion*, edited by P. Clayton and P. Davies, 244-256. Oxford: Oxford University Press, 2006.

[47](#) Fodor, J. A. "Special Sciences, or the Disunity of Science as a Working Hypothesis." *Synthese* 28, no. 2 (1974): 97-115.

[48](#) Gillett, C. *Reduction and Emergence in Science and Philosophy*. Cambridge: Cambridge University Press, 2016.

[49](#) Humphreys, P. "How Properties Emerge." *Philosophy of Science* 64, no. 1 (1997): 1-17.

[50](#) Kim, J. *Supervenience and Mind: Selected Philosophical Essays*. Cambridge: Cambridge University Press, 1993.

[51](#) Kim, J. *Mind in a Physical World: An Essay on the Mind-Body Problem and Mental Causation*. Cambridge, MA: MIT Press, 1998.

[52](#) Kim, J. "Making Sense of Emergence." *Philosophical Studies* 95, no. 1/2 (1999): 3-36.

[53](#) Kim, J. *Physicalism, or Something Near Enough*. Princeton: Princeton University Press, 2005.

[54](#) List, C., and P. Menzies. "Nonreductive Physicalism and the Limits of the Exclusion Principle." *The Journal of Philosophy* 106, no. 9 (2009): 475-502.

[55](#) Mitchell, S. D. *Unsimple Truths: Science, Complexity, and Policy*. Chicago: University of Chicago Press, 2009.

[56](#) Nagel, E. *The Structure of Science: Problems in the Logic of Scientific Explanation*. New York: Harcourt, Brace & World, 1961.

[57](#) O'Connor, T. "Emergent Properties." *American Philosophical Quarterly* 31, no. 2 (1994): 91-104.

[58](#) O'Connor, T., and H. Y. Wong. "The Metaphysics of Emergence." *Noûs* 39, no. 4 (2005): 658-678.

[59](#) Oppenheim, P., and H. Putnam. "Unity of Science as a Working Hypothesis." In *Minnesota Studies in the Philosophy of Science, Volume II: Concepts, Theories, and the Mind-Body Problem*, edited by H. Feigl, M. Scriven, and G. Maxwell, 3-36. Minneapolis: University of Minnesota Press, 1958.

[60](#) Polger, T. W., and L. A. Shapiro. *The Multiple Realization Book*. Oxford: Oxford University Press, 2016.

[61](#) Shoemaker, S. *Physical Realization*. Oxford: Oxford University Press, 2007.

[62](#) Silberstein, M., and J. McGeever. "The Search for Ontological Emergence." *The Philosophical Quarterly* 49, no. 195 (1999): 182-200.

[63](#) Wilson, J. M. *Metaphysical Emergence*. Oxford: Oxford University Press, 2021.

[64](#) Wimsatt, W. C. *Re-Engineering Philosophy for Limited Beings: Piecewise Approximations to Reality*. Cambridge, MA: Harvard University Press, 2007.

[65](#) Yablo, S. "Mental Causation." *The Philosophical Review* 101, no. 2 (1992): 245-280.

IV. Non-Causal Explanation, Mathematical Explanation, Structural Constraint, and Permissibility

[66](#) Baker, A. "Are There Genuine Mathematical Explanations of Physical Phenomena?" *Mind* 114, no. 454 (2005): 223-238.

[67](#) Batterman, R. W. "Multiple Realizability and Universality." *The British Journal for the Philosophy of Science* 51, no. 1 (2000): 115-145.

[68](#) Batterman, R. W. *The Devil in the Details: Asymptotic Reasoning in Explanation, Reduction, and Emergence*. Oxford: Oxford University Press, 2002.

[69](#) Batterman, R. W. "On the Explanatory Role of Mathematics in Empirical Science." *The British Journal for the Philosophy of Science* 61, no. 1 (2010): 1-25.

[70](#) Batterman, R. W. *A Middle Way: A Non-Fundamental Approach to Many-Body Physics*. Oxford: Oxford University Press, 2021.

[71](#) Batterman, R. W., and C. C. Rice. "Minimal Model Explanations." *Philosophy of Science* 81, no. 3 (2014): 349-376.

[72](#) Colyvan, M. *The Indispensability of Mathematics*. Oxford: Oxford University Press, 2001.

[73](#) Huneman, P. "Topological Explanations and Robustness in Biological Sciences." *Synthese* 177, no. 2 (2010): 213-245.

[74](#) Lange, M. "What Makes a Scientific Explanation Distinctively Mathematical?" *The British Journal for the Philosophy of Science* 64, no. 3 (2013): 485-511.

[75](#) Lange, M. *Because Without Cause: Non-Causal Explanations in Science and Mathematics*. Oxford: Oxford University Press, 2017.

[76](#) Morrison, M. *Reconstructing Reality: Models, Mathematics, and Simulations*. Oxford: Oxford University Press, 2015.

[77](#) Pincock, C. *Mathematics and Scientific Representation*. Oxford: Oxford University Press, 2012.

[78](#) Reutlinger, A., and J. Saatsi, eds. *Explanation Beyond Causation: Philosophical Perspectives on Non-Causal Explanations*. Oxford: Oxford University Press, 2018.

[79](#) Rice, C. C. "Moving Beyond Causes: Optimality Models and Scientific Explanation." *Noûs* 49, no. 3 (2015): 589-615.

[80](#) Steiner, M. "Mathematical Explanation." *Philosophical Studies* 34, no. 2 (1978): 135-151.

V. Complex Systems, Organizational Form, and Stable Patterns

[81](#) Bar-Yam, Y. *Dynamics of Complex Systems*. Reading, MA: Addison-Wesley, 1997.

[82](#) Cilliers, P. *Complexity and Postmodernism: Understanding Complex Systems*. London: Routledge, 1998.

[83](#) Holland, J. H. *Emergence: From Chaos to Order*. Reading, MA: Addison-Wesley, 1998.

[84](#) Kauffman, S. A. *The Origins of Order: Self-Organization and Selection in Evolution*. Oxford: Oxford University Press, 1993.

[85](#) Laughlin, R. B. *A Different Universe: Reinventing Physics from the Bottom Down*. New York: Basic Books, 2005.

[86](#) Prigogine, I., and I. Stengers. *Order out of Chaos: Man's New Dialogue with Nature*. New York: Bantam Books, 1984.

[87](#) Simon, H. A. "The Architecture of Complexity." *Proceedings of the American Philosophical Society* 106, no. 6 (1962): 467-482.

VI. Quantum Entanglement, Nonlocality, and the Foundations of Quantum Theory

[88](#) Albert, D. Z. *Quantum Mechanics and Experience*. Cambridge, MA: Harvard University Press, 1992.

[89](#) Aspect, A., J. Dalibard, and G. Roger. "Experimental Test of Bell's Inequalities Using Time-Varying Analyzers." *Physical Review Letters* 49, no. 25 (1982): 1804-1807.

[90](#) Bell, J. S. "On the Einstein Podolsky Rosen Paradox." *Physics Physique Fizika* 1, no. 3 (1964): 195-200.

[91](#) Bell, J. S. *Speakable and Unspeakable in Quantum Mechanics*. 2nd ed. Cambridge: Cambridge University Press, 2004.

[92](#) Bohm, D. "A Suggested Interpretation of the Quantum Theory in Terms of 'Hidden' Variables. I." *Physical Review* 85, no. 2 (1952): 166-179.

[93](#) Bohm, D. "A Suggested Interpretation of the Quantum Theory in Terms of 'Hidden' Variables. II." *Physical Review* 85, no. 2 (1952): 180-193.

[94](#) Brunner, N., D. Cavalcanti, S. Pironio, V. Scarani, and S. Wehner. "Bell Nonlocality." *Reviews of Modern Physics* 86, no. 2 (2014): 419-478.

[95](#) Bub, J. *Interpreting the Quantum World*. Cambridge: Cambridge University Press, 1997.

[96](#) Clauser, J. F., M. A. Horne, A. Shimony, and R. A. Holt. "Proposed Experiment to Test Local Hidden-Variable Theories." *Physical Review Letters* 23, no. 15 (1969): 880-884.

[97](#) Cushing, J. T., and E. McMullin, eds. *Philosophical Consequences of Quantum Theory: Reflections on Bell's Theorem*. Notre Dame: University of Notre Dame Press, 1989.

[98](#) Einstein, A., B. Podolsky, and N. Rosen. "Can Quantum-Mechanical Description of Physical Reality Be Considered Complete?" *Physical Review* 47, no. 10 (1935): 777-780.

[99](#) Fine, A. "Hidden Variables, Joint Probability, and the Bell Inequalities." *Physical Review Letters* 48, no. 5 (1982): 291-295.

[100](#) Fuchs, C. A., N. D. Mermin, and R. Schack. "An Introduction to QBism with an Application to the Locality of Quantum Mechanics." *American Journal of Physics* 82, no. 8 (2014): 749-754.

[101](#) Healey, R. *The Philosophy of Quantum Mechanics: An Interactive Interpretation*. Cambridge: Cambridge University Press, 1989.

[102](#) Howard, D. "Holism, Separability, and the Metaphysical Implications of the Bell Experiments." In *Philosophical Consequences of Quantum Theory: Reflections on Bell's Theorem*, edited by J. T. Cushing and E. McMullin, 224-253. Notre Dame: University of Notre Dame Press, 1989.

[103](#) Jarrett, J. P. "On the Physical Significance of the Locality Conditions in the Bell Arguments." *Noûs* 18, no. 4 (1984): 569-589.

[104](#) Maudlin, T. *Quantum Non-Locality and Relativity: Metaphysical Intimations of Modern Physics*. 3rd ed. Malden, MA: Wiley-Blackwell, 2011.

[105](#) Mermin, N. D. "Bringing Home the Atomic World: Quantum Mysteries for Anybody." *American Journal of Physics* 49, no. 10 (1981): 940-943.

[106](#) Norsen, T. *Foundations of Quantum Mechanics: An Exploration of the Physical Meaning of Quantum Theory*. Cham: Springer, 2017.

[107](#) Rovelli, C. "Relational Quantum Mechanics." *International Journal of Theoretical Physics* 35, no. 8 (1996): 1637-1678.

[108](#) Timpson, C. G. *Quantum Information Theory and the Foundations of Quantum Mechanics*. Oxford: Oxford University Press, 2013.

[109](#) Wallace, D. *The Emergent Multiverse: Quantum Theory according to the Everett Interpretation*. Oxford: Oxford University Press, 2012.

[110](#) Wiseman, H. M. "The Two Bell's Theorems of John Bell." *Journal of Physics A: Mathematical and Theoretical* 47, no. 42 (2014): 424001.

VII. Structural Realism, Scientific Realism, and Ontological Layering

[111](#) Chakravartty, A. *A Metaphysics for Scientific Realism: Knowing the Unobservable*. Cambridge: Cambridge University Press, 2007.

[112](#) Esfeld, M., and V. Lam. "Ontic Structural Realism as a Metaphysics of Objects." In *Scientific Structuralism*, edited by A. Bokulich and P. Bokulich. Dordrecht: Springer, [year/page details continue in next entries](#).

[112](#) Esfeld, M., and V. Lam. "Ontic Structural Realism as a Metaphysics of Objects." In *Scientific Structuralism*, edited by A. Bokulich and P. Bokulich, 143-159. Dordrecht: Springer, 2011.

[113](#) French, S. *The Structure of the World: Metaphysics and Representation*. Oxford: Oxford University Press, 2014.

[114](#) French, S., and J. Ladyman. "Remodelling Structural Realism: Quantum Physics and the Metaphysics of Structure." *Synthese* 136, no. 1 (2003): 31-56.

[115](#) Frigg, R., and I. Votsis. "Everything You Always Wanted to Know about Structural Realism but Were Afraid to Ask." *European Journal for Philosophy of Science* 1, no. 2 (2011): 227-276.

[116](#) Ladyman, J. "What Is Structural Realism?" *Studies in History and Philosophy of Science* 29, no. 3 (1998): 409-424.

[117](#) Ladyman, J., and D. Ross. *Every Thing Must Go: Metaphysics Naturalized*. Oxford: Oxford University Press, 2007.

[118](#) Psillos, S. *Scientific Realism: How Science Tracks Truth*. London: Routledge, 1999.

[119](#) Worrall, J. "Structural Realism: The Best of Both Worlds?" *Dialectica* 43, no. 1/2 (1989): 99-124.

VIII. Grounding, Ontological Dependence, and Metaphysical Foundations

[120](#) Bird, A. *Nature's Metaphysics: Laws and Properties*. Oxford: Oxford University Press, 2007.

[121](#) Correia, F., and B. Schnieder, eds. *Metaphysical Grounding: Understanding the Structure of Reality*. Cambridge: Cambridge University Press, 2012.

[122](#) Ellis, B. *Scientific Essentialism*. Cambridge: Cambridge University Press, 2001.

[123](#) Fine, K. "Ontological Dependence." *Proceedings of the Aristotelian Society* 95 (1995): 269-290.

[124](#) Fine, K. "The Question of Realism." *Philosophers' Imprint* 1, no. 1 (2001): 1-30.

[125](#) Heil, J. *The Universe as We Find It*. Oxford: Oxford University Press, 2012.

[126](#) Lowe, E. J. *The Four-Category Ontology: A Metaphysical Foundation for Natural Science*. Oxford: Oxford University Press, 2006.

[127](#) Mumford, S. *Laws in Nature*. London: Routledge, 2004.

[128](#) Mumford, S., and R. L. Anjum. *Getting Causes from Powers*. Oxford: Oxford University Press, 2011.

[129](#) Rosen, G. "Metaphysical Dependence: Grounding and Reduction." In *Modality: Metaphysics, Logic, and Epistemology*, edited by B. Hale and A. Hoffmann, 109-136. Oxford: Oxford University Press, 2010.

[130](#) Schaffer, J. "On What Grounds What." In *Metametaphysics: New Essays on the Foundations of Ontology*, edited by D. J. Chalmers, D. Manley, and R. Wasserman, 347-383. Oxford: Oxford University Press, 2009.

[131](#) Schaffer, J. "Monism: The Priority of the Whole." *The Philosophical Review* 119, no. 1 (2010): 31-76.

[132](#) Tahko, T. E. *An Introduction to Metametaphysics*. Cambridge: Cambridge University Press, 2015.

[133](#) Wilson, J. M. "No Work for a Theory of Grounding." *Inquiry* 57, no. 5/6 (2014): 535-579.

IX. Scientific Models, Theoretical Maturity, and Methodological Reflection

[134](#) Giere, R. N. *Explaining Science: A Cognitive Approach*. Chicago: University of Chicago Press, 1988.

[135](#) Giere, R. N. *Science without Laws*. Chicago: University of Chicago Press, 1999.

[136](#) Laudan, L. *Progress and Its Problems: Toward a Theory of Scientific Growth*. Berkeley: University of California Press, 1977.

[137](#) Longino, H. E. *Science as Social Knowledge: Values and Objectivity in Scientific Inquiry*. Princeton: Princeton University Press, 1990.

[138](#) Morgan, M. S., and M. Morrison, eds. *Models as Mediators: Perspectives on Natural and Social Science*. Cambridge: Cambridge University Press, 1999.

[139](#) Suppes, P. "A Comparison of the Meaning and Uses of Models in Mathematics and the Empirical Sciences." *Synthese* 12, no. 2/3 (1960): 287-301.

[140](#) Winsberg, E. Science in the Age of Computer Simulation. Chicago: University of Chicago Press, 2010.

X. Translations and Introductory Works for Chinese Readers

[141](#) Popper, K. The Logic of Scientific Discovery. Translated by Cha Ruqiang and Qiu Renzong. Beijing: China Academy of Art Press, 2008.

[142](#) Duhem, P. The Aim and Structure of Physical Theory. Translated by Li Xingmin. Beijing: The Commercial Press, 2011.

[143](#) Kuhn, T. S. The Structure of Scientific Revolutions. Translated by Jin Wulun and Hu Xinhe. Beijing: Peking University Press, 2012.

[144](#) van Fraassen, B. C. The Scientific Image. Translated by Zheng Xiangfu. Shanghai: Shanghai Translation Publishing House, 2002.

[145](#) Hacking, I. Representing and Intervening: Introductory Topics in the Philosophy of Natural Science. Translated by Wang Wei. Beijing: Science Press, 2011.

Appendix

Appendix A:

Criteria Table for the Three Explanatory Roles

The core conceptual framework of this book is a threefold distinction among explanatory roles: causal production, structural constraint, and permissibility. These are not three mutually exclusive types of explanation, but roles taken on by explanation in response to different dimensions of a problem. A given explanation may contain more than one of these roles, though one usually plays the dominant part. To avoid conceptual confusion in the main text, the criteria for the three are summarized here.

Explanatory Role	Core Question	Primary Object	Condition of Success	Common Miswriting
Causal production	How does a certain result occur?	Events, processes, state-transitions, difference-making relations	Provides a generative pathway and shows which factors make a difference to the occurrence of the result	Treating all explanation as if it had to be a production-chain
Structural constraint	Which paths, configurations, or states are restricted or excluded?	Boundary conditions, organizational relations, symmetries, topological structures, realizational conditions	Explains why certain changes are unreachable, infeasible, unstable, or inconsistent	Treating limiting conditions as mere background
Permissibility	Which patterns can stably exist?	Stable patterns, realizable configurations, spaces of possibility, attractor-like structures	Explains why a certain kind of form can be maintained, recur, or remain identifiable under given constraints	Rewriting stable patterns as single event-results

It must be emphasized that although structural constraint and permissibility are closely related, they should not be completely conflated. Structural constraint focuses on exclusion, limitation, and boundary. Permissibility focuses on what kinds of patterns can stably hold within those boundaries. The former answers “what

cannot happen, or cannot stably happen"; the latter answers "what can take stable form in reality."

Appendix B:

Roadmap of the Book's Overall Argument

The overall argument of this book can be summarized in four stages.

The first stage is problem-diagnosis. Chapters 1 and 2 argue that a number of disputes in contemporary philosophy of science, though distributed across different domains, share a common structure of explanatory contraction. Emergence, higher-level explanation, mechanistic explanation, quantum entanglement, and unknown phenomena are often rewritten as problems about causal production. This tendency is what the book calls explanatory monism.

The second stage is conceptual reconstruction. Chapter 3 introduces the three explanatory roles of causal production, structural constraint, and permissibility. Causal production answers how a result occurs; structural constraint answers which paths are restricted or excluded; permissibility answers which patterns can stably exist. The aim of this framework is not to replace causal explanation with a new type of explanation, but to dissolve the falsely imposed competitive relation among different explanatory roles.

The third stage is application to specific problems. Chapter 4 rewrites emergence as a problem of structural novelty. Chapter 5 releases higher-level explanation from causal competition. Chapter 6 defends the scientific standing of non-predictive frameworks. Chapter 7 rewrites the demand for nonlocal causation in quantum entanglement as a problem of structural inseparability. Chapter 8, in turn, transforms the ontological anxiety provoked by unknown phenomena into a problem of explanatory discipline.

The fourth stage is ontological synthesis. Chapter 9 argues that if structural constraint and permissibility display stability, indispensability, counterfactual constraint, and realizational support across multiple theoretical contexts, then they

should not be downgraded to mere pragmatic convenience, but should support a restrained layered structural realism. Chapter 10 then reviews the book's contributions and limitations, showing that its aim is not ontological inflation, but a re-stratification of explanatory roles and ontological commitments.

Appendix C:

The Book's Relation to Major Theoretical Traditions

Theoretical Tradition	Core Contribution	What the Book Takes Up	What the Book Revises
Kim-style exclusion arguments	Reveals the pressure between higher-level causation and physical closure	Acknowledges that higher-level explanation must confront causal closure and exclusion-pressure	Rejects the inference from causal sufficiency directly to explanatory sufficiency
Mechanistic explanation	Emphasizes how parts, activities, and organization produce phenomena	Recognizes the central place of productive mechanistic explanation in science	Rejects expanding mechanistic explanation into the sole paradigm for all explanation
Woodwardian interventionism	Characterizes causal explanation in terms of counterfactuals and intervention-relations	Draws on its fine-grained analysis of difference-making	Holds that not all explanation is centered on intervention-dependence
Theories of non-causal explanation	Show that mathematics, symmetry, constraint, and the like can also explain	Provide background resources for structural constraint and permissibility	Attempts to integrate scattered cases into a unified framework of explanatory roles
Structural realism	Emphasizes the place of structure in scientific realism	Supports recognition of the reality of structural stability	Places greater emphasis on a restrained inference from explanatory role to ontological recognition
Grounding theory	Discusses ontological dependence-relations among facts	Helps explain the dependence of higher-level facts on lower-level realization	Does not reduce higher-level reality entirely to grounding-relations

Appendix D:

Clarifications of Several Concepts in the Chapter on Quantum Entanglement

Chapter 7 of this book does not attempt to propose a new interpretation of quantum mechanics, nor does it seek to render a final verdict among Bohmian theory, the many-worlds interpretation, relational quantum mechanics, QBism, and other interpretive approaches. Its aim is more limited: to analyze why quantum entanglement is so often rapidly translated into a demand for nonlocal causal influence, and to indicate the possibility of an alternative structural rewriting.

To avoid misunderstanding, the relevant concepts are distinguished as follows.

Nonlocality: Usually refers to the fact that certain correlations in quantum theory cannot be explained by classical local hidden-variable models. It should not be simply equated with controllable superluminal signal-transmission.

No-signaling: Means that the correlations exhibited in quantum entanglement cannot be used for controllable superluminal information-transmission. No-signaling does not remove the philosophical depth of entanglement, but it does limit the way it may be understood as an ordinary transmission-mechanism.

Nonseparability: Refers to the fact that the state of a composite system cannot be reduced to a simple combination of the independent states of its constituent parts. Entangled states challenge the classical part-whole picture precisely in this sense.

Structural inseparability: The concept used in this book. It means that the joint outcome space in an entangled system cannot be understood as a simple product of mutually independent local spaces of permissibility. This concept does not deny the importance of Bell's theorem. Rather, it emphasizes that the explanatory pressure raised by entanglement need not first be a problem of causal

transmission, but may instead be understood as a problem of structural constraint within a joint space of possibility.

Appendix E:

Concise Table of the Layered Structure of Restrained Layered Structural Realism

(with “Radiance-Will Ontology” as a secondary designation)

The restrained layered structural realism proposed in Chapter 9 of this book, also rhetorically referred to as “Radiance-Will Ontology,” is not meant to add several independent world-layers to reality, nor to claim that each layer is a self-sufficient entity. What is meant by “levels” here are distinctions in the different functional aspects of reality’s organization. To avoid misunderstanding, they are summarized as follows.

Level	Name	Basic Meaning	Main Function
L0	Unmanifest base layer	The background of real potentiality not yet entered into stable manifest form	Indicates the supporting background prior to manifestation
L1	Will layer	Organizational directedness in a non-psychological sense	Indicates that reality is not arbitrarily open, but has structural direction
L2	Radiance layer	Stable manifestness	Indicates that a certain pattern continues to manifest at a particular scale
L3	Topological layer	Structure of relations, boundaries, and connections	Indicates the organizational form of structural constraint
L4	Manifested-structure layer	Stable patterns identifiable within reality	Indicates higher-level facts and forms of organization
L5	Structural logic and convergence layer	Reflexive integration of structural relations across different levels	Indicates theory’s conceptual grasp of reality’s structure

It must be emphasized that the above levels are not an inventory of entities, but distinctions among organizational functions. The book’s ontological claim is a restrained one: only when some pattern satisfies stability, indispensability, counterfactual constraint, and realizational support does it qualify for ontological recognition.

Appendix F: The Relation Between Restrained Layered Structural Realism and Other Ontological Positions

Comparison Target	Common Ground	Key Difference	The Book's Position
Reductionism / physicalism	Acknowledges that higher levels depend on lower-level realization	Rejects equating dependence with elimination	An anti-flattening dependence view
Strong emergentism	Acknowledges that higher-level facts are not illusions	Does not introduce new causal powers	A non-inflationary realism about higher-level reality
Structural realism	Acknowledges that structure has real standing	Starts from explanatory roles rather than from scientific realism as such	An explanation-oriented structural realism
Process ontology	Emphasizes that reality is not a static heap of objects	Gives greater emphasis to the way process is organized by boundaries and permissibility	A structured process view
Grounding theory	Focuses on layered dependence-relations	Adds criteria of explanatory indispensability and stable manifestation	A level-realism with explanatory criteria
Crude pluralism	Acknowledges multiple levels and multiple scales	Does not accept that any useful distinction is automatically real	A threshold-based multilayer realism
Substance ontology	Shares concern with the structure of existence	Does not treat object-inventory as central	A structurally organized ontology

Glossary

Explanatory monism

The tendency to understand explanation as a single task, especially by ultimately contracting it into an account of causal production. It may appear in the forms of evaluative monism, role monism, and competitive monism. Its problem does not lie in valuing causal explanation, but in elevating causal production into the sole standard by which all explanation is judged.

Causal production

One of the explanatory roles. It answers how a certain event, state, or result is brought into reality through a concrete process. Its core includes generative pathways and difference-making relations. Causal production is not identical with micro-reduction; higher-level processes may also be causally productive.

Structural constraint

One of the explanatory roles. It answers which paths, configurations, states, or forms of change are restricted or excluded by a system's structure. Structural constraint does not directly "push" a result into occurrence, but specifies the boundaries within which generative processes may unfold.

Permissibility

One of the explanatory roles. It answers which patterns, configurations, or forms of organization can stably exist under given structural constraint. Permissibility is not abstract logical possibility, but realizability and maintainability under specific structural conditions.

Multi-role conception of explanation

The explanatory framework proposed in this book. Its basic claim is that explanation is not a single productive task, but includes at least three mutually irreducible roles: causal production, structural constraint, and permissibility.

Explanatory miswriting

A case in which one kind of explanatory question is wrongly rewritten as another. For example, a question about structural stability may be rewritten as a question about new causal powers, or a question about the structure of a joint outcome space may be rewritten as a question about distant transmission.

Structural novelty

The book's rewritten concept for the problem of emergence. It means that the novelty of higher-level phenomena lies not in new causal powers, but in the system's exhibiting new organizational relations, stable patterns, and spaces of permissibility.

Higher-level explanation

Explanation offered at a scale of organization above lower-level physical realization, concerning patterns, functions, stability, relations of constraint, or system behavior. Higher-level explanation need not justify itself by competing with lower-level causal explanation for the same place.

Causal exclusion argument

A tradition of argument associated especially with Jægwon Kim. Its basic pressure is that if lower-level physical causes already sufficiently explain a result, then higher-level causes seem either redundant or to lead to overdetermination. This book argues that the central problem with exclusion arguments is their rewriting of causal sufficiency as explanatory sufficiency.

Mechanistic explanation

A form of explanation that shows how a phenomenon is produced through parts, activities, and their organization. The book acknowledges the importance of mechanistic explanation, but rejects extending it into the sole paradigm for all explanation.

Non-predictive framework

A theoretical framework whose value lies not primarily in generating new predictions directly, but in reorganizing problem-space, clarifying explanatory roles, limiting illegitimate ontological inferences, and providing methodological discipline.

Ontological robustness

The capacity, when faced with unknown phenomena, anomalous experiential objects, or theoretical pressure, not to immediately resort to new entities, new forces, or new levels, but first to distinguish whether the problem lies at the causal, structural, model, or ontological level.

Ontological stability

The capacity of a set of ontological commitments to preserve structural layering, conceptual boundaries, and the order of revision when facing new phenomena and explanatory pressure, without undergoing uncontrolled inflation. It differs from ontological robustness: ontological stability emphasizes the structural stability of the resulting state of commitments, whereas ontological robustness emphasizes resilience under pressure and discipline in revision.

Ontological inflation

The premature introduction of new entities, new causal powers, or new ontological levels before explanatory difficulty has been sufficiently analyzed. It often results from mistaking explanatory insufficiency for insufficiency in what exists.

Structural inseparability

The concept used in this book to rewrite the problem of quantum entanglement. It means that a joint outcome space cannot be decomposed into a simple combination of independent local spaces of permissibility. It is not a new physical theory, but a way of rewriting the explanatory problem.

Restrained layered structural realism

The more neutral ontological name introduced in Chapter 9. It holds that reality possesses multiple organizational scales and structural levels, but this recognition is restricted: only structural patterns that satisfy stability, indispensability, counterfactual constraint, and realizational support should receive ontological recognition.

Radiance

In this book, “radiance” does not refer to physical light or mystical aura, but to the stable manifestation of reality’s structure at a certain organizational scale. When a pattern has “radiance,” it means that it is not an arbitrary naming, but something that stably manifests within reality’s organization and becomes an unavoidable object of explanation.

Will

In this book, “will” does not refer to subjective psychology, personal intention, or supernatural purpose, but to directedness within reality’s organization. It indicates that reality does not arbitrarily permit every configuration, but exhibits structural direction through constraint, exclusion, stability, and conditions of realization.

Radiance-Will Ontology

The book’s rhetorically charged secondary name for restrained layered structural realism. It seeks to characterize the stable manifestation and organizational directionality found in reality, and to explain why higher-level facts, structural constraint, and permissibility require neither extra causal powers nor reduction to mere descriptive convenience.

Realizational support

The condition that a higher-level pattern or structural fact is not free-floating, but must be supported by some real organizational structure. Realizational support is used to prevent arbitrary conceptual distinctions from being ontologized.

Counterfactual constraint

The fact that a structure or pattern not only holds in the actual case, but also supports judgments about what changes would or would not be possible if conditions were different. It is one of the important conditions for moving from explanatory role to ontological recognition.

Postscript

The original impulse for writing this book came from a recurring question: why do so many difficulties in philosophy and in scientific theory ultimately get reformulated as causal questions? Whether the issue is emergence, higher-level explanation, mechanistic analysis, quantum entanglement, or the ontological anxiety provoked by unknown phenomena, discussion so often returns to the same direction: how exactly some result is produced, whether some level possesses an independent causal power, whether some anomaly requires a new entity or a new mode of operation. These are, of course, important questions. But they are not always the only important ones.

What this book has tried to do is to secure a clear place for another class of explanatory tasks without denying the importance of causal explanation. In many cases, what we truly need to understand is not how a result is additionally produced, but why certain paths are excluded, why certain patterns can stably exist, and why certain higher-level structures retain explanatory standing across multiple conditions of realization. It is around this judgment that the book gradually developed its threefold explanatory framework of causal production, structural constraint, and permissibility.

This book did not begin as a “book” in any straightforward sense. It began as a set of draft essays on explanation, emergence, and ontological stability. I tried at various points to compress its central concerns into the structure of a single article, and I went through repeated rounds of submission, revision, rejection, and rewriting around related themes. The process was not smooth, but it gradually made one fact clearer: the problem that concerned me was not a local question

that could be quickly settled within a short paper, but a set of structural problems bound up with one another.

The article form demands concentrated claims, limited length, clear boundaries, and quickly gathered conclusions. The problems addressed in this book, however, kept revealing deeper internal connections. Once one begins discussing explanatory monism, one cannot avoid the differentiation of explanatory roles; once one discusses the differentiation of explanatory roles, one cannot avoid emergence and higher-level explanation; once one discusses emergence and higher-level explanation, one is led further into the question of ontological stability; and once one enters the problem of ontological stability, the relations among structure, manifestation, levels, and reality inevitably reappear. In this way, what had originally been a scattered set of article-projects gradually became, through repeated compression and re-expansion, the present work.

For that reason, this book is not simply an expanded version of a single paper, nor merely a substitute response to unsuccessful submissions. On the contrary, it is the result of something that took shape only gradually through repeated attempts: when a problem cannot be adequately unfolded within a short argument, what it may truly require is not greater compression, but a more complete structure. The points that once seemed to obstruct the writing later became the clues by which the book's structure was generated. It was precisely in the space between these repeatedly rewritten problems that the book's ten-chapter architecture gradually emerged.

During the writing process, the greatest difficulty was not introducing concepts, but limiting them. The more a theoretical framework tries to move across multiple domains, the more easily it risks overexpansion. This book deals with emergence, higher-level explanation, mechanistic explanation, non-predictive frameworks, quantum entanglement, unknown phenomena, and ontological layering; the breadth of that range carries its own risk. In order to avoid turning the

framework into a theory of everything, the book repeatedly emphasizes that it does not seek to propose a new empirical scientific theory, nor to replace concrete scientific research. Its real concern is with the relation among explanatory form, problem structure, and ontological commitment.

Accordingly, many of the book's claims should be understood in a limited sense. Structural constraint is not a mysterious force. Permissibility is not arbitrary possibility. Radiance-Will Ontology is not an attempt to add supernatural purpose or personified will to the world. All of these are intended to express a single more basic thought: reality is not merely the flat unfolding of a single causal chain, but appears across different scales as the joint organization of generative processes, structural boundaries, stable patterns, and conditions of manifestation.

This book could not have been completed without many existing intellectual resources. Work on causal explanation, mechanistic explanation, higher-level causation, non-causal explanation, structural realism, and the foundations of quantum theory provided the necessary background. Although this book attempts to reorganize these traditions, that attempt is not a creation out of nothing, but a further movement within an already existing field of problems and literature. If there is anything of value in the book, it should first be credited to the depth of those problems themselves; if there are shortcomings, they belong entirely to the limitations of the author's own judgment, expression, and argumentative ability.

Much in this book remains unfinished. As an open explanatory framework, it necessarily leaves many points in need of further formalization. The three explanatory roles can be formalized more fully. The discussion of quantum entanglement in Chapter 7 still needs sustained engagement with more specialized work in the foundations of quantum theory. The ontological articulation proposed in Chapter 9 also requires more detailed comparison with traditions such as grounding, realization, and structural realism. This book is therefore better

understood as the proposal of a systematic framework than as the final resolution of all related questions.

The auxiliary use of artificial intelligence tools in the writing of this book has already been specifically addressed in the “Note to the Reader: On the Book’s Mode of Composition and Responsibility,” and need not be repeated here. It must nonetheless be emphasized again that all of the book’s views, conceptual choices, argumentative arrangements, final judgments, and possible errors are the author’s own responsibility. Any tool can participate only in organization, feedback, and expression; it cannot replace the author’s responsibility for the content of thought.

If this book is able to help readers slow down, even slightly, when confronting theoretical difficulty; if it can make them less eager to turn every problem into a problem of causal production, less eager to introduce new ontological commitments at the first sign of explanatory insufficiency, and more ready to ask instead whether the issue before them is one of generation, constraint, or permissibility, then its principal aim will already have been achieved.

Finally, I am grateful to all readers willing to engage this work with patience. The significance of philosophical writing does not lie only in offering conclusions, but also in offering a way of asking questions anew. If this book has value, it may lie precisely here: in its attempt to reopen some questions that had seemed already fixed, and to make the relation among explanation, structure, and being thinkable again.

About the Author

The author has long engaged with issues in the philosophy of science, theories of explanation, the problem of emergence, higher-level explanation, and ontological structure, with particular concern for the relations among explanatory form, problem structure, and ontological commitment in contemporary theoretical discussions.

This book, *From Causal Production to Structural Constraint: Explanatory Roles, Emergence, and Ontological Stability*, presents in a concentrated form the author's systematic reflections in recent years on the differentiation of explanatory roles. Without denying the central place of causal explanation, the author seeks to distinguish anew the three explanatory roles of causal production, structural constraint, and permissibility, and on that basis to analyze emergence, higher-level explanation, quantum entanglement, unknown phenomena, and ontological stability.

The author's aim in writing is not to propose an empirical model that competes with specific scientific theories, but to clarify the forms of questioning at work in theoretical discussion, to reveal the explanatory miswritings underlying certain long-standing debates, and to explore a more layered expression of the philosophy of science and ontology.